

# Machine Learning and Optimization Theory

Lecturer: Göktuğ Güvercin

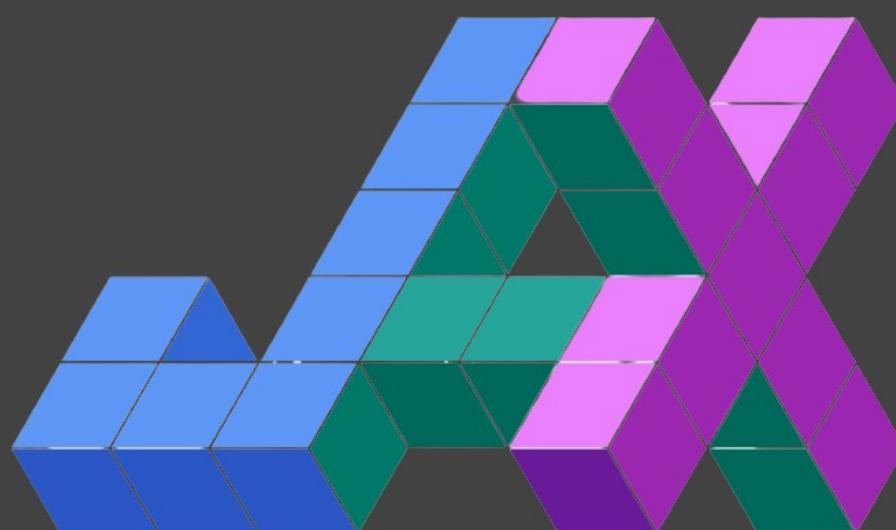
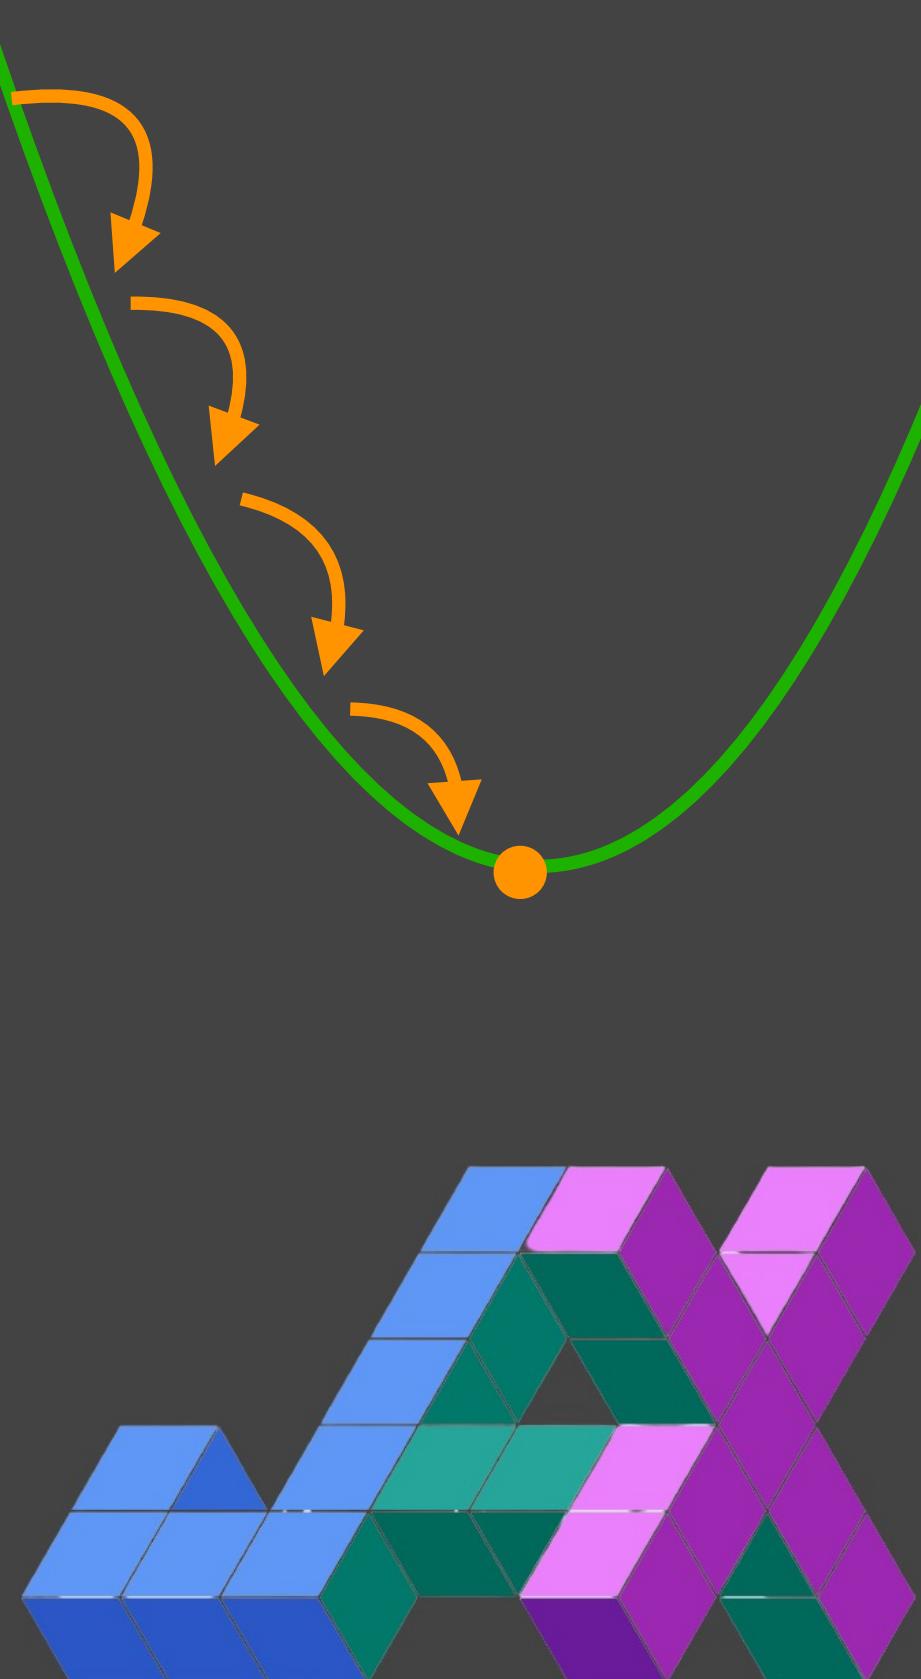
Organizer: BASIRA-CENTER

Program: 100TUNAI



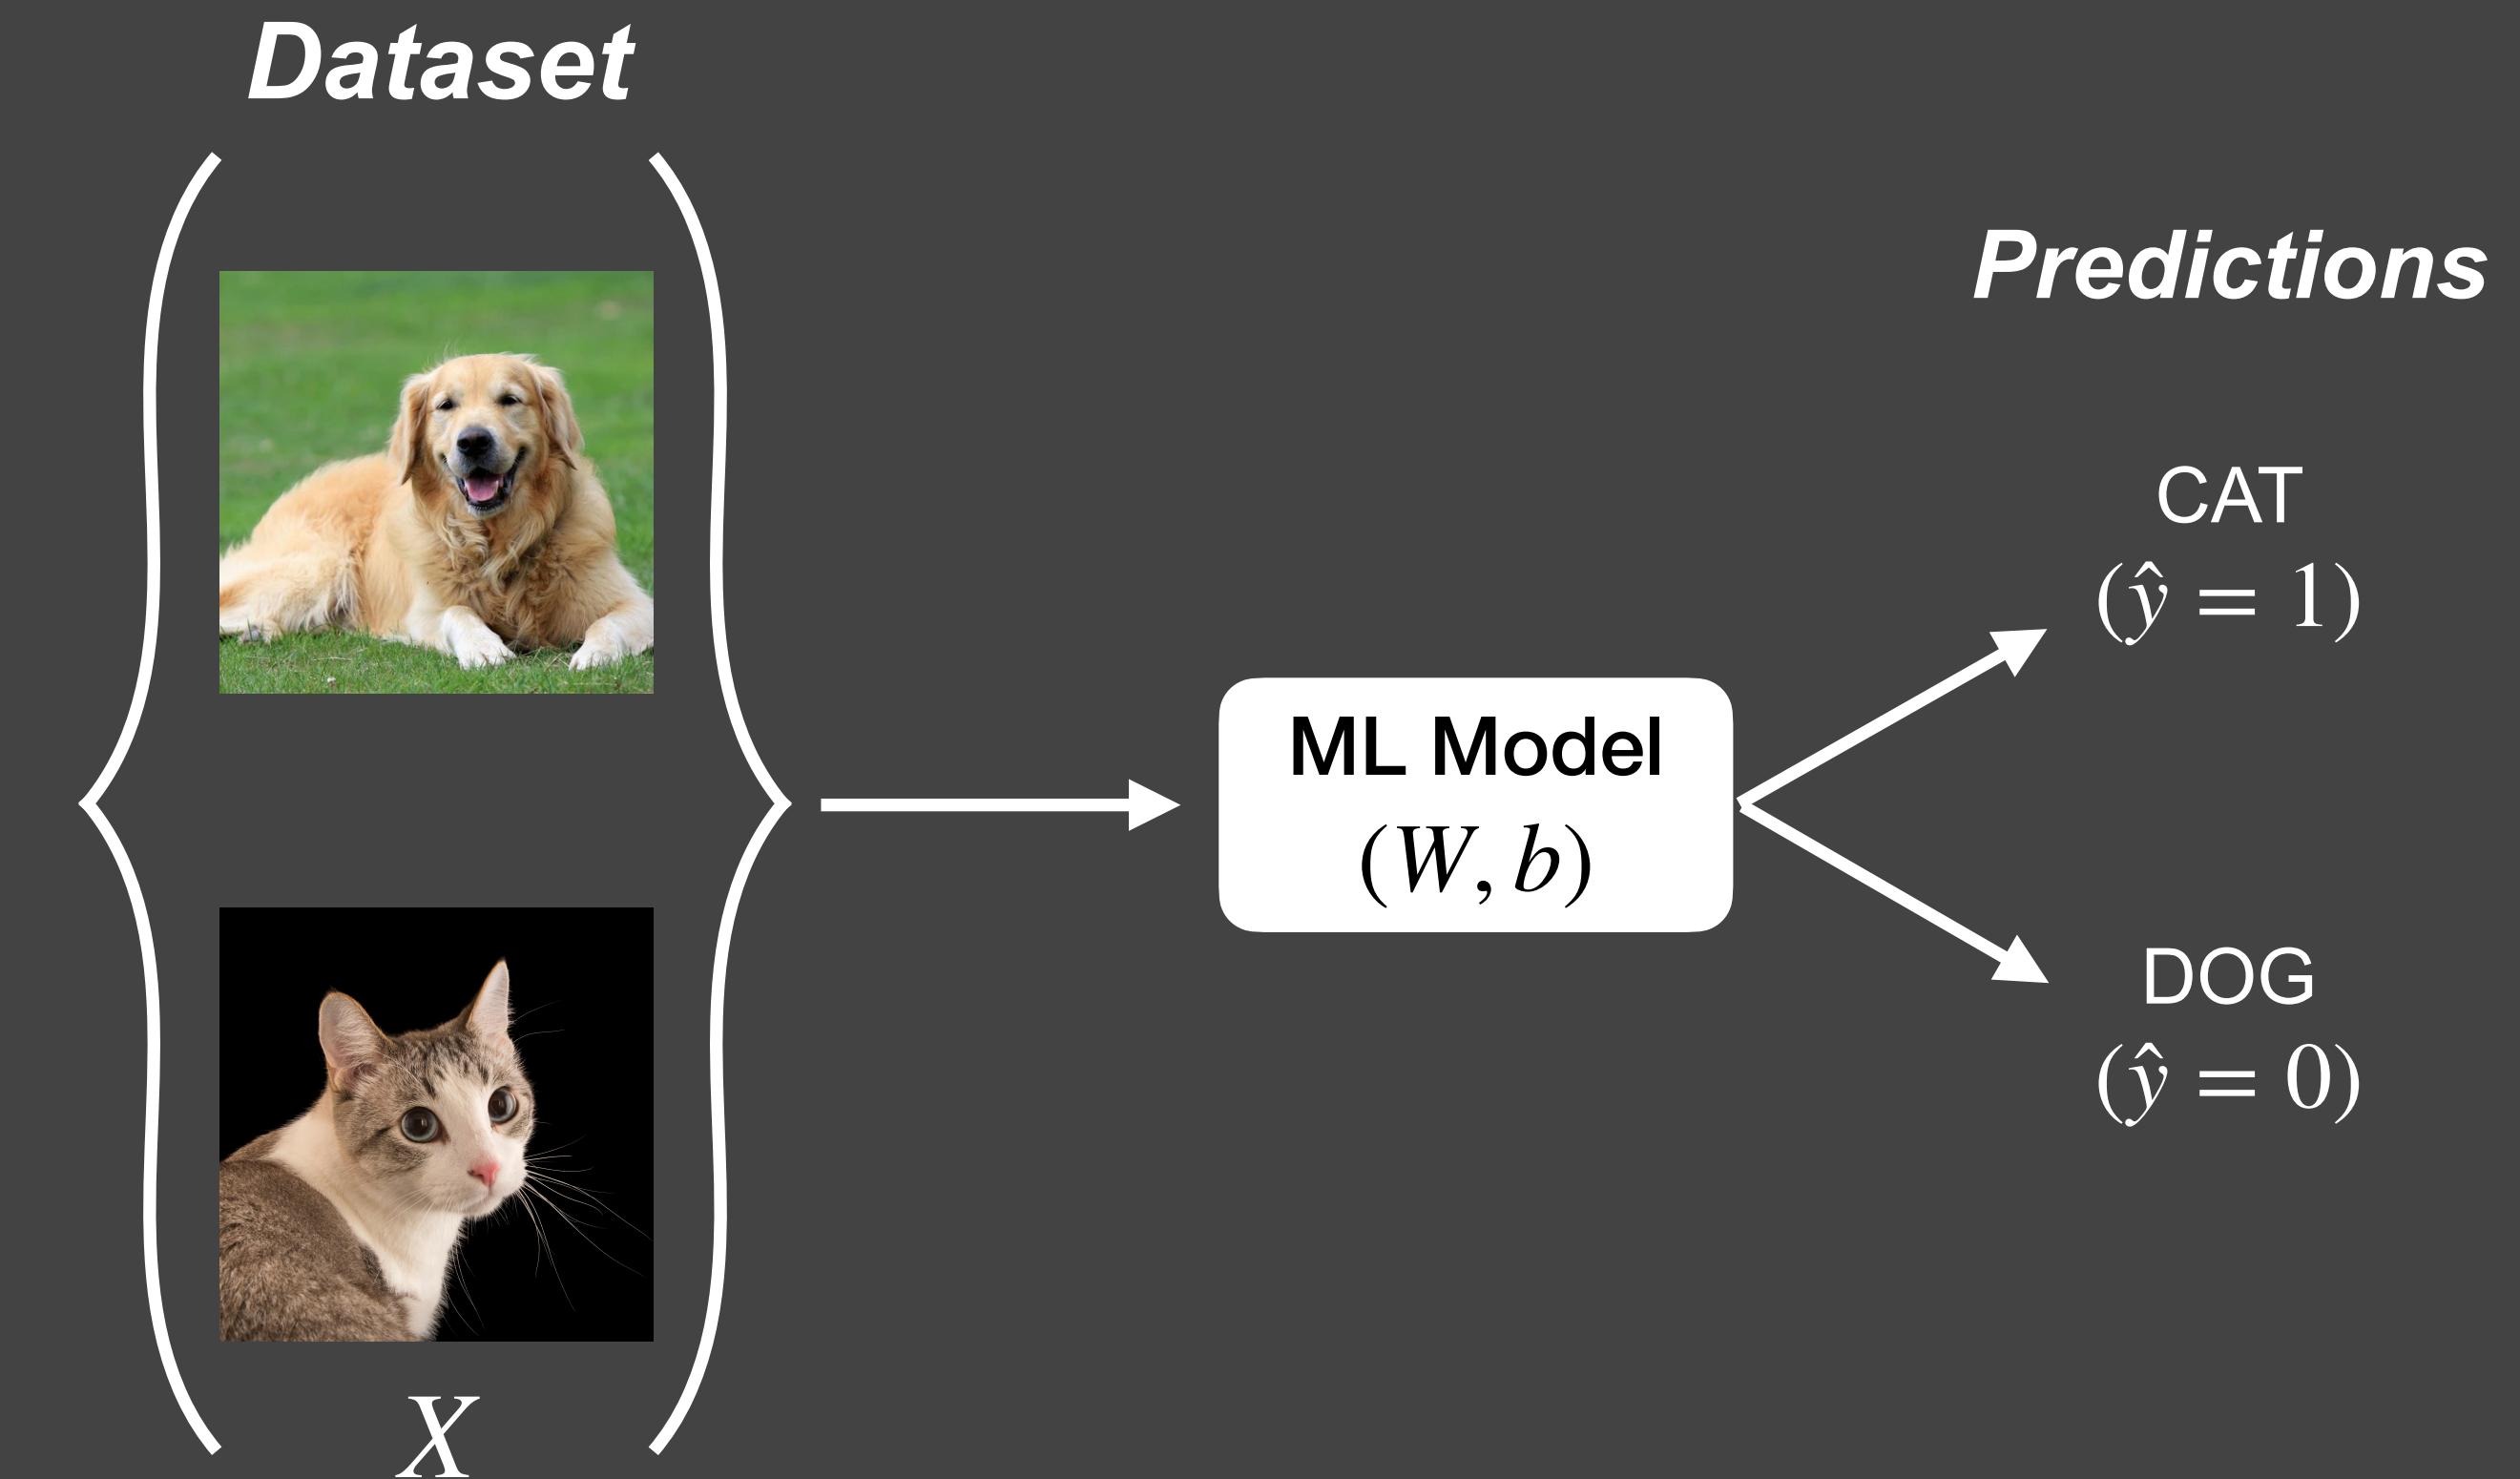
# Table of Contents

- Optimization Theory
  - Taxonomy of Machine Learning
  - Data Notations
  - Model Notations
  - Stationary Points
  - Differentiation and Gradients
  - Gradient Descent
  - Convexity of Functions
- JAX
  - Fundamentals
  - Numpy
  - Composable Function Transformations



# Optimization Theory - Taxonomy of Machine Learning

- ML models to learn predictive tasks
  - House price prediction (*regression*)
  - Cat-Dog classification (*classification*)
- Requirement of dataset
  - Data samples  $X$
  - Ground-truth labels  $y$
- Optimization of model parameters
  - Weights  $W$
  - Biases  $b$



# Optimization Theory - Taxonomy of Machine Learning

- Loss function for ML models
  - Performance measure for learning models
  - Comparison of model predictions with ground-truth labels
  - Quantification similarity between predictions and ground-truth labels

# Optimization Theory - Taxonomy of Machine Learning

- Loss function for ML models
  - Performance measure for learning models
  - Comparison of model predictions with ground-truth labels
  - Quantification similarity between predictions and ground-truth labels

	Model Predictions	Ground-Truth Labels	Quality of Predictions	Loss Function
Model 1	18.2	18.4	Good	Low
Model 2	13.8	18.4	Bad	High

# Optimization Theory - Taxonomy of Machine Learning

- Loss function for ML models
  - Performance measure for learning models
  - Comparison of model predictions with ground-truth labels
  - Quantification similarity between predictions and ground-truth labels
- Our purpose
  - Increasing quality of model predictions
  - Decreasing error rate (cost) of model predictions

	Model Predictions	Ground-Truth Labels	Quality of Predictions	Loss Function
Model 1	18.2	18.4	Good	Low
Model 2	13.8	18.4	Bad	High

# Optimization Theory - Taxonomy of Machine Learning

- Loss function for ML models
  - Performance measure for learning models
  - Comparison of model predictions with ground-truth labels
  - Quantification similarity between predictions and ground-truth labels
- Our purpose
  - Increasing quality of model predictions
  - Decreasing error rate (cost) of model predictions

	Model Predictions	Ground-Truth Labels	Quality of Predictions	Loss Function
Model 1	18.2	18.4	Good	Low
Model 2	13.8	18.4	Bad	High

# Optimization Theory - Taxonomy of Machine Learning

## *Dataset*



$X$

# Optimization Theory - Taxonomy of Machine Learning

*Dataset*



$X$



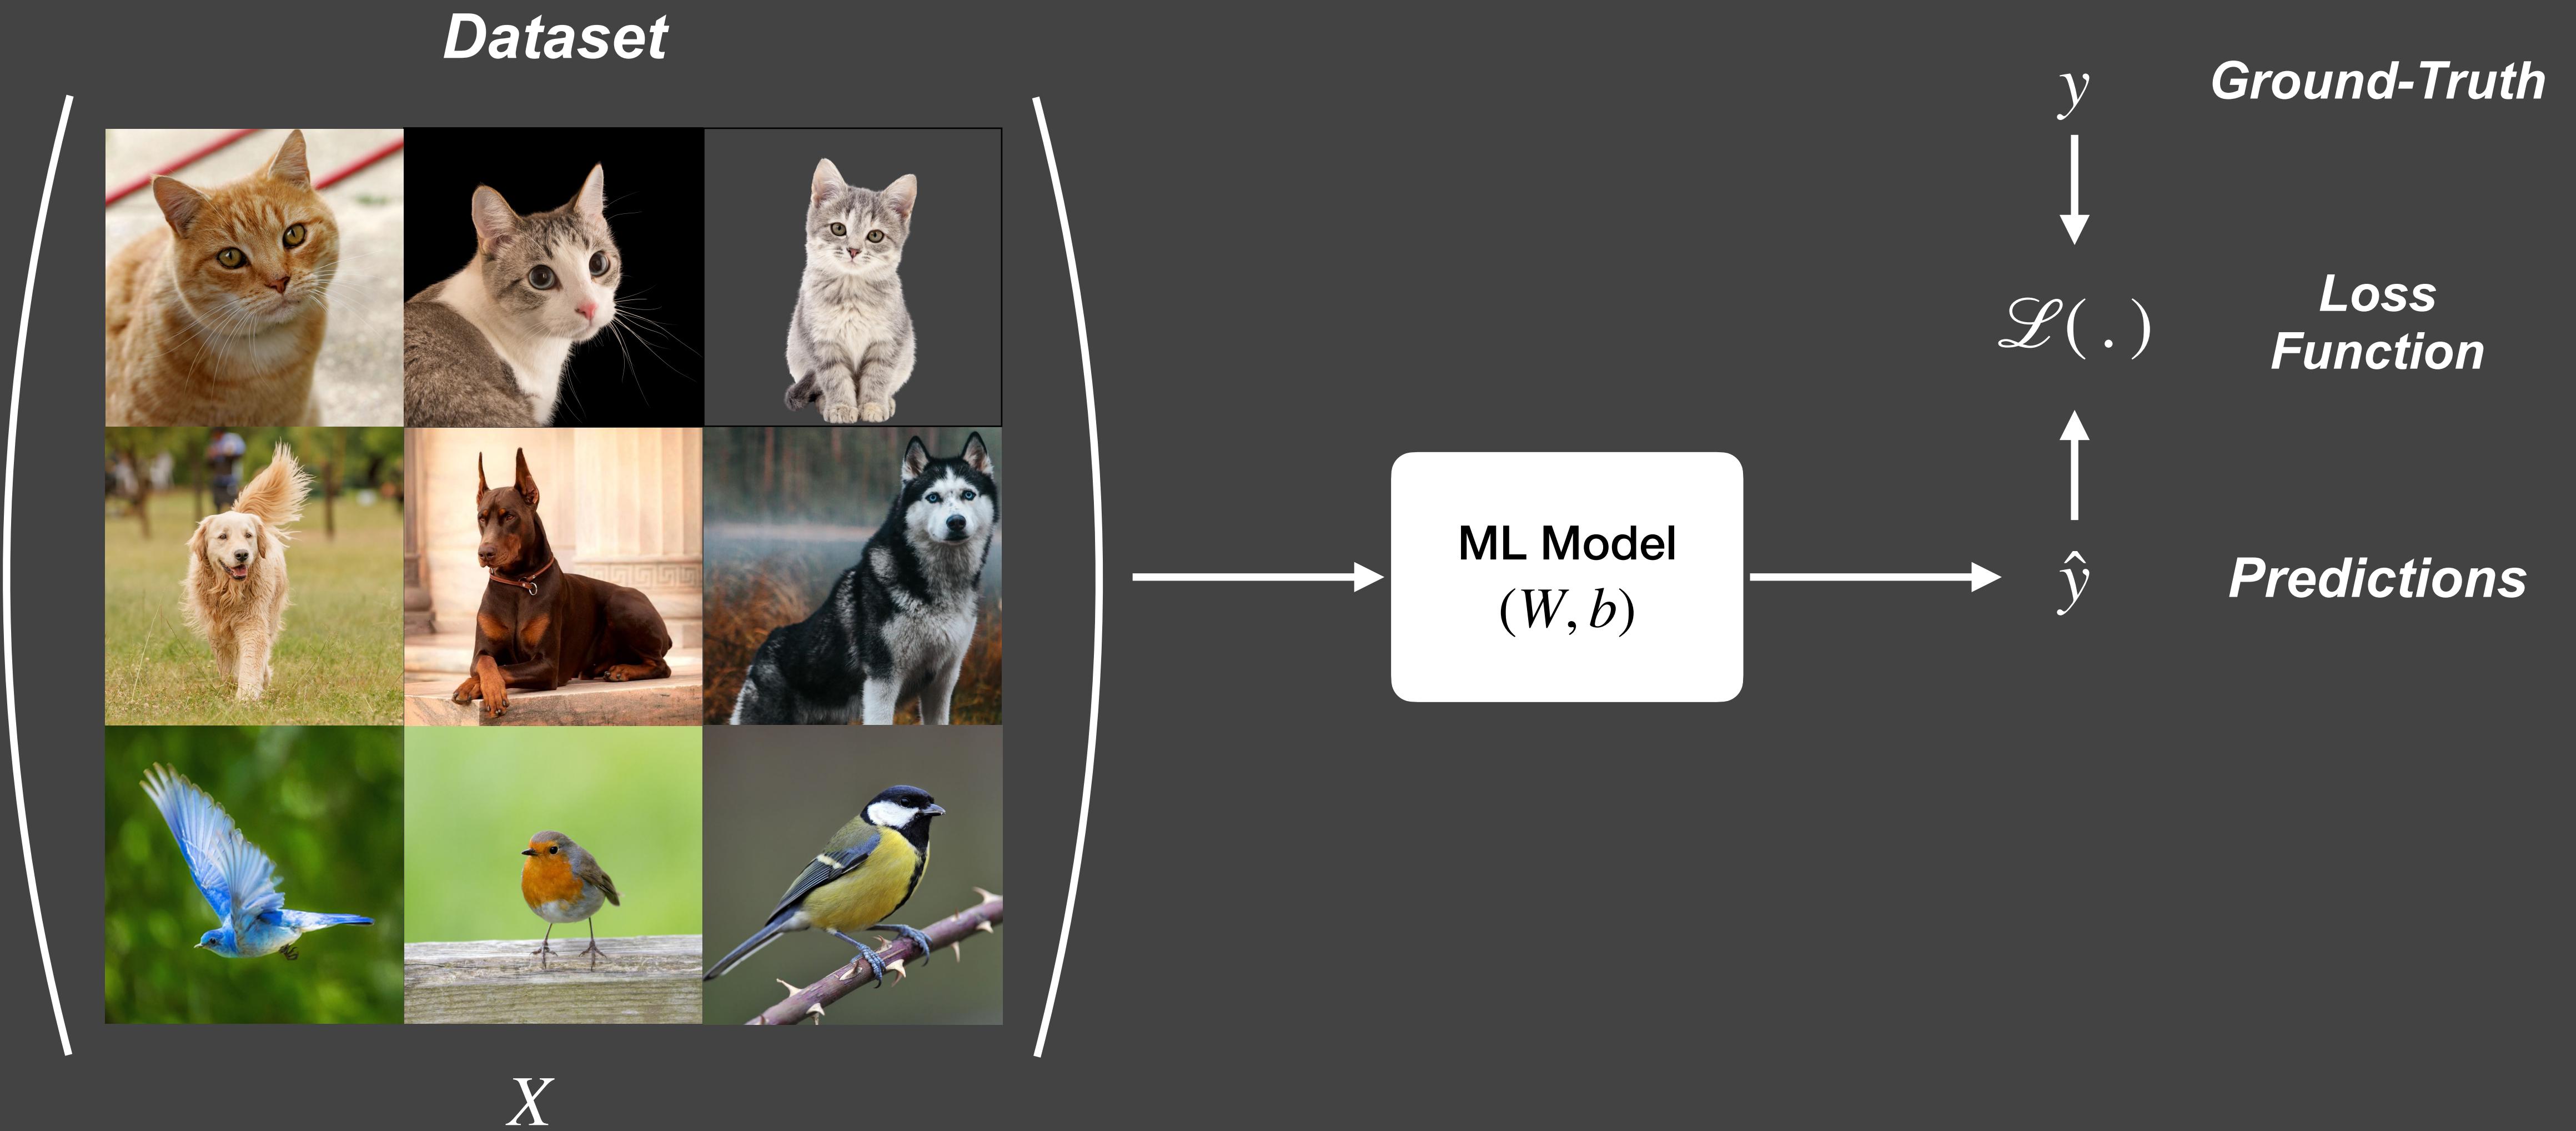
ML Model  
 $(W, b)$



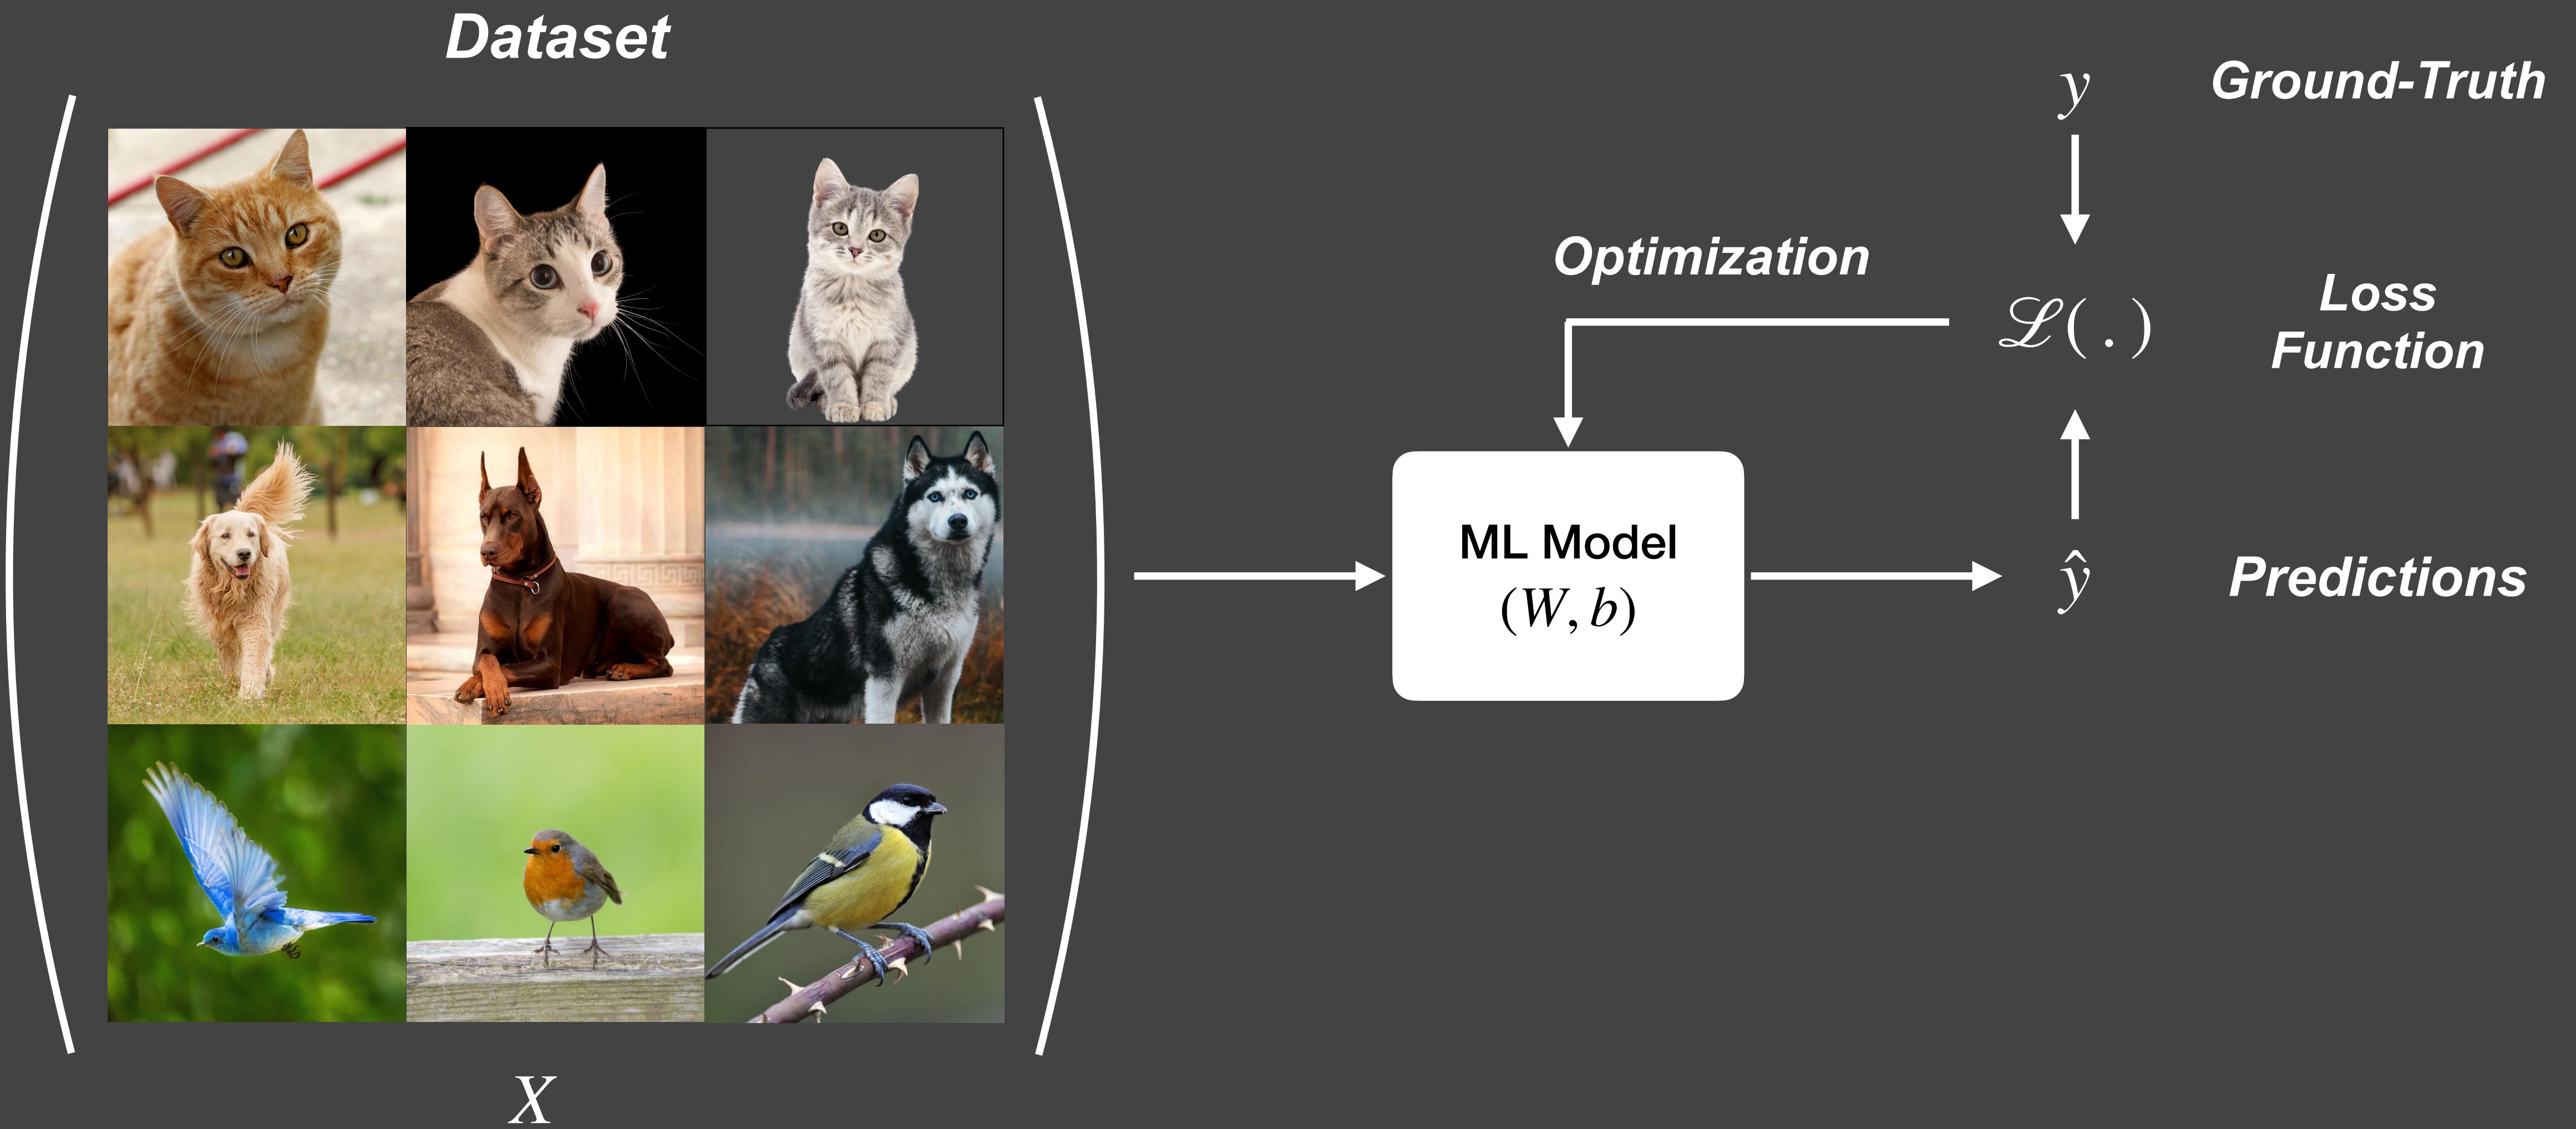
$\hat{y}$

*Predictions*

# Optimization Theory - Taxonomy of Machine Learning



# Optimization Theory - Taxonomy of Machine Learning



# Optimization Theory - Taxonomy of Machine Learning

## Dataset

- Resource for updating and optimizing model parameters
- Larger dataset, better model

## Model and Loss Functions

- Model: Mathematical functions to learn relationship between features and labels
- Loss Function: Mathematical functions to compare predictions and labels

## Optimization Algorithm

- Updating and optimizing model parameters

# Optimization Theory - Taxonomy of Machine Learning

## Dataset

- Resource for updating and optimizing model parameters
- Larger dataset, better model

## Model and Loss Functions

- Model: Mathematical functions to learn relationship between features and labels
- Loss Function: Mathematical functions to compare predictions and labels

## Optimization Algorithm

- Updating and optimizing model parameters

# Optimization Theory - Data Notations

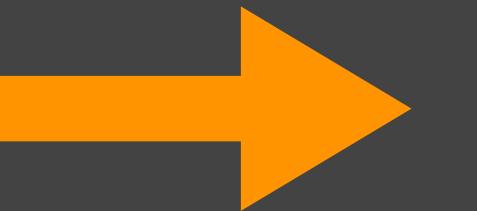
- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$

# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$

# Optimization Theory - Data Notations

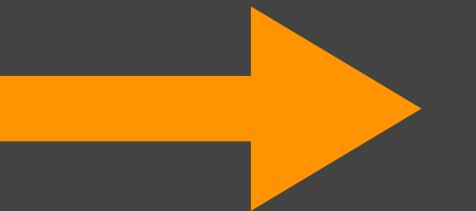
- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$



$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$

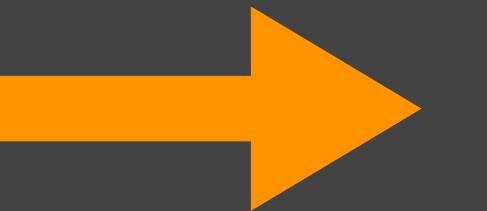


$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

$x^i \rightarrow \text{Sample index}$   
 $x_j \rightarrow \text{Feature index}$

# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$

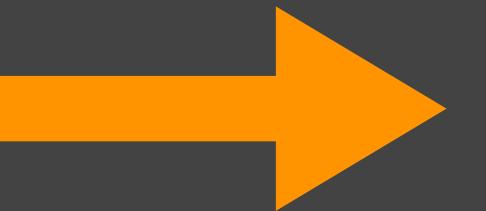


$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

- Dataset  $D = (X, Y)$ 
  - Matrix of  $N$  data samples  $X$
  - Vector of  $N$  labels  $Y$

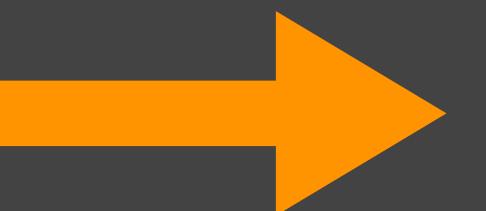
# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$



$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

- Dataset  $D = (X, Y)$ 
  - Matrix of  $N$  data samples  $X$
  - Vector of  $N$  labels  $Y$

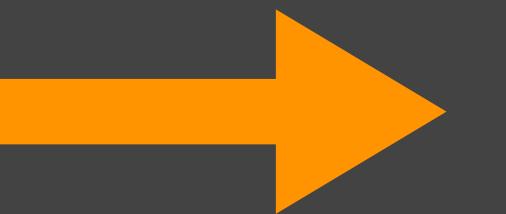


$$D = \{(x^1, y^1), (x^2, y^2), \dots, (x^N, y^N)\}$$

$$D = (X, Y) = \left( \begin{bmatrix} x_1^1 & x_2^1 & \cdots & x_d^1 \\ x_1^2 & x_2^2 & \cdots & x_d^2 \\ \vdots & \vdots & \ddots & \vdots \\ x_1^N & x_2^N & \cdots & x_d^N \end{bmatrix}, \begin{bmatrix} y^1 \\ y^2 \\ \vdots \\ y^N \end{bmatrix} \right)$$

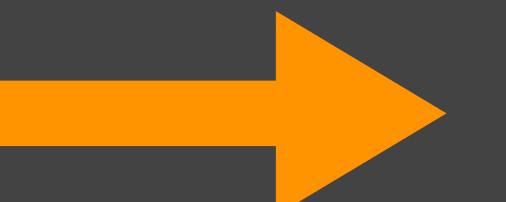
# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$



$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

- Dataset  $D = (X, Y)$ 
  - Matrix of  $N$  data samples  $X$
  - Vector of  $N$  labels  $Y$

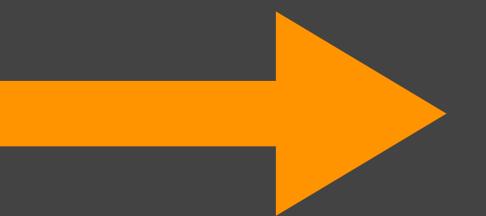


$$D = \{(x^1, y^1), (x^2, y^2), \dots, (x^N, y^N)\}$$

$$D = (X, Y) = \left( \begin{array}{cccc} x_1^1 & x_2^1 & \cdots & x_d^1 \\ x_1^2 & x_2^2 & \cdots & x_d^2 \\ \vdots & \vdots & \ddots & \vdots \\ x_1^N & x_2^N & \cdots & x_d^N \end{array} \right), \begin{bmatrix} y^1 \\ y^2 \\ \vdots \\ y^N \end{bmatrix}$$

# Optimization Theory - Data Notations

- Data sample  $(x^i, y^i)$ 
  - Feature vector of  $d$  components  $x^i$
  - Scalar ground-truth label  $y^i$



$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \rightarrow y$$

- Dataset  $D = (X, Y)$ 
  - Matrix of  $N$  data samples  $X$
  - Vector of  $N$  labels  $Y$



$$D = \{(x^1, y^1), (x^2, y^2), \dots, (x^N, y^N)\}$$

$$D = (X, Y) = \left( \begin{bmatrix} x_1^1 & x_2^1 & \cdots & x_d^1 \\ x_1^2 & x_2^2 & \cdots & x_d^2 \\ \vdots & \vdots & \ddots & \vdots \\ x_1^N & x_2^N & \cdots & x_d^N \end{bmatrix}, \begin{bmatrix} y^1 \\ y^2 \\ \vdots \\ y^N \end{bmatrix} \right)$$

# Optimization Theory - Taxonomy of Machine Learning

## Dataset

- Resource for updating and optimizing model parameters
- Larger dataset, better model

## Model and Loss Functions

- Model: Mathematical functions to learn relationship between features and labels
- Loss Function: Mathematical functions to compare predictions and labels

## Optimization Algorithm

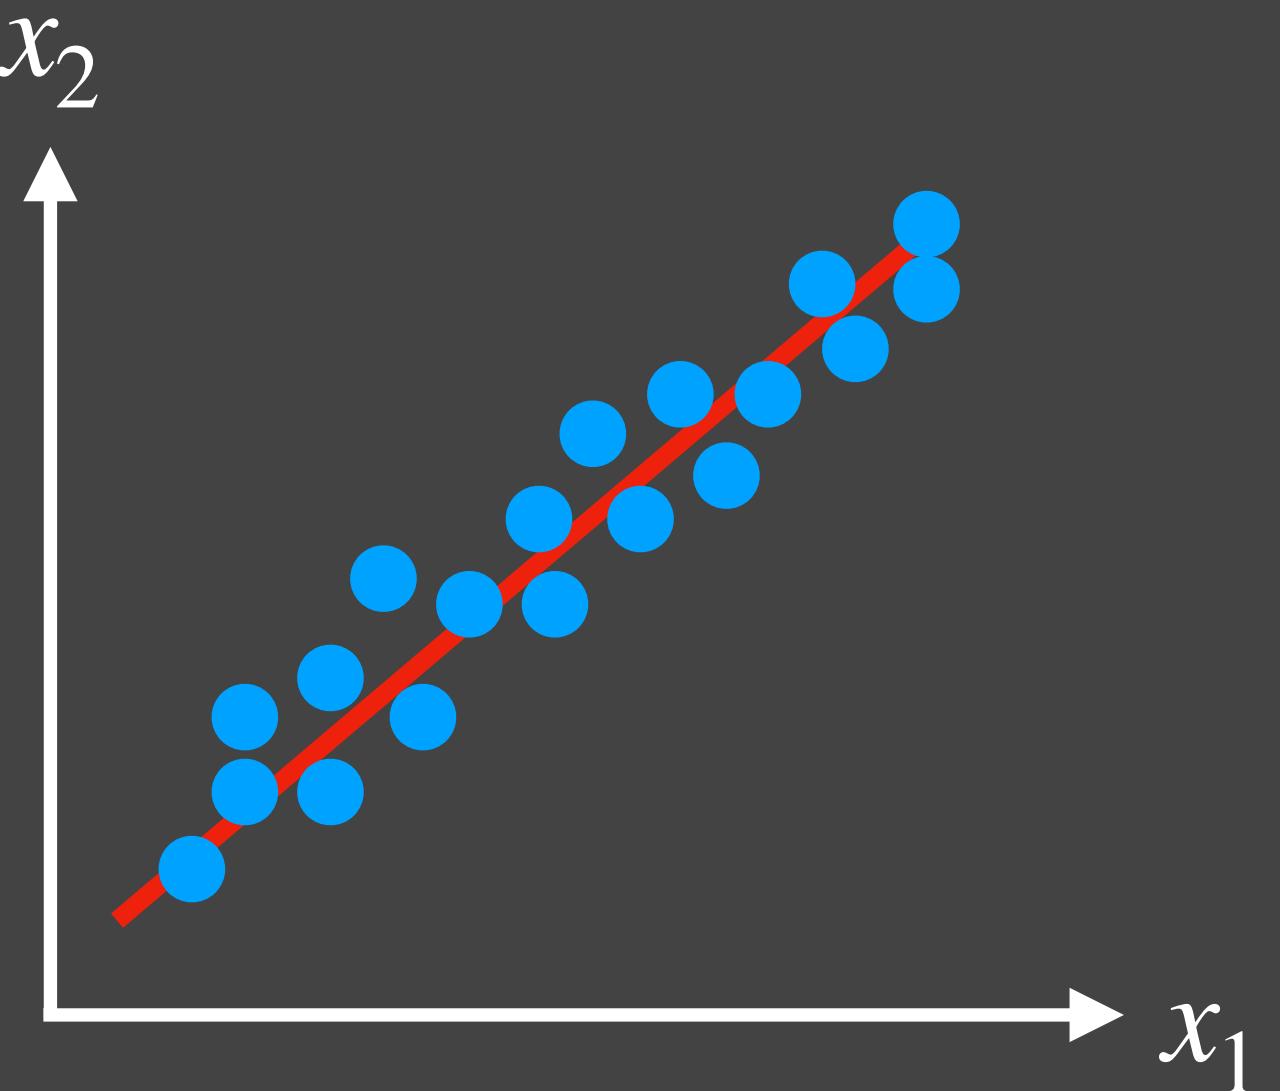
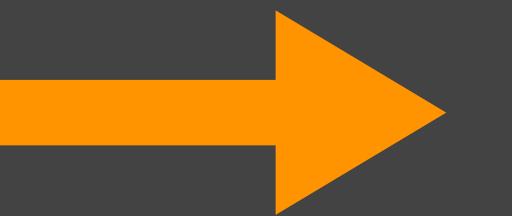
- Updating and optimizing model parameters

# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

$$x \in R^d, \quad w \in R^d, \quad b \in R$$

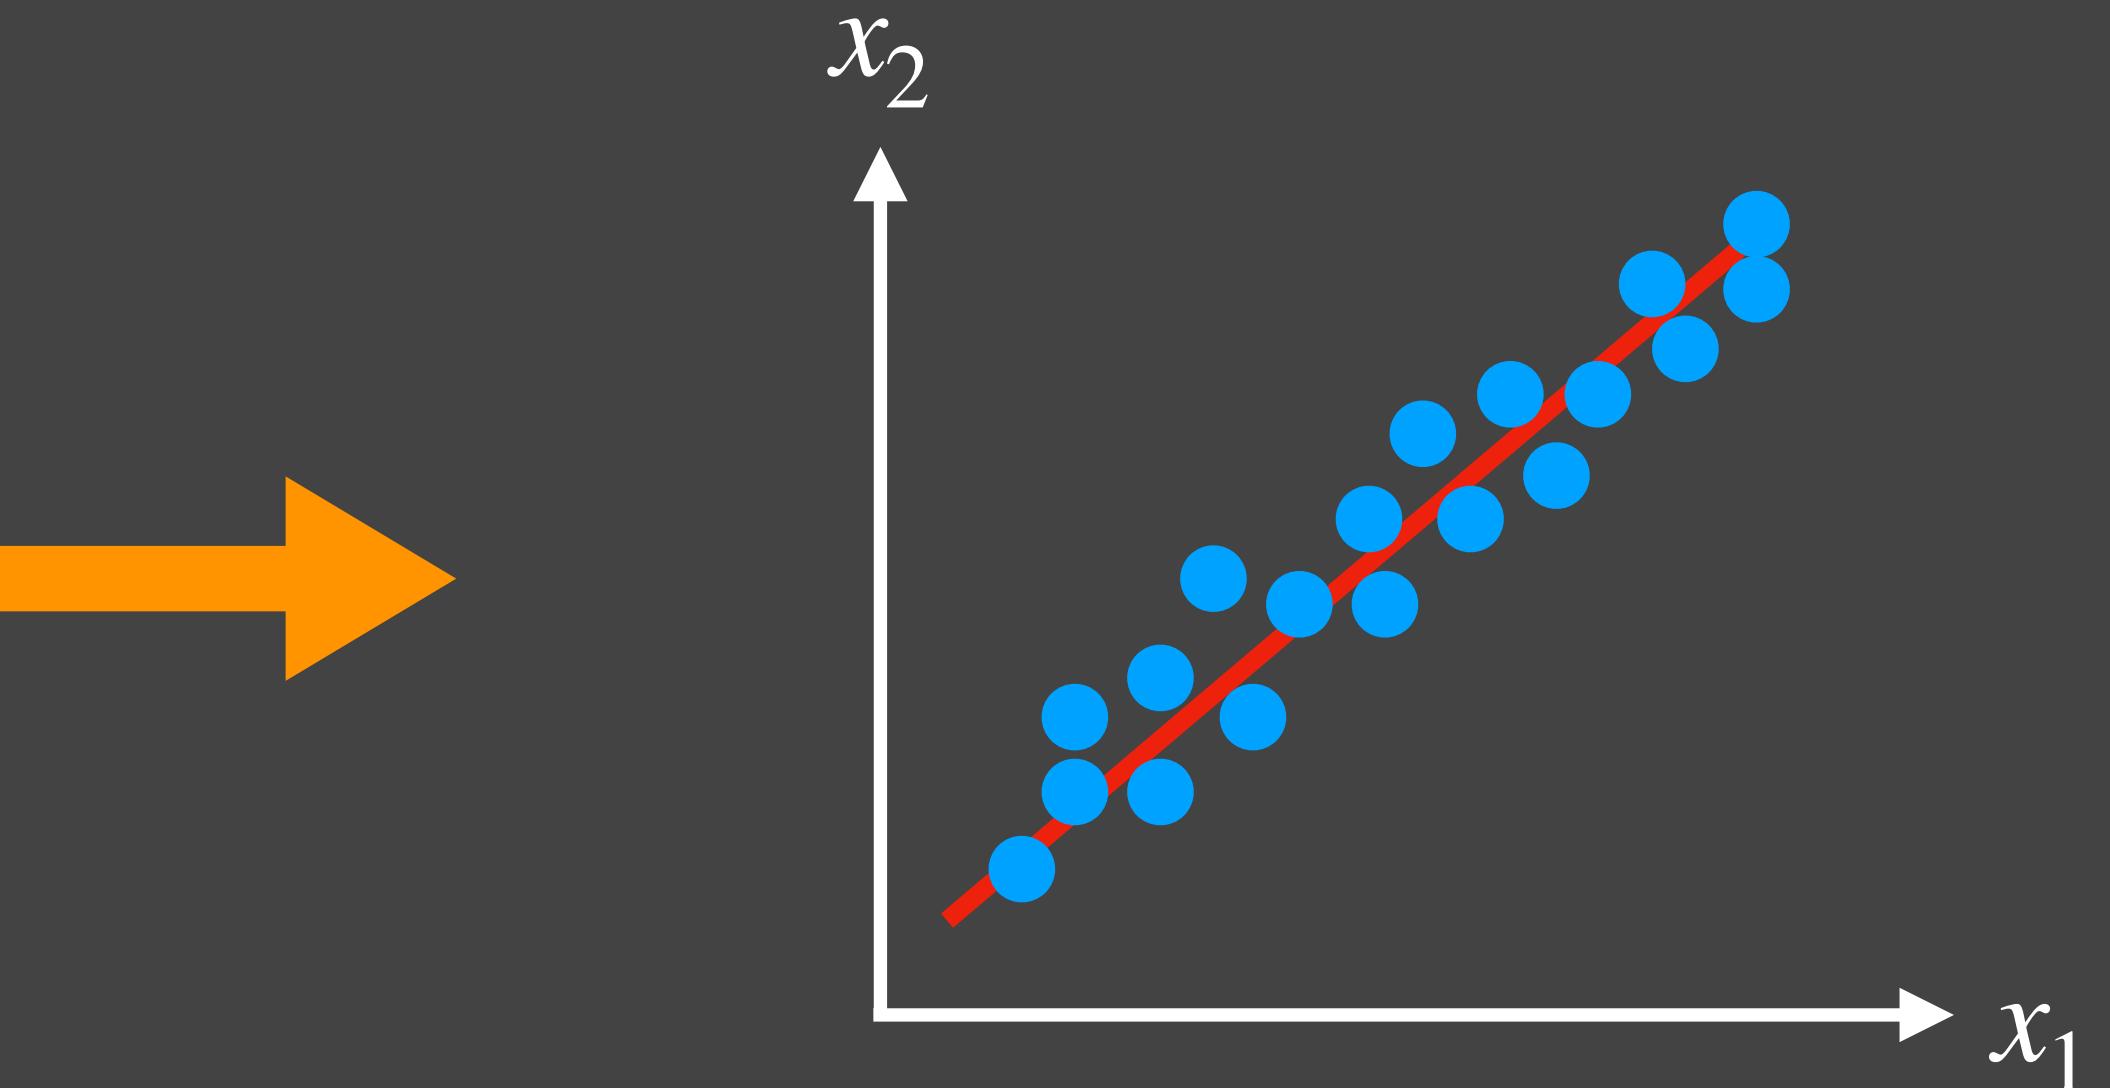


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

$$x \in R^d, \quad w \in R^d, \quad b \in R$$

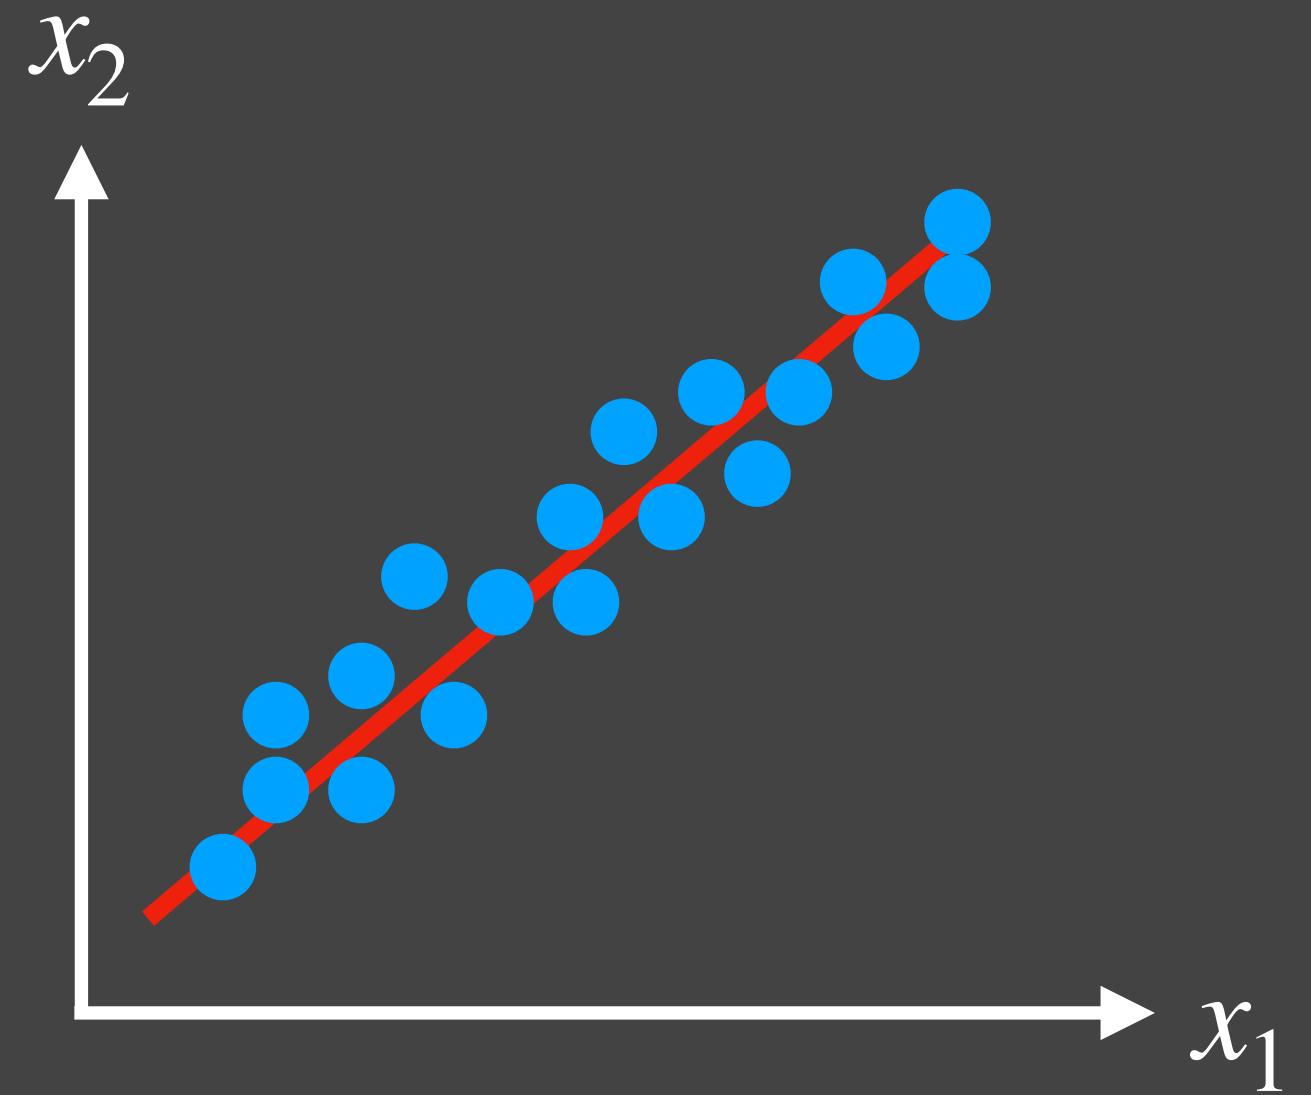
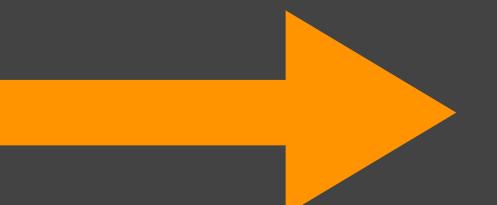


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

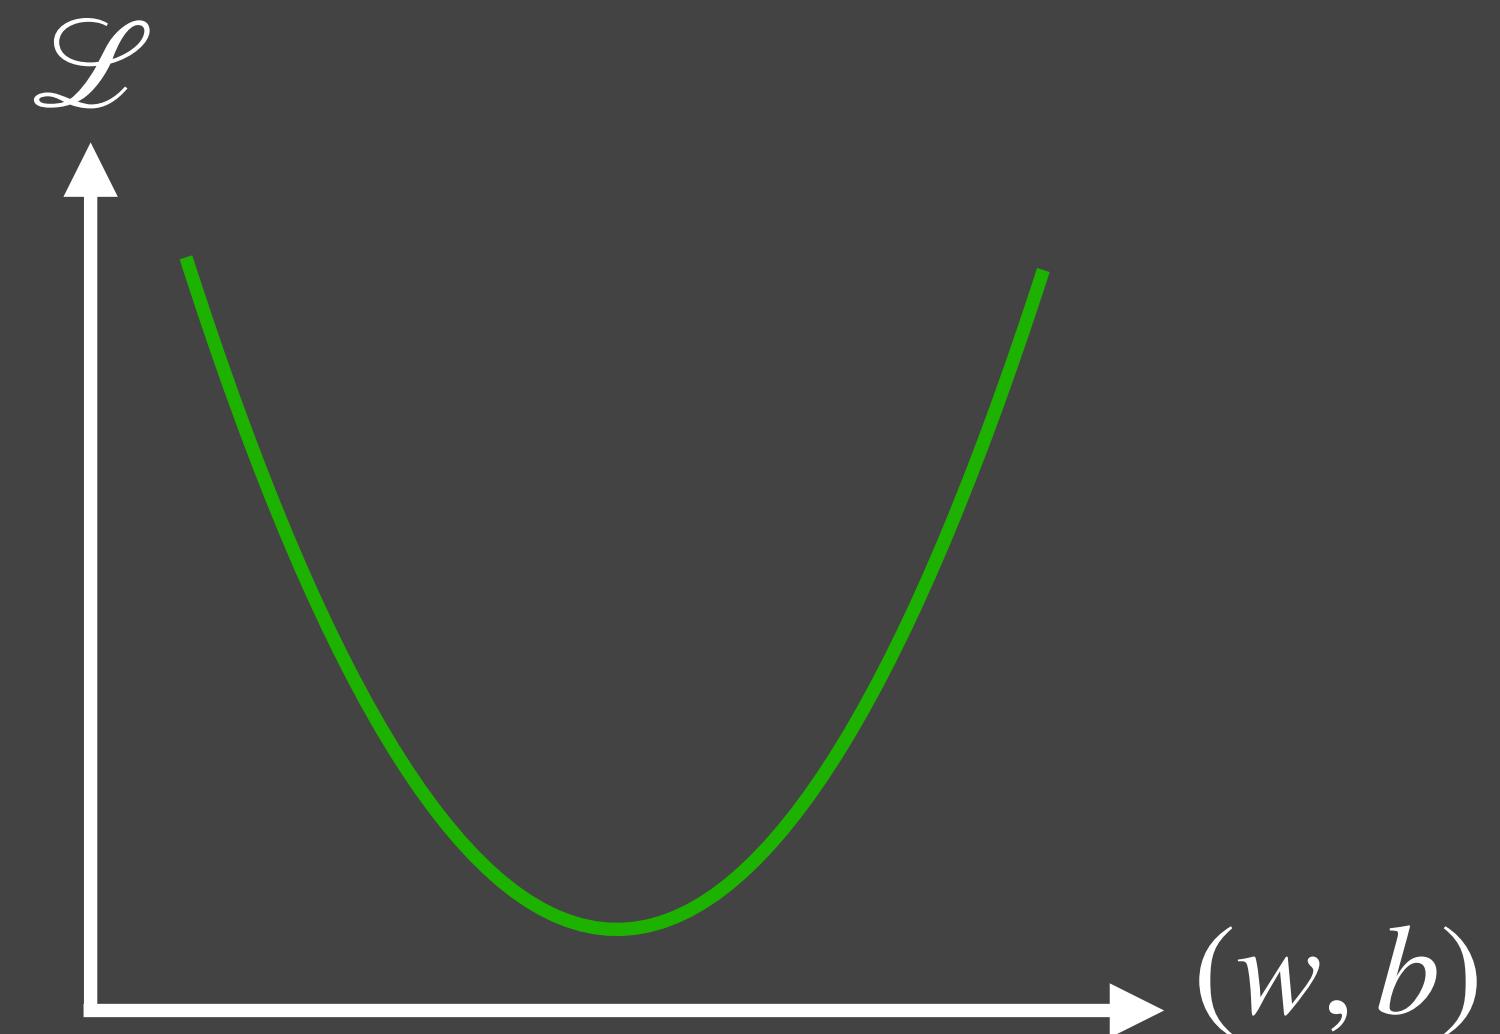
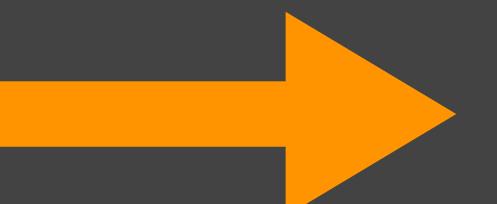
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## Loss Function

$$\mathcal{L}(w, b) = \frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2$$

$$x^i \in R^d, \quad y^i \in R$$

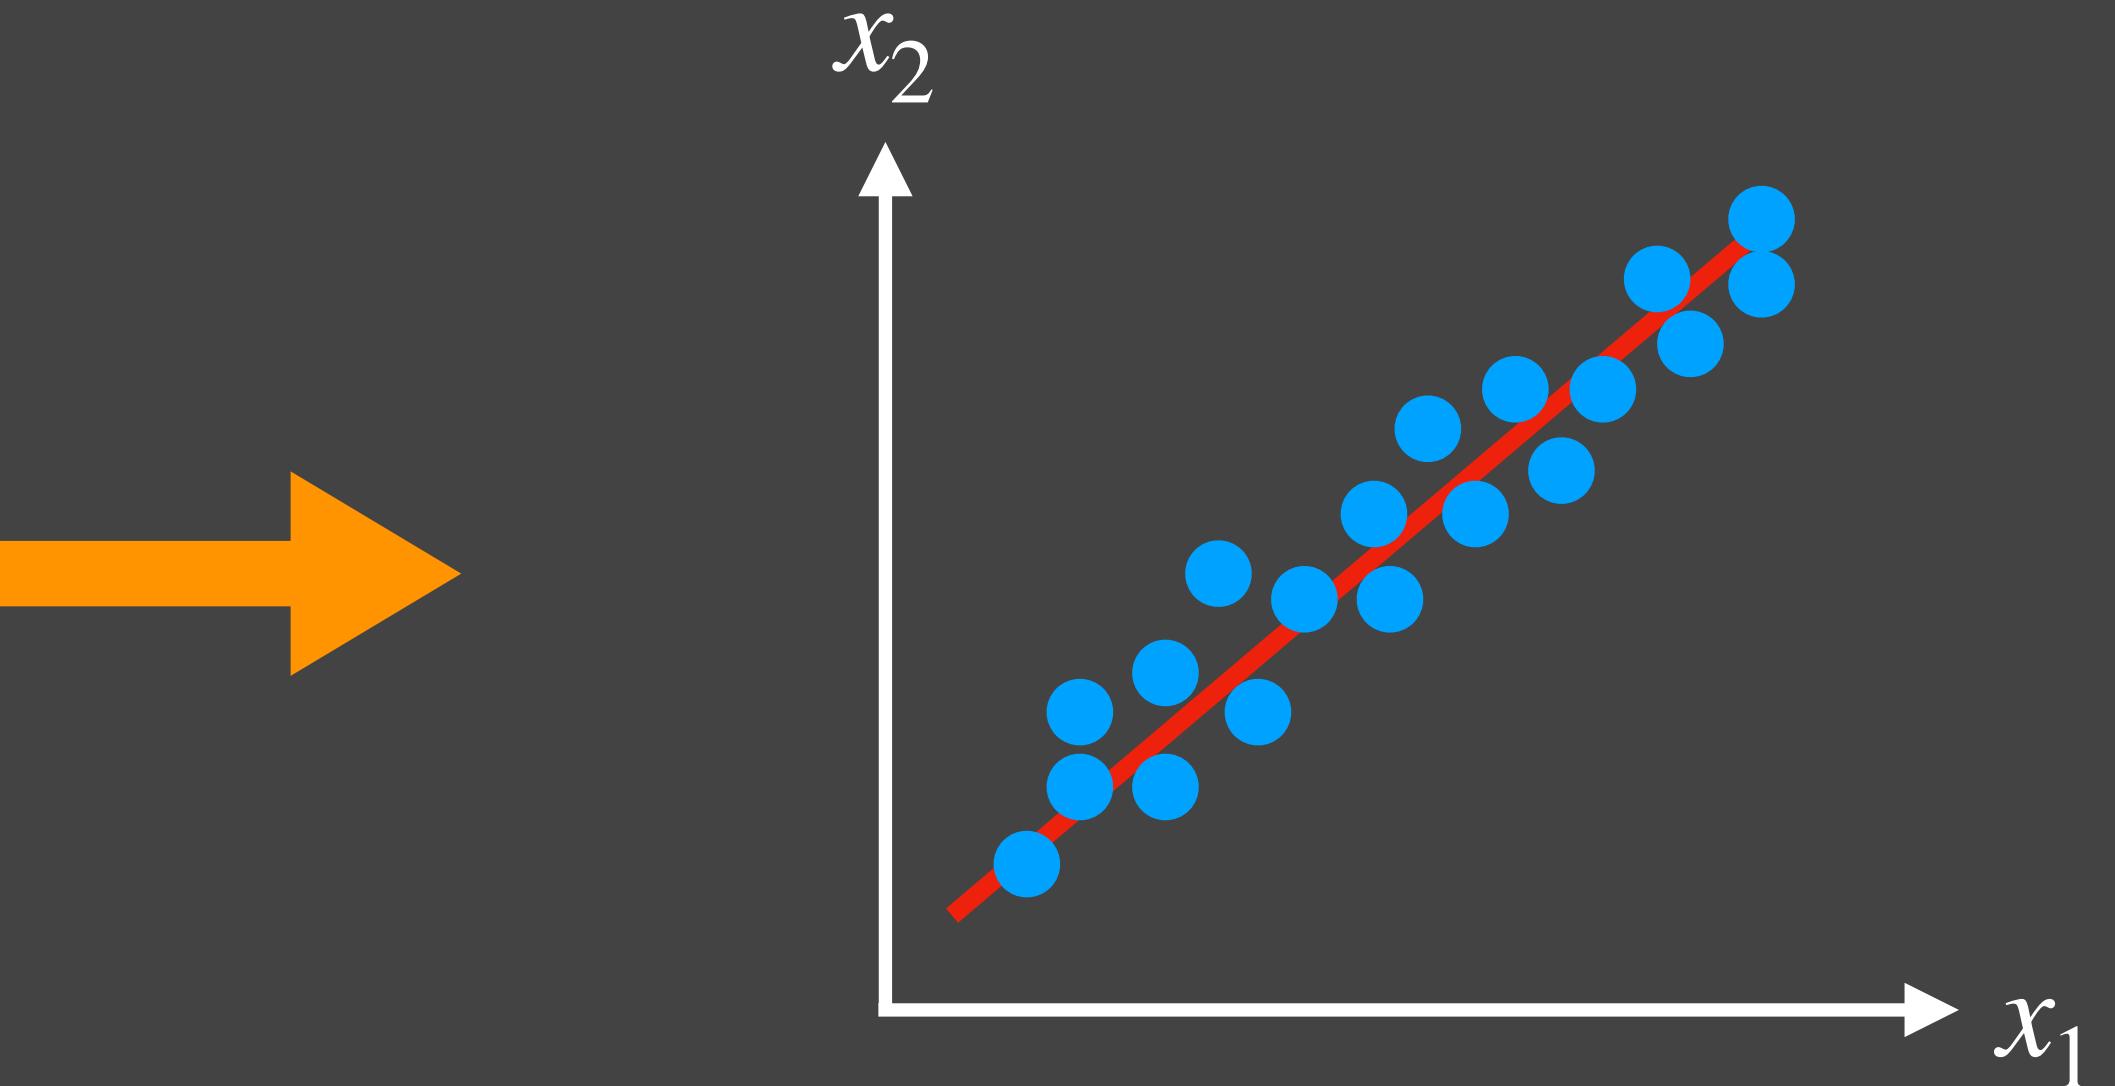


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

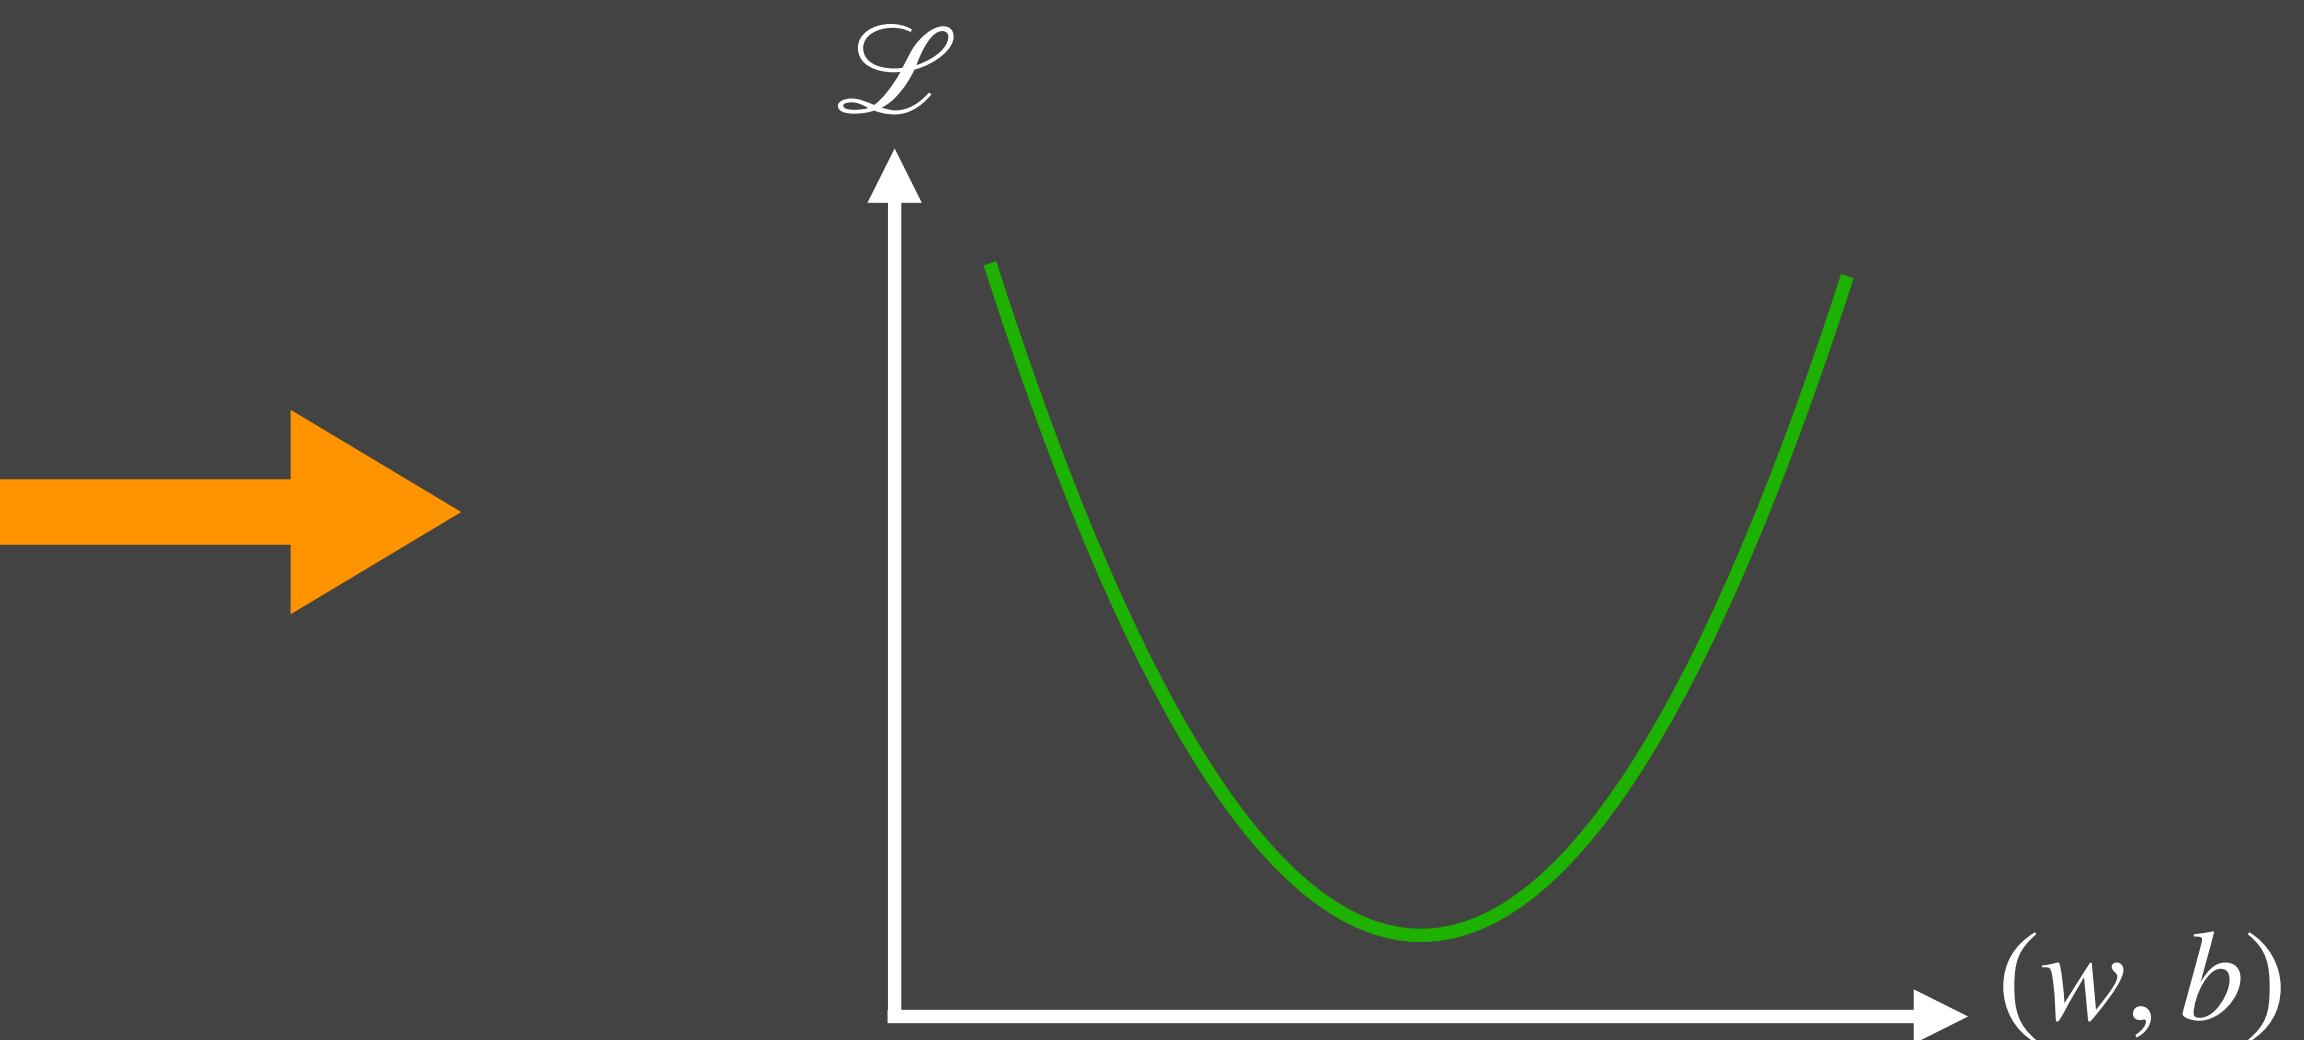
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## Loss Function

$$\mathcal{L}(w, b) = \frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2$$

$$x^i \in R^d, \quad y^i \in R$$

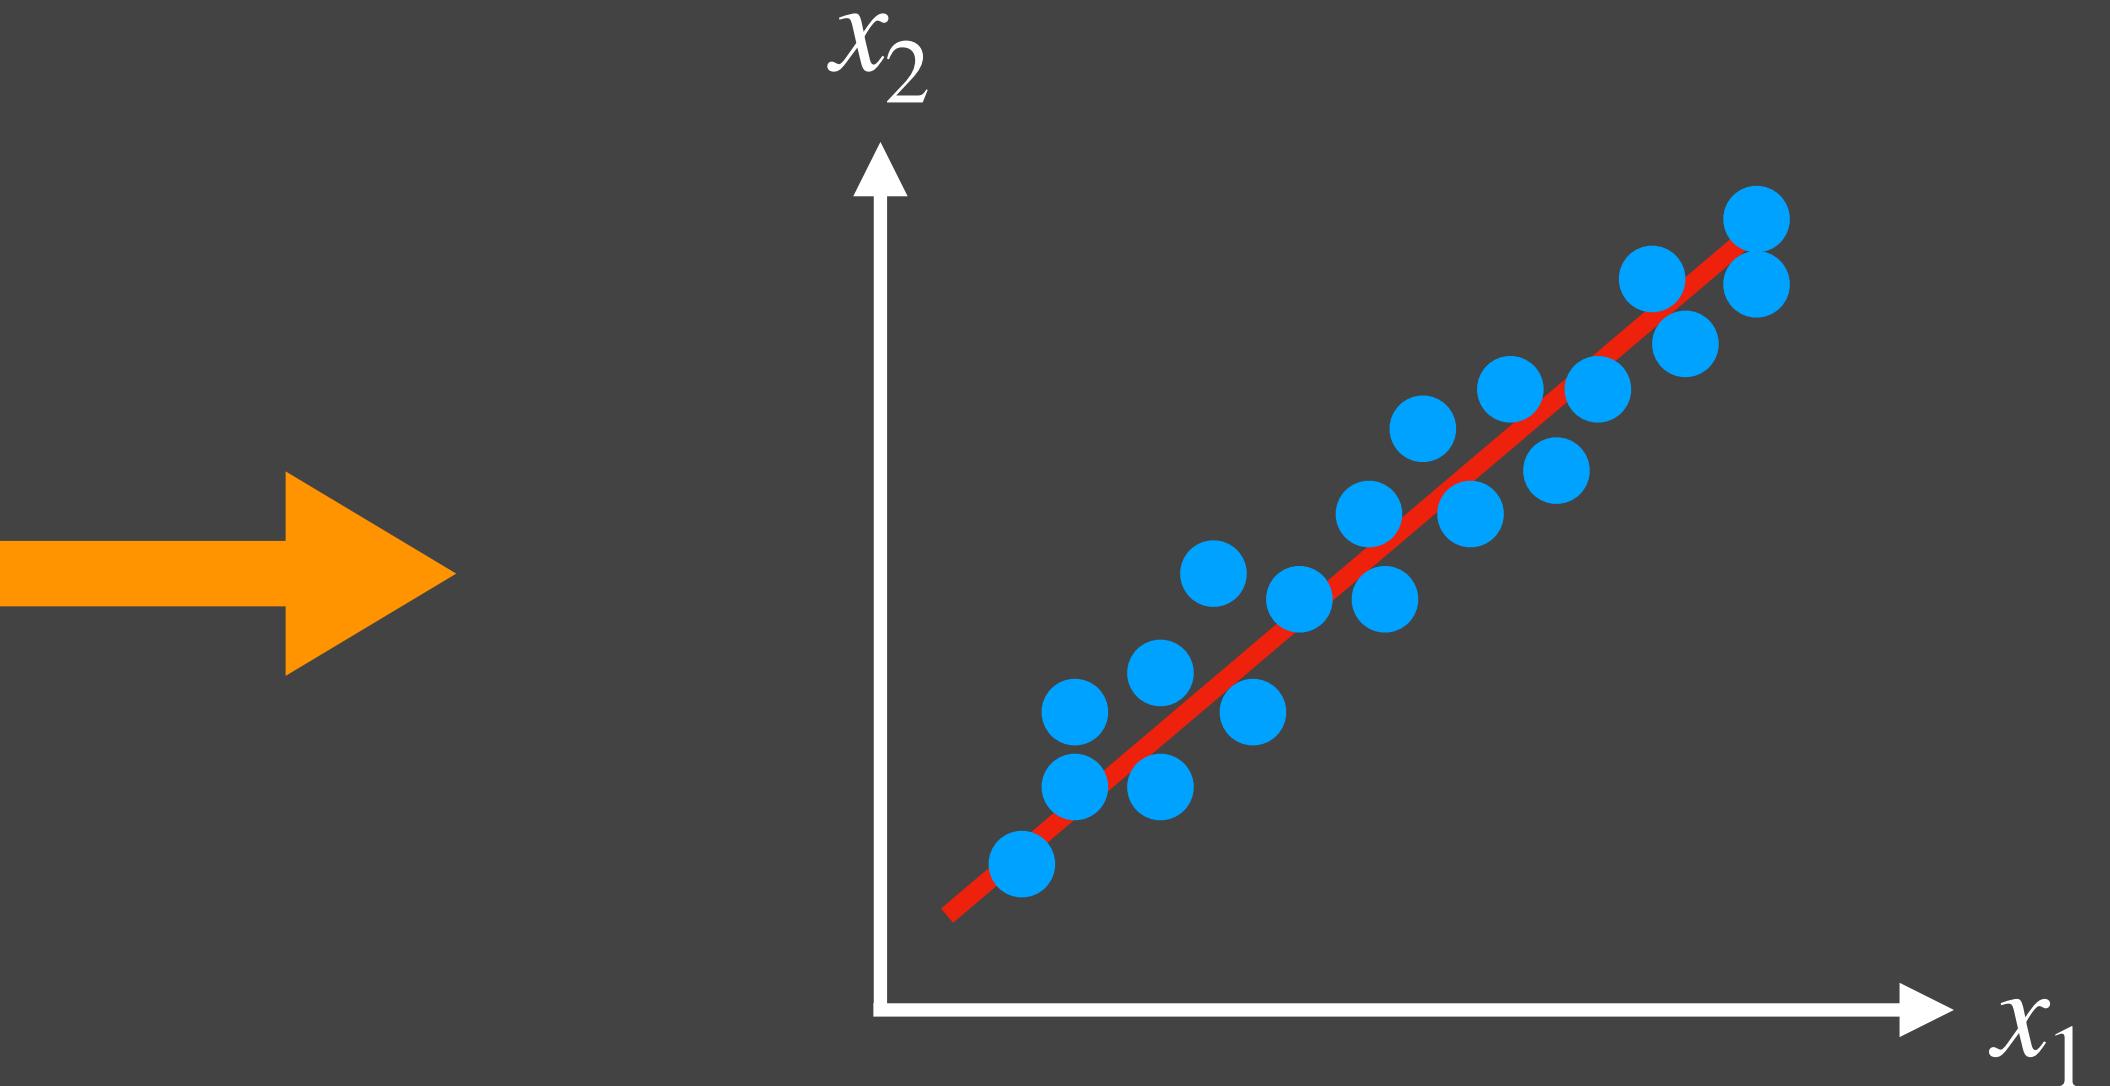


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

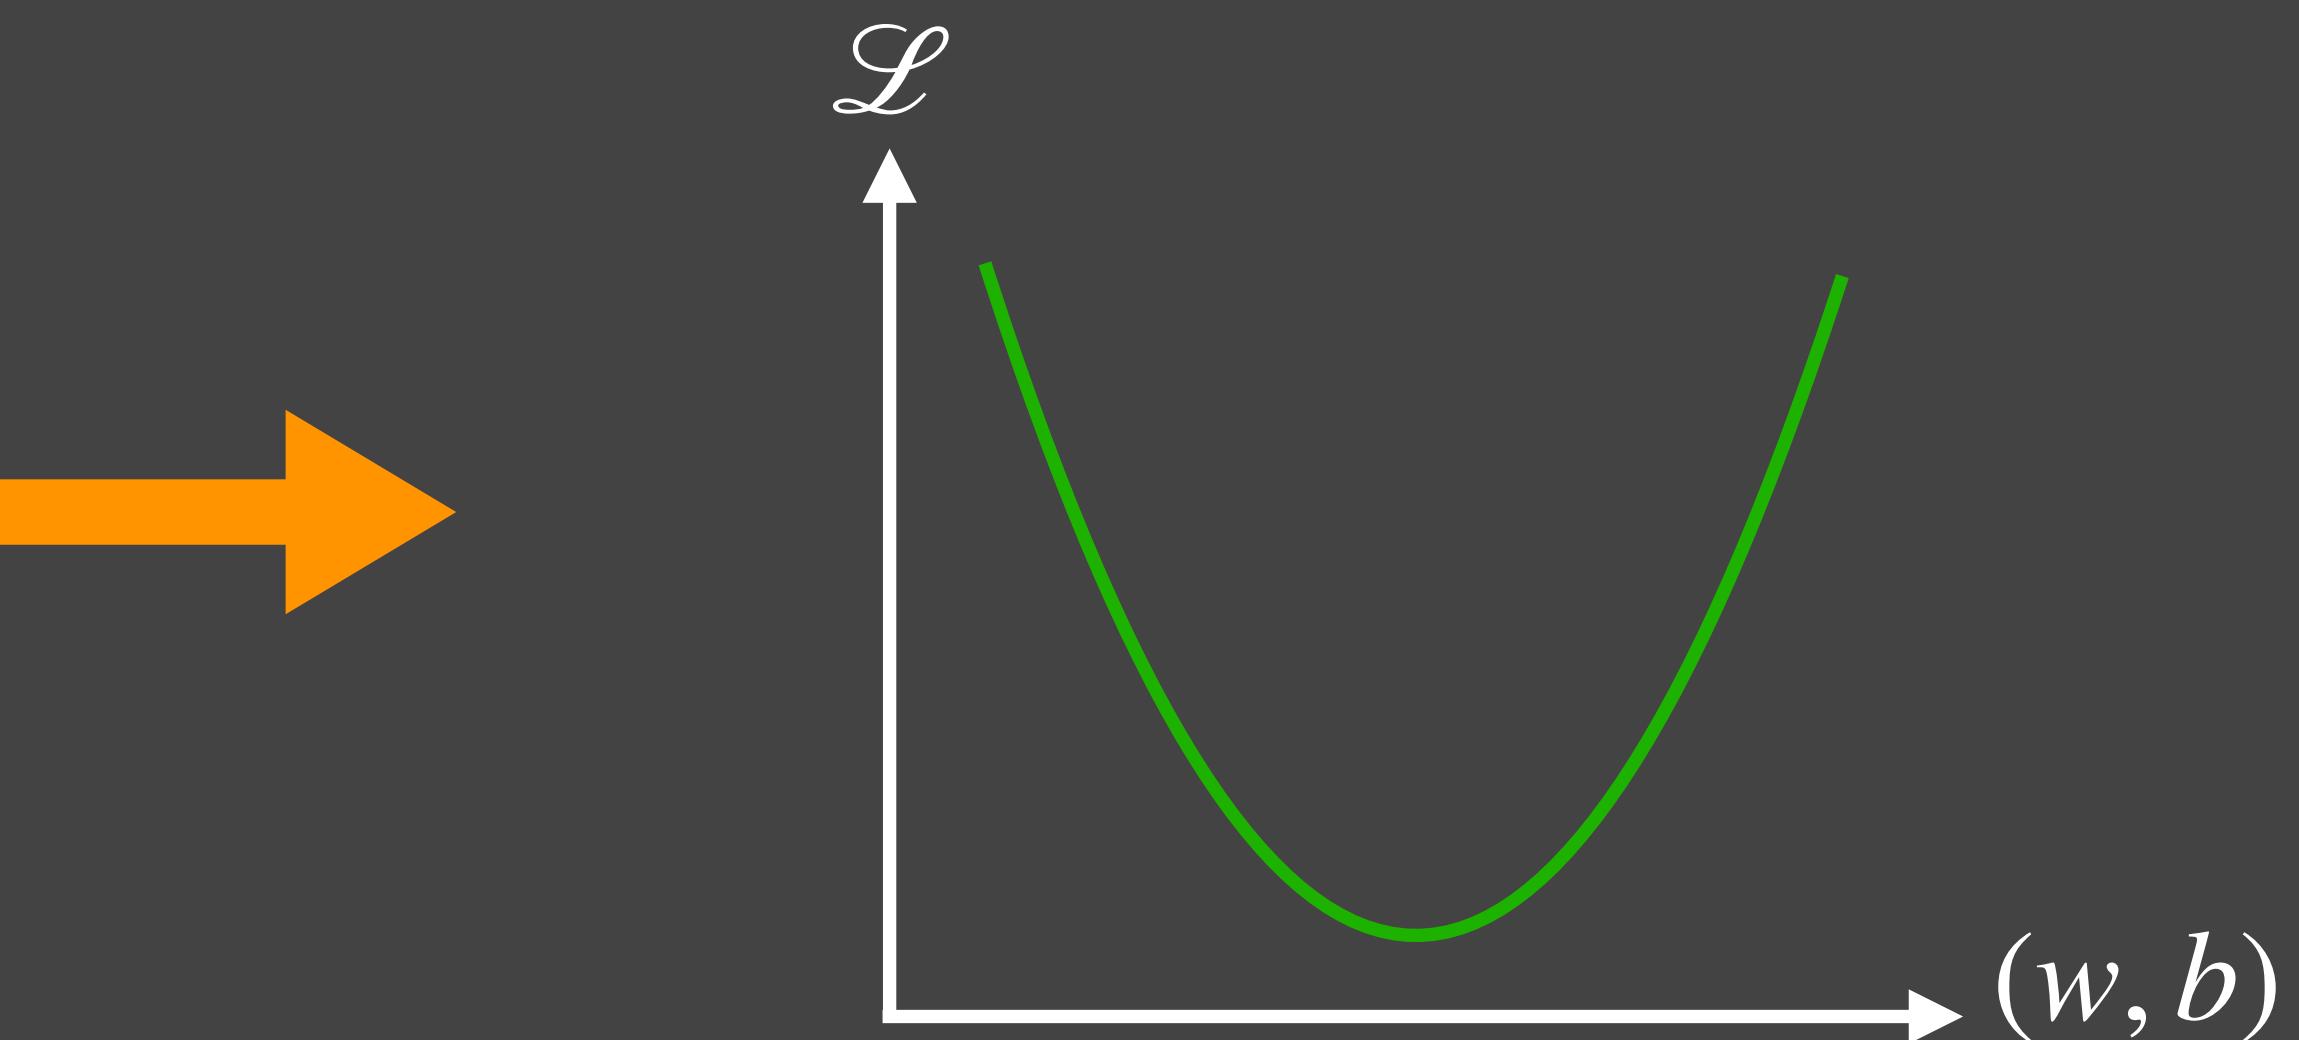
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## Loss Function

$$\mathcal{L}(w, b) = \frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2$$

$$x^i \in R^d, \quad y^i \in R$$

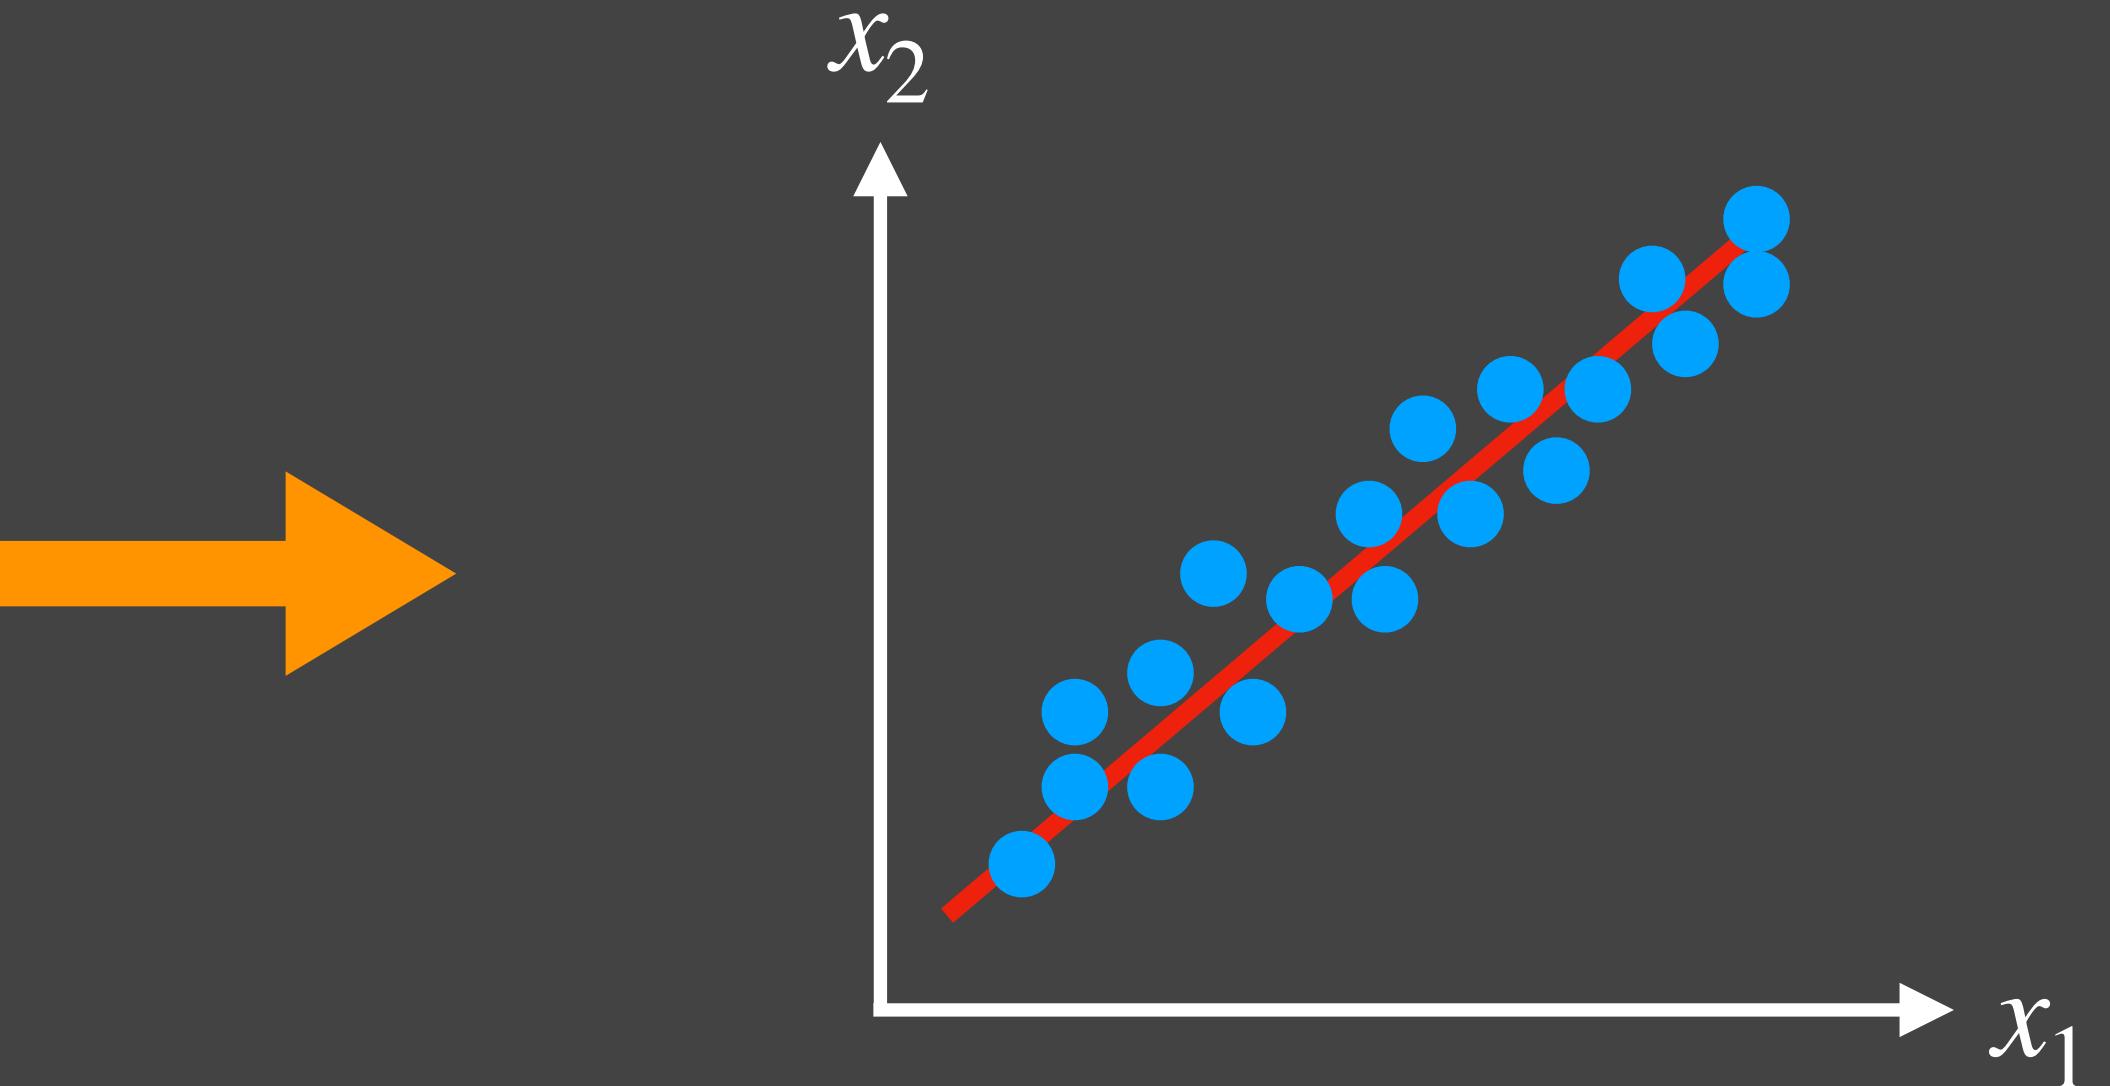


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

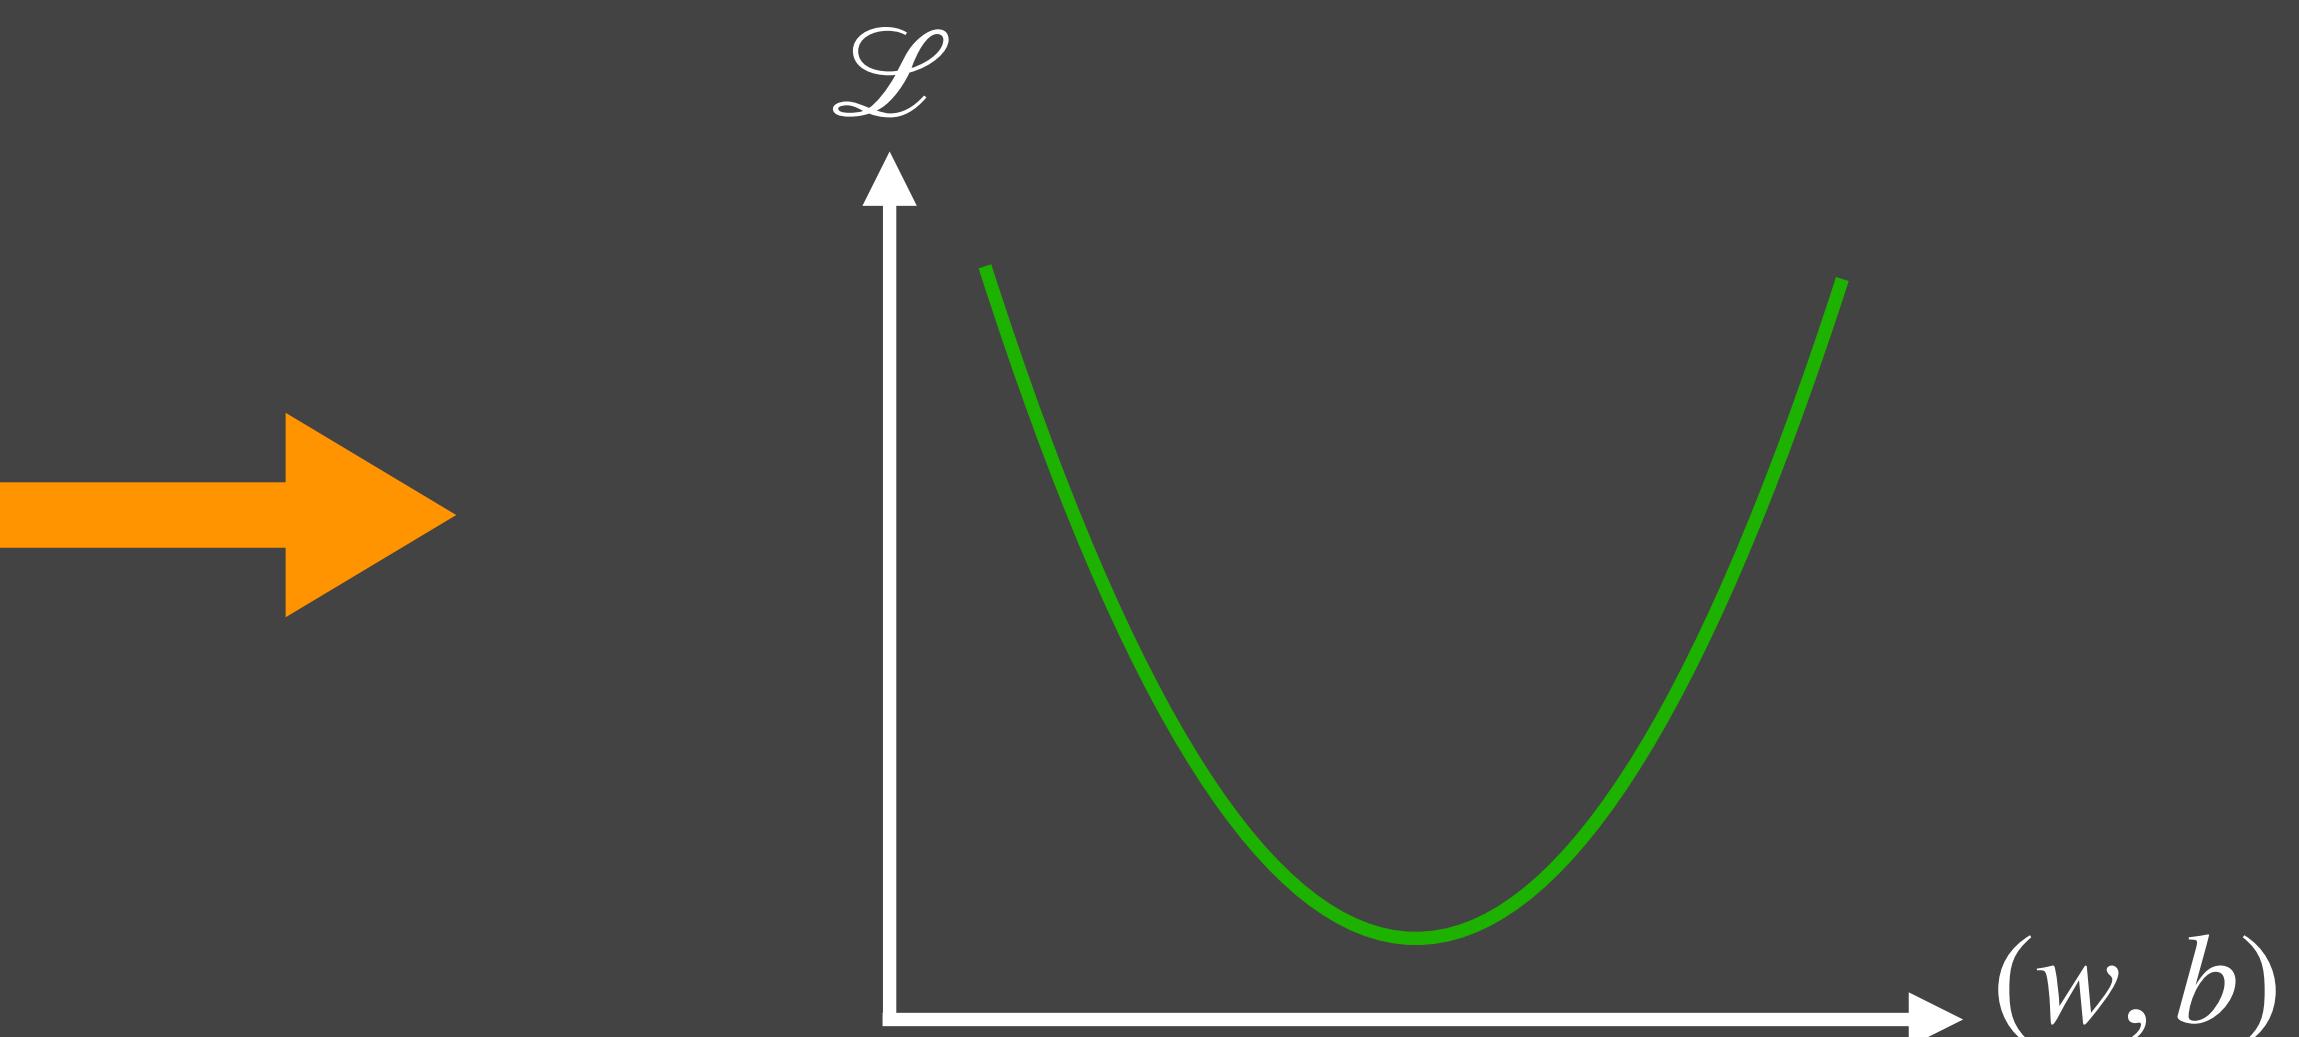
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## Loss Function

$$\mathcal{L}(w, b) = \boxed{\frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2}$$

$$x^i \in R^d, \quad y^i \in R$$

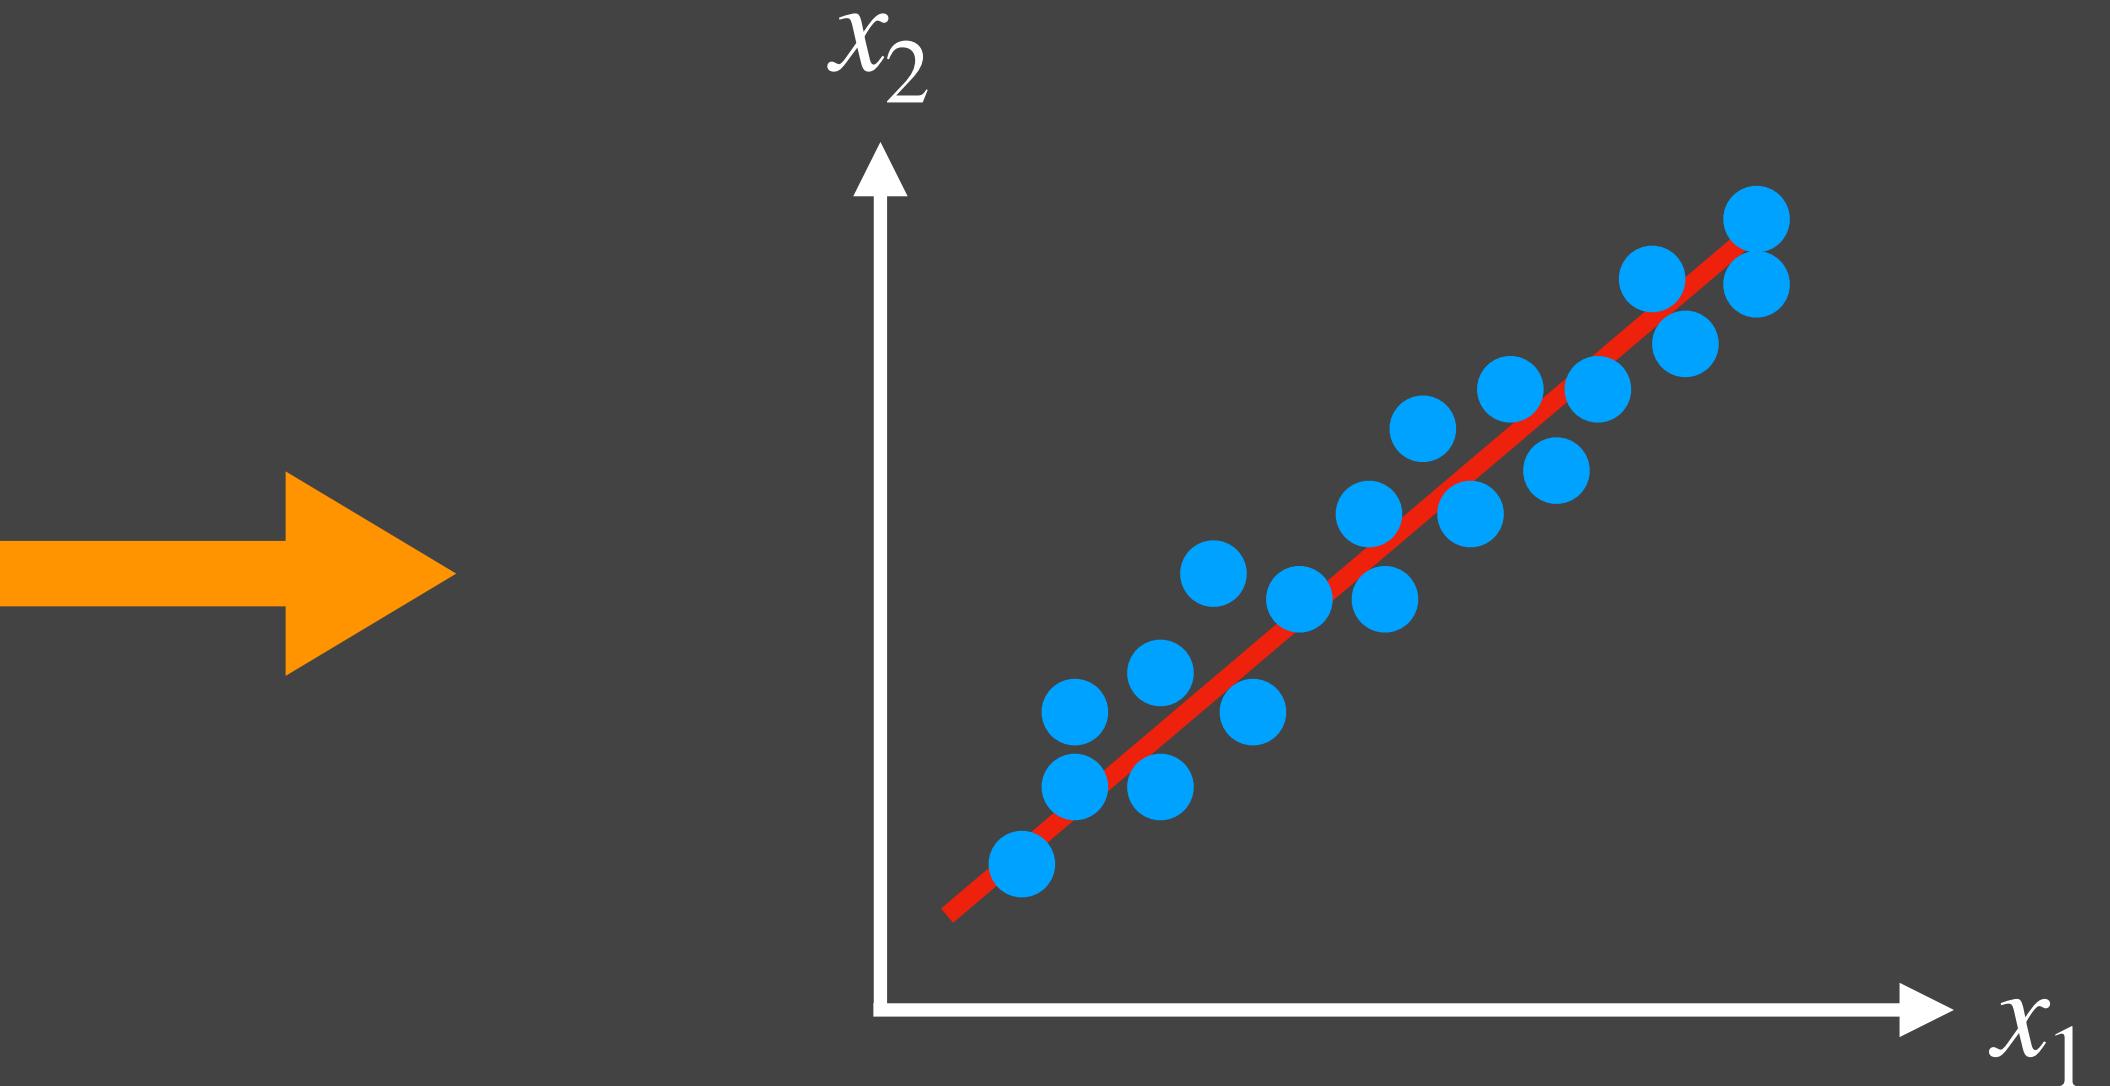


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

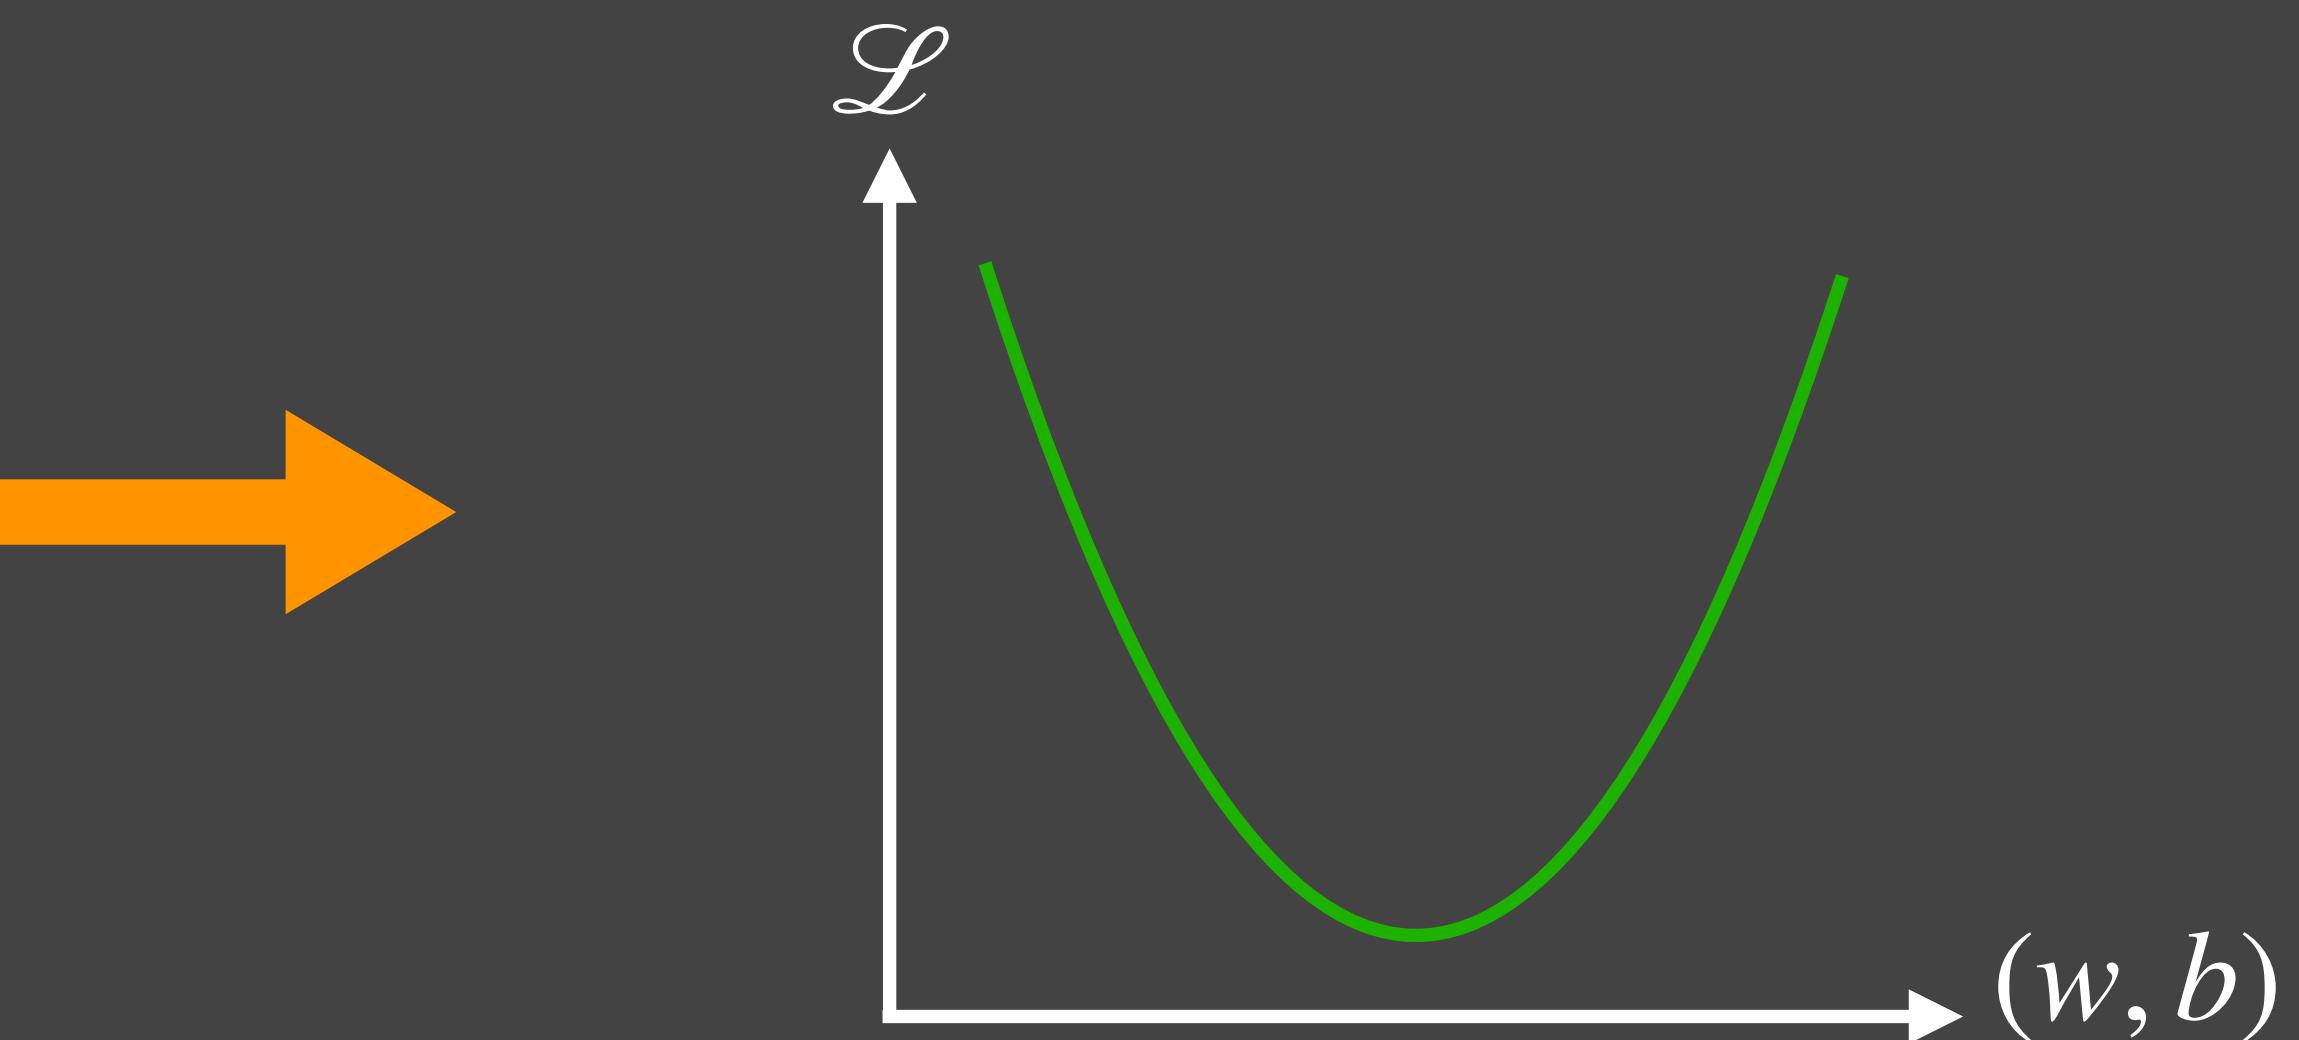
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## Loss Function

$$\mathcal{L}(w, b) = \frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2$$

$$x^i \in R^d, \quad y^i \in R$$

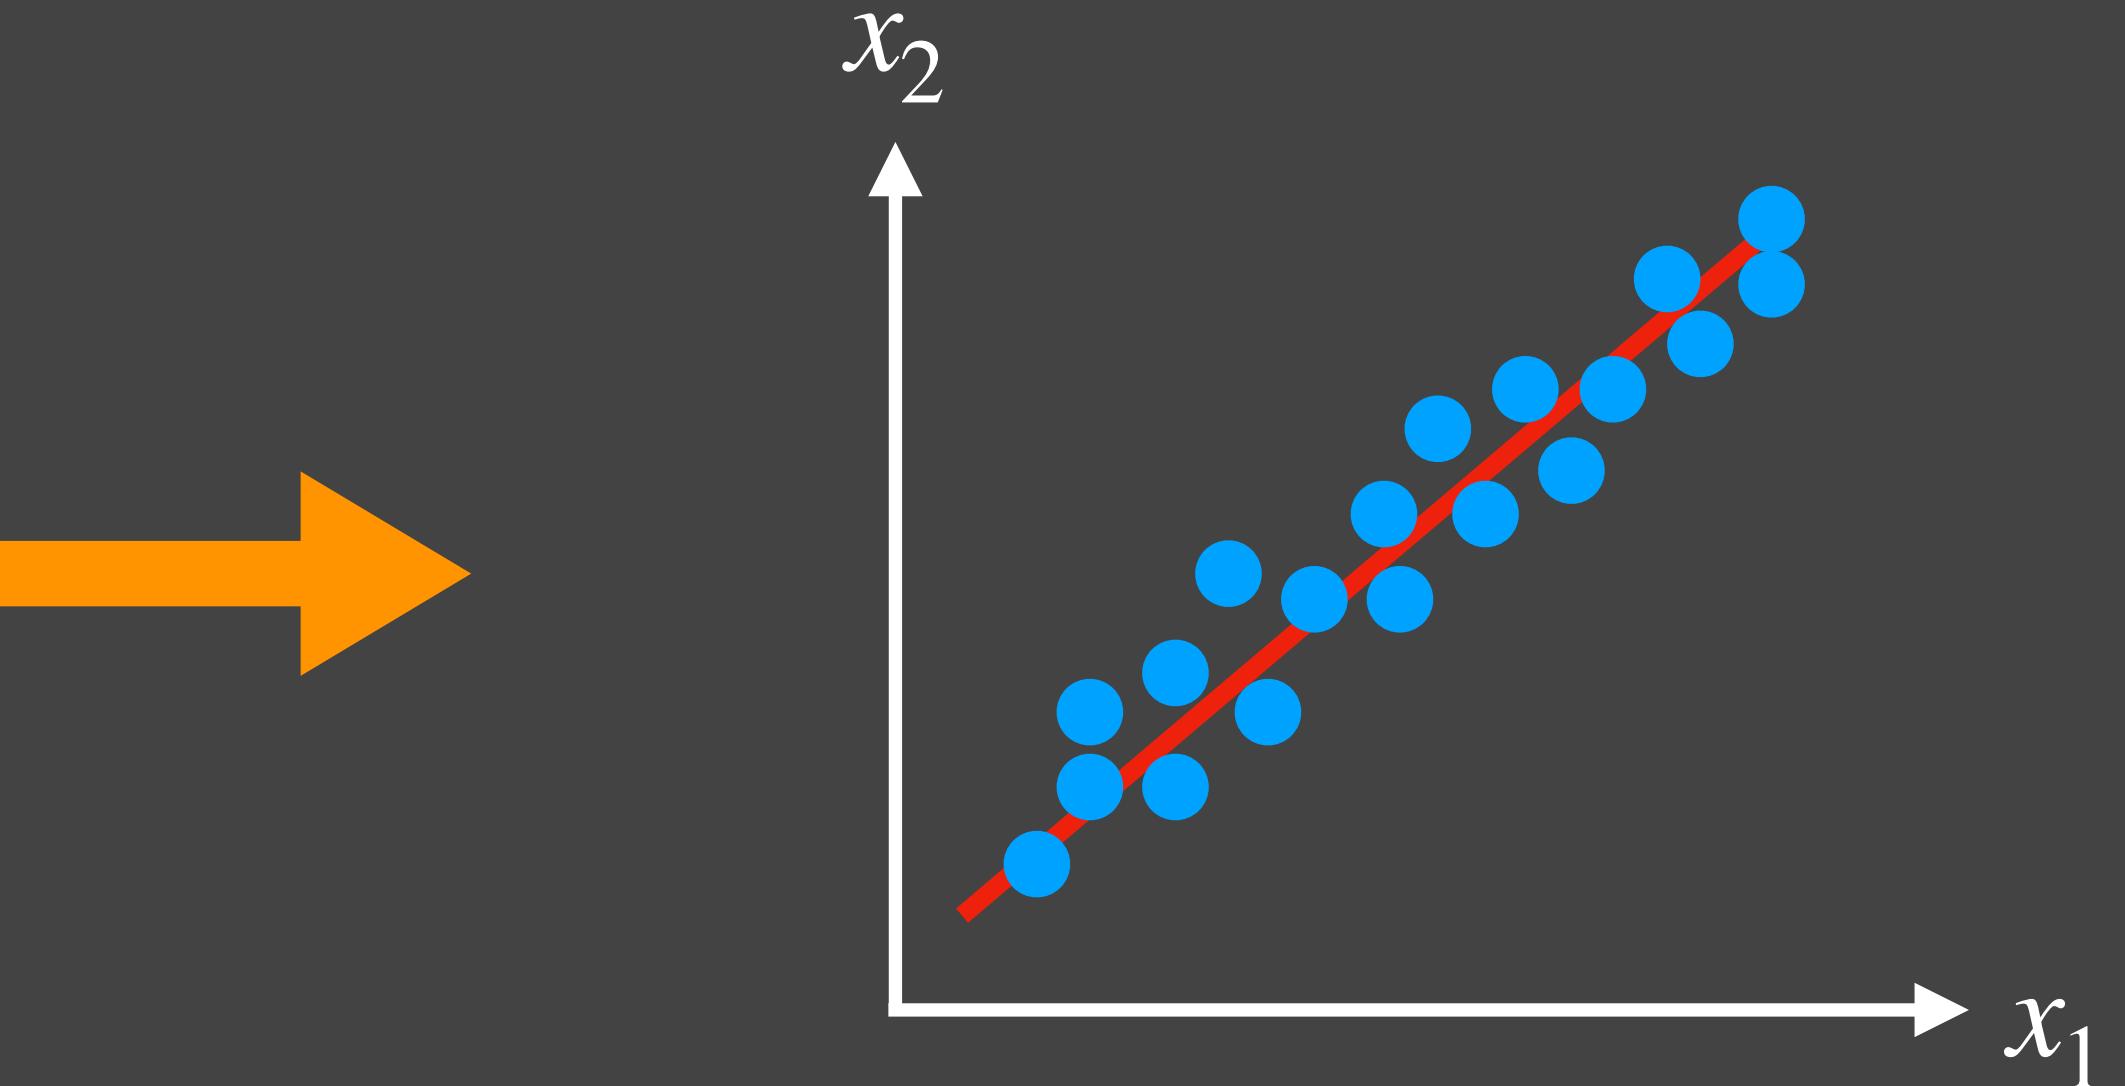


# Optimization Theory - Model Notations

## Model Function

$$f_{w,b}(x) = x^T \cdot w + b$$

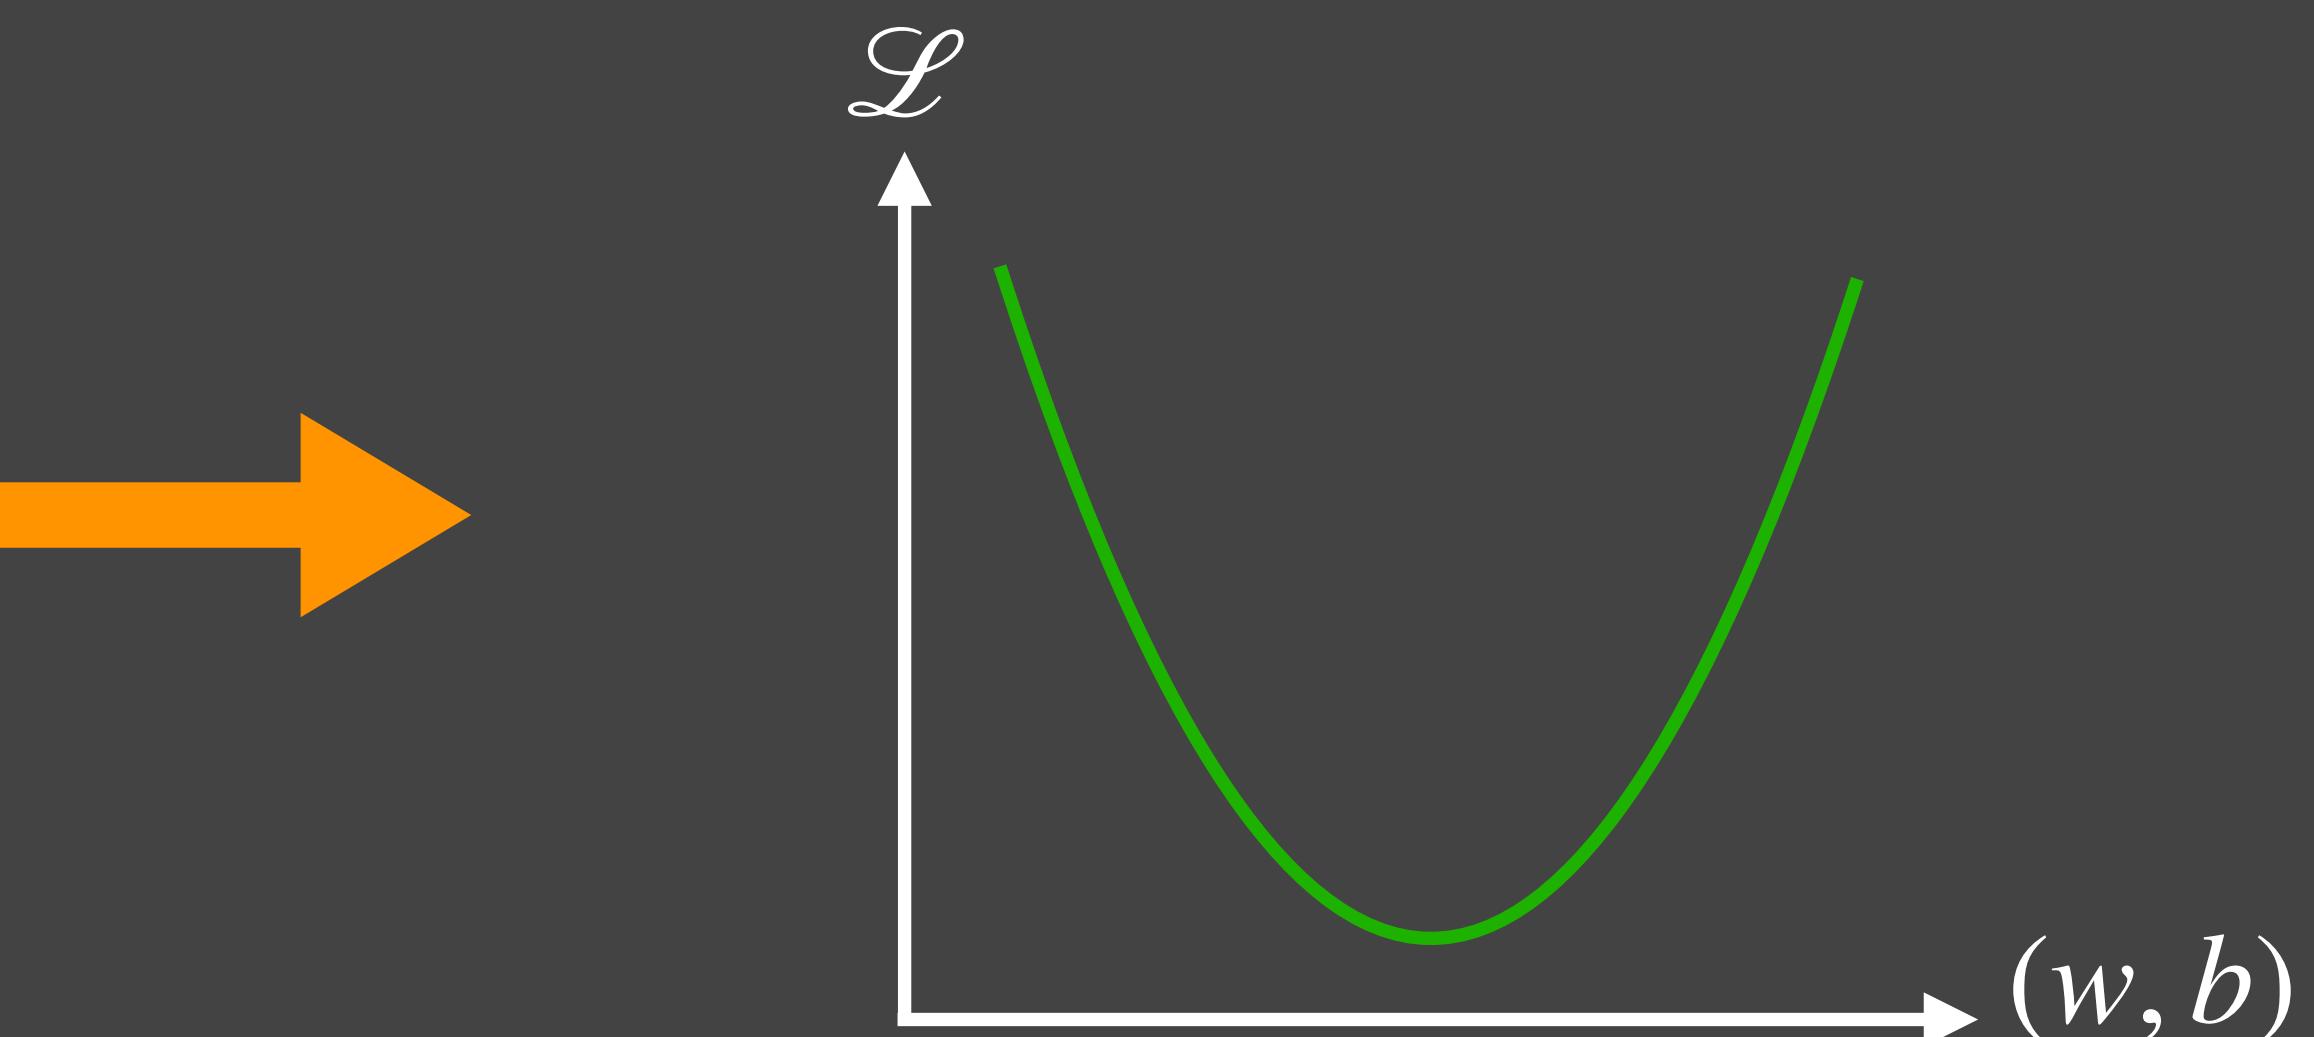
$$x \in R^d, \quad w \in R^d, \quad b \in R$$



## MSE Loss

$$\mathcal{L}(w, b) = \frac{1}{N} \sum_{i=1}^N (f_{w,b}(x^i) - y^i)^2$$

$$x^i \in R^d, \quad y^i \in R$$



# Optimization Theory - Taxonomy of Machine Learning

## Dataset

- Resource for updating and optimizing model parameters
- Larger dataset, better model

## Model and Loss Functions

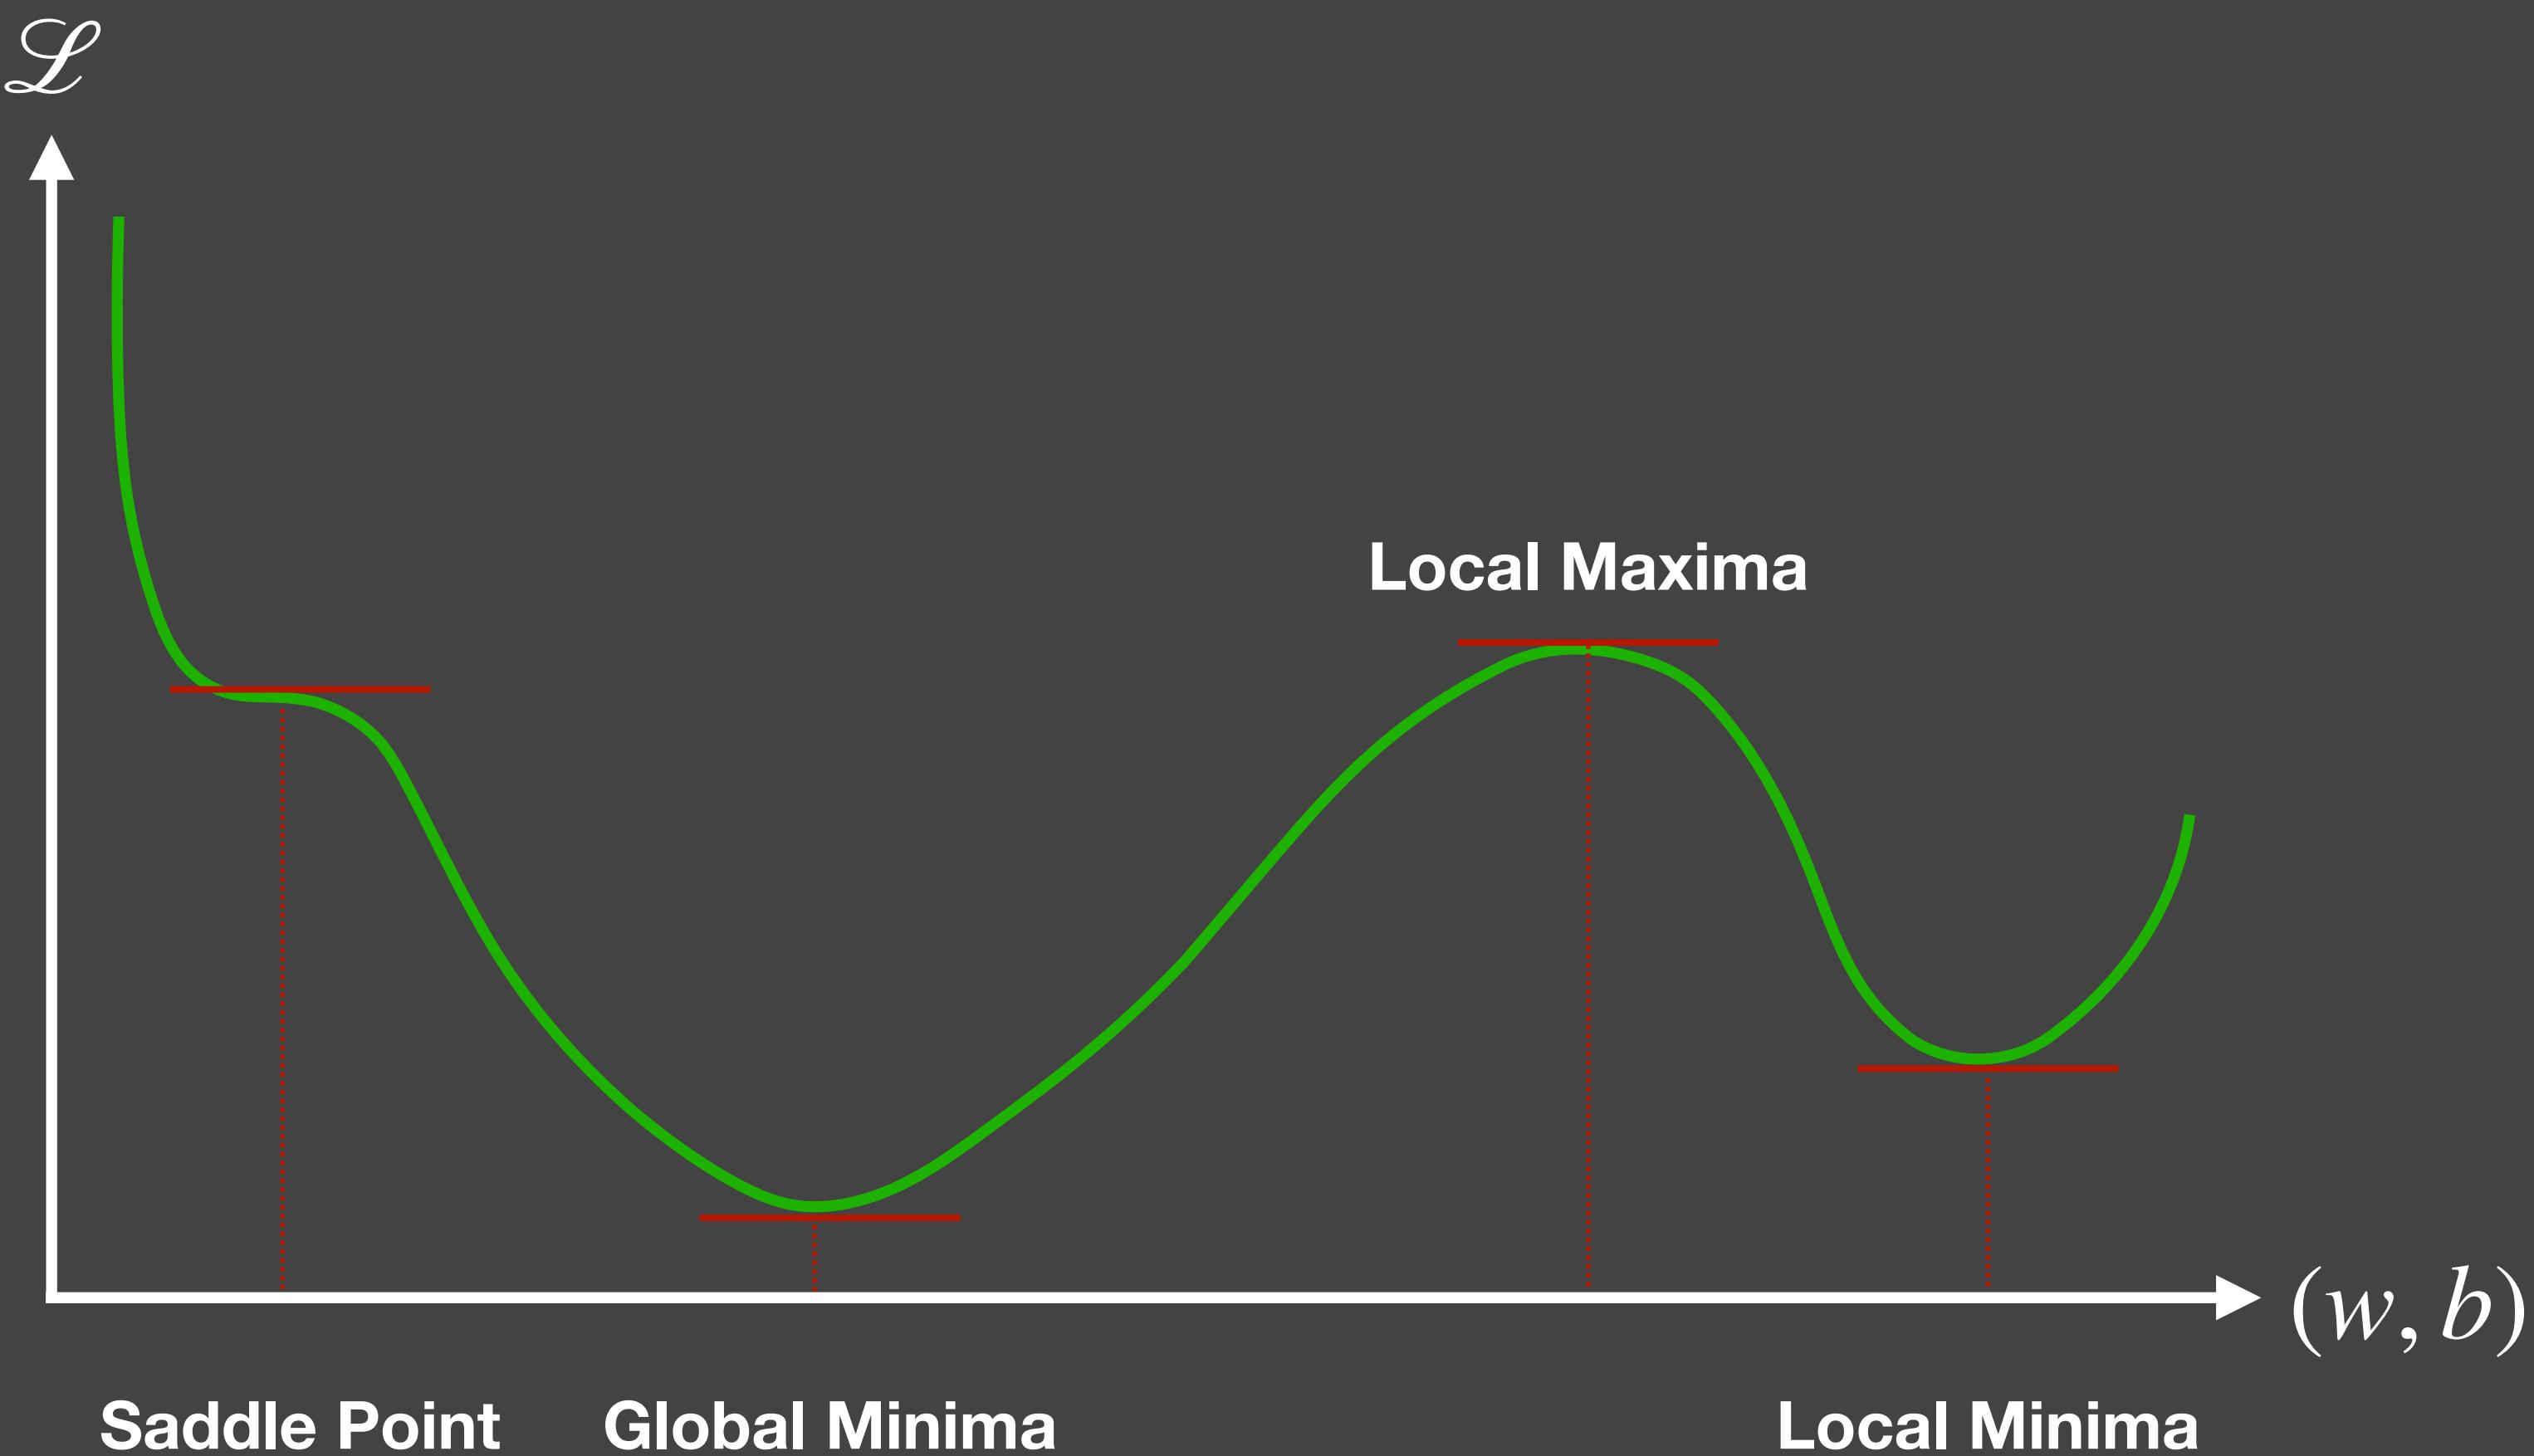
- Model: Mathematical functions to learn relationship between features and labels
- Loss Function: Mathematical functions to compare predictions and labels

## Optimization Algorithm

- Updating and optimizing model parameters

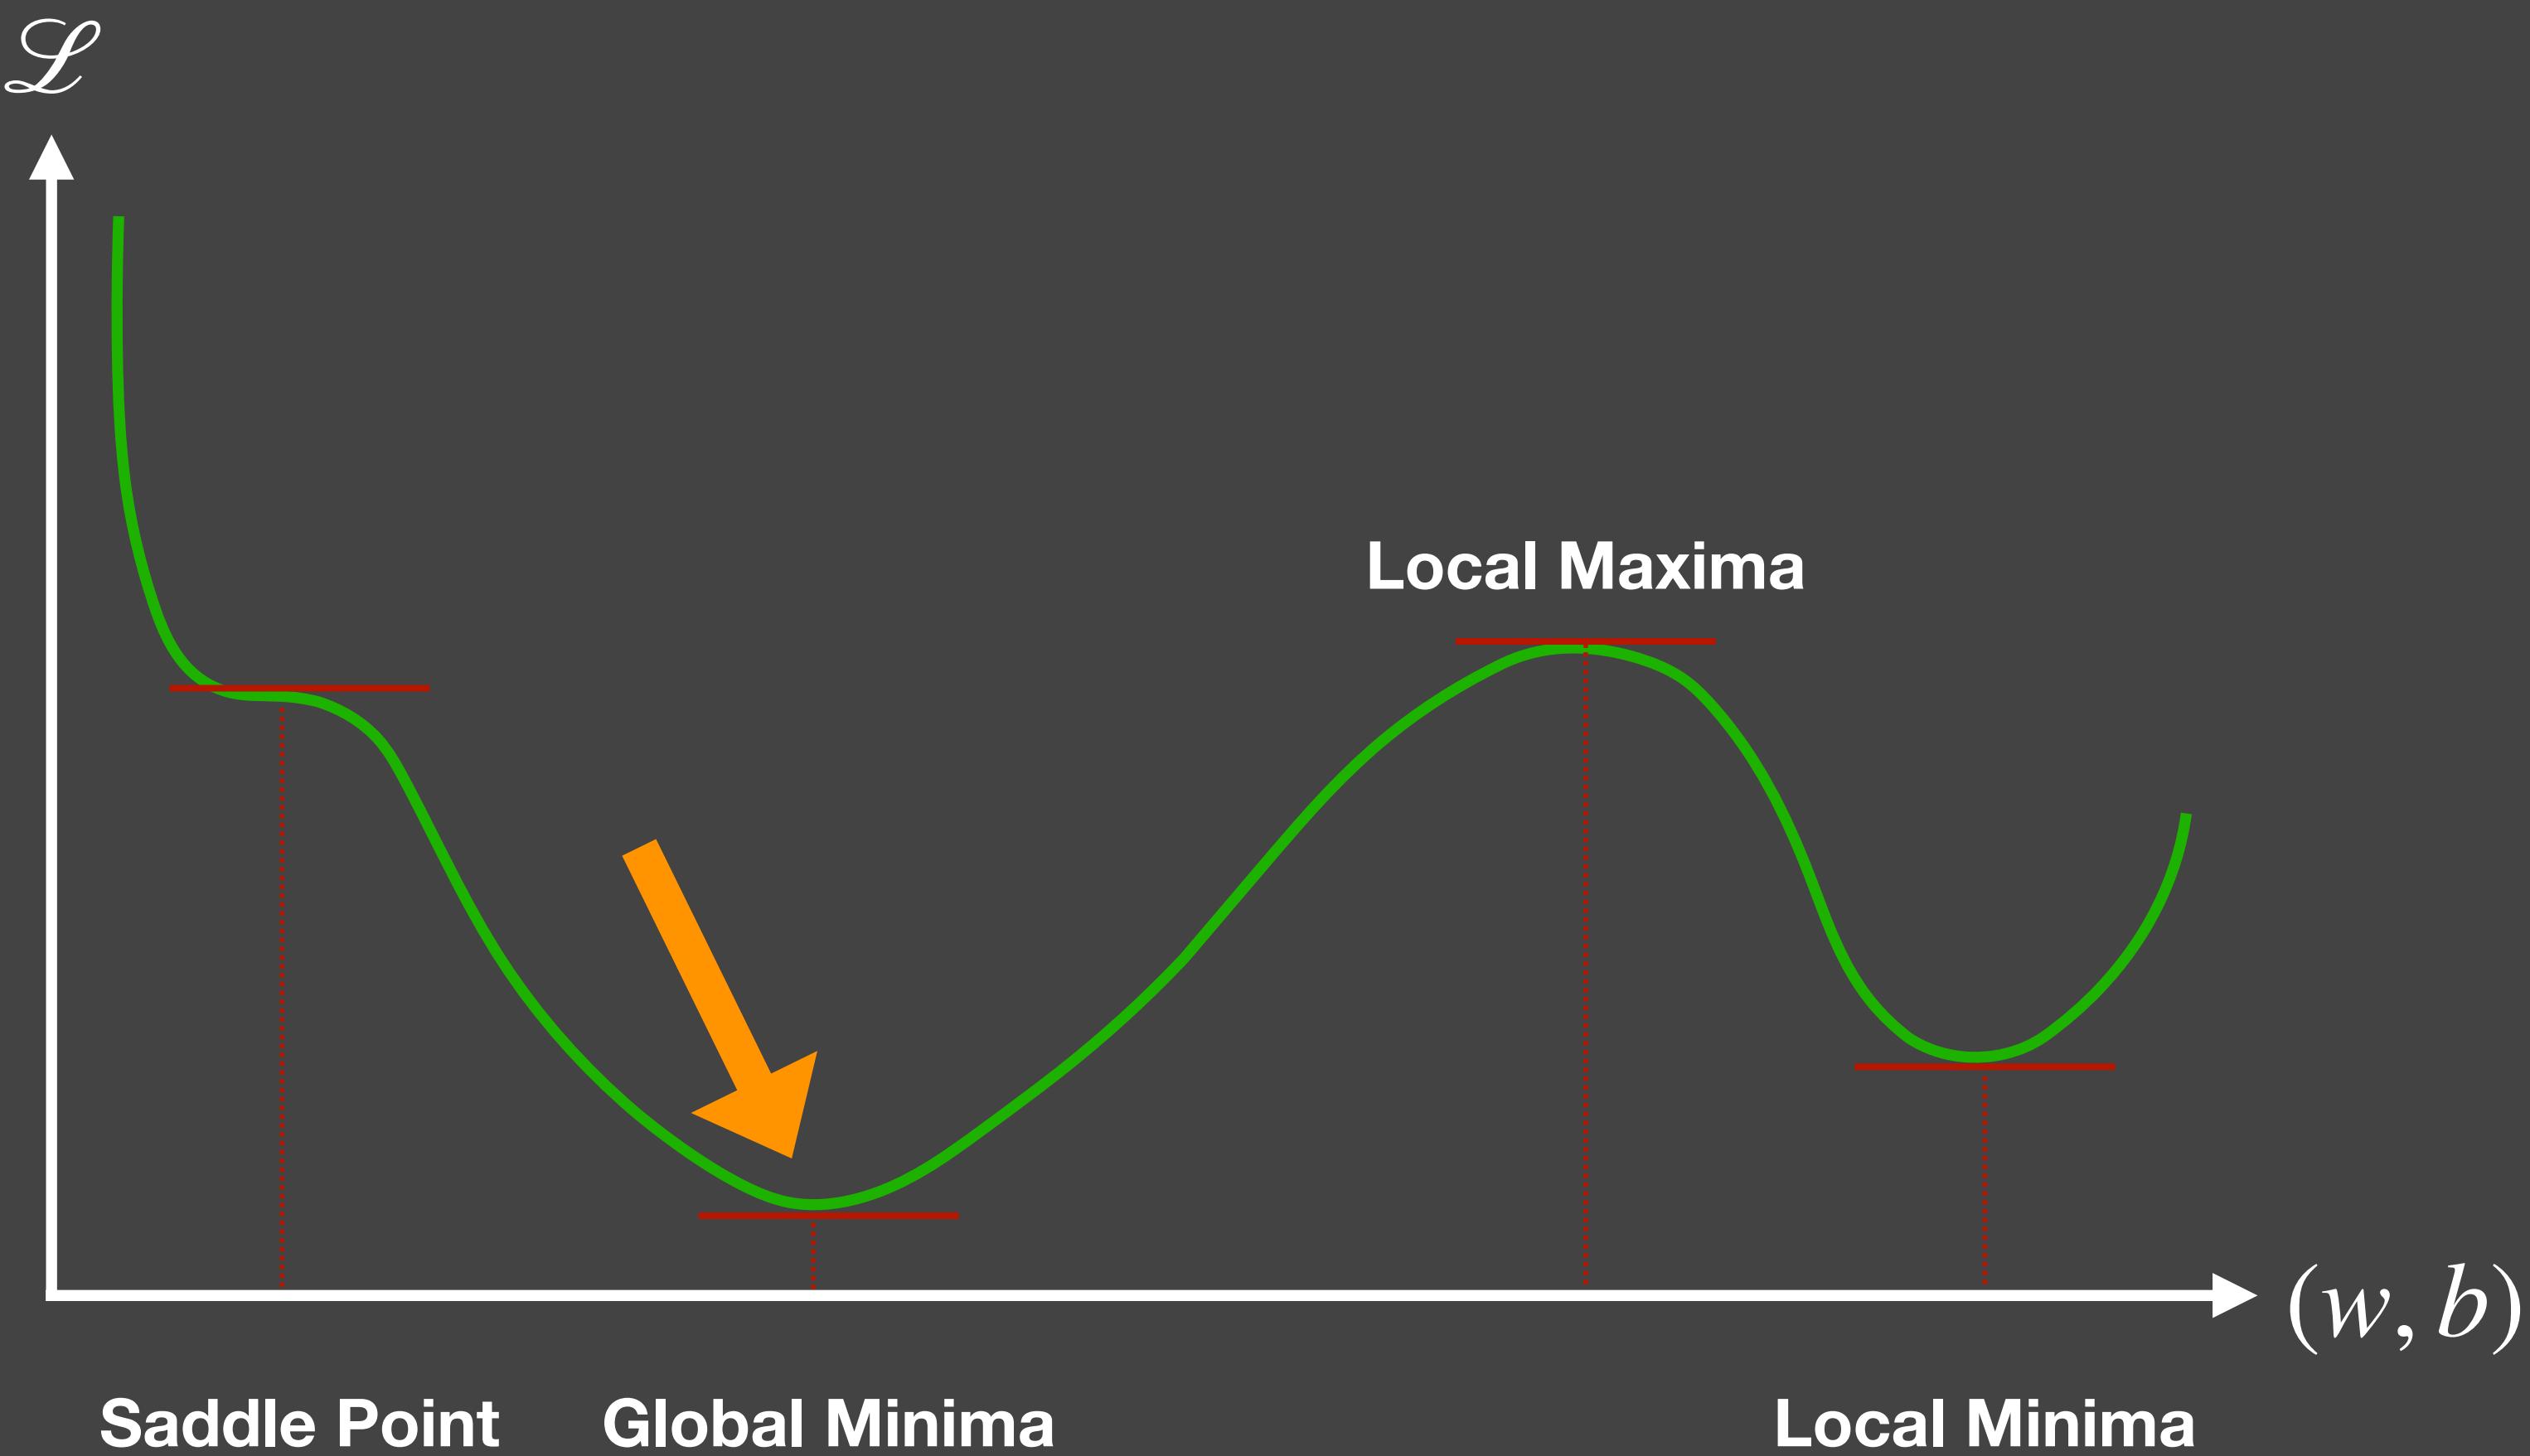
# Optimization Theory - Stationary Points

- Stationary Points
  - Minima, maxima, saddle points
  - Derivative is zero



# Optimization Theory - Stationary Points

- Stationary Points
  - Minima, maxima, saddle points
  - Derivative is zero

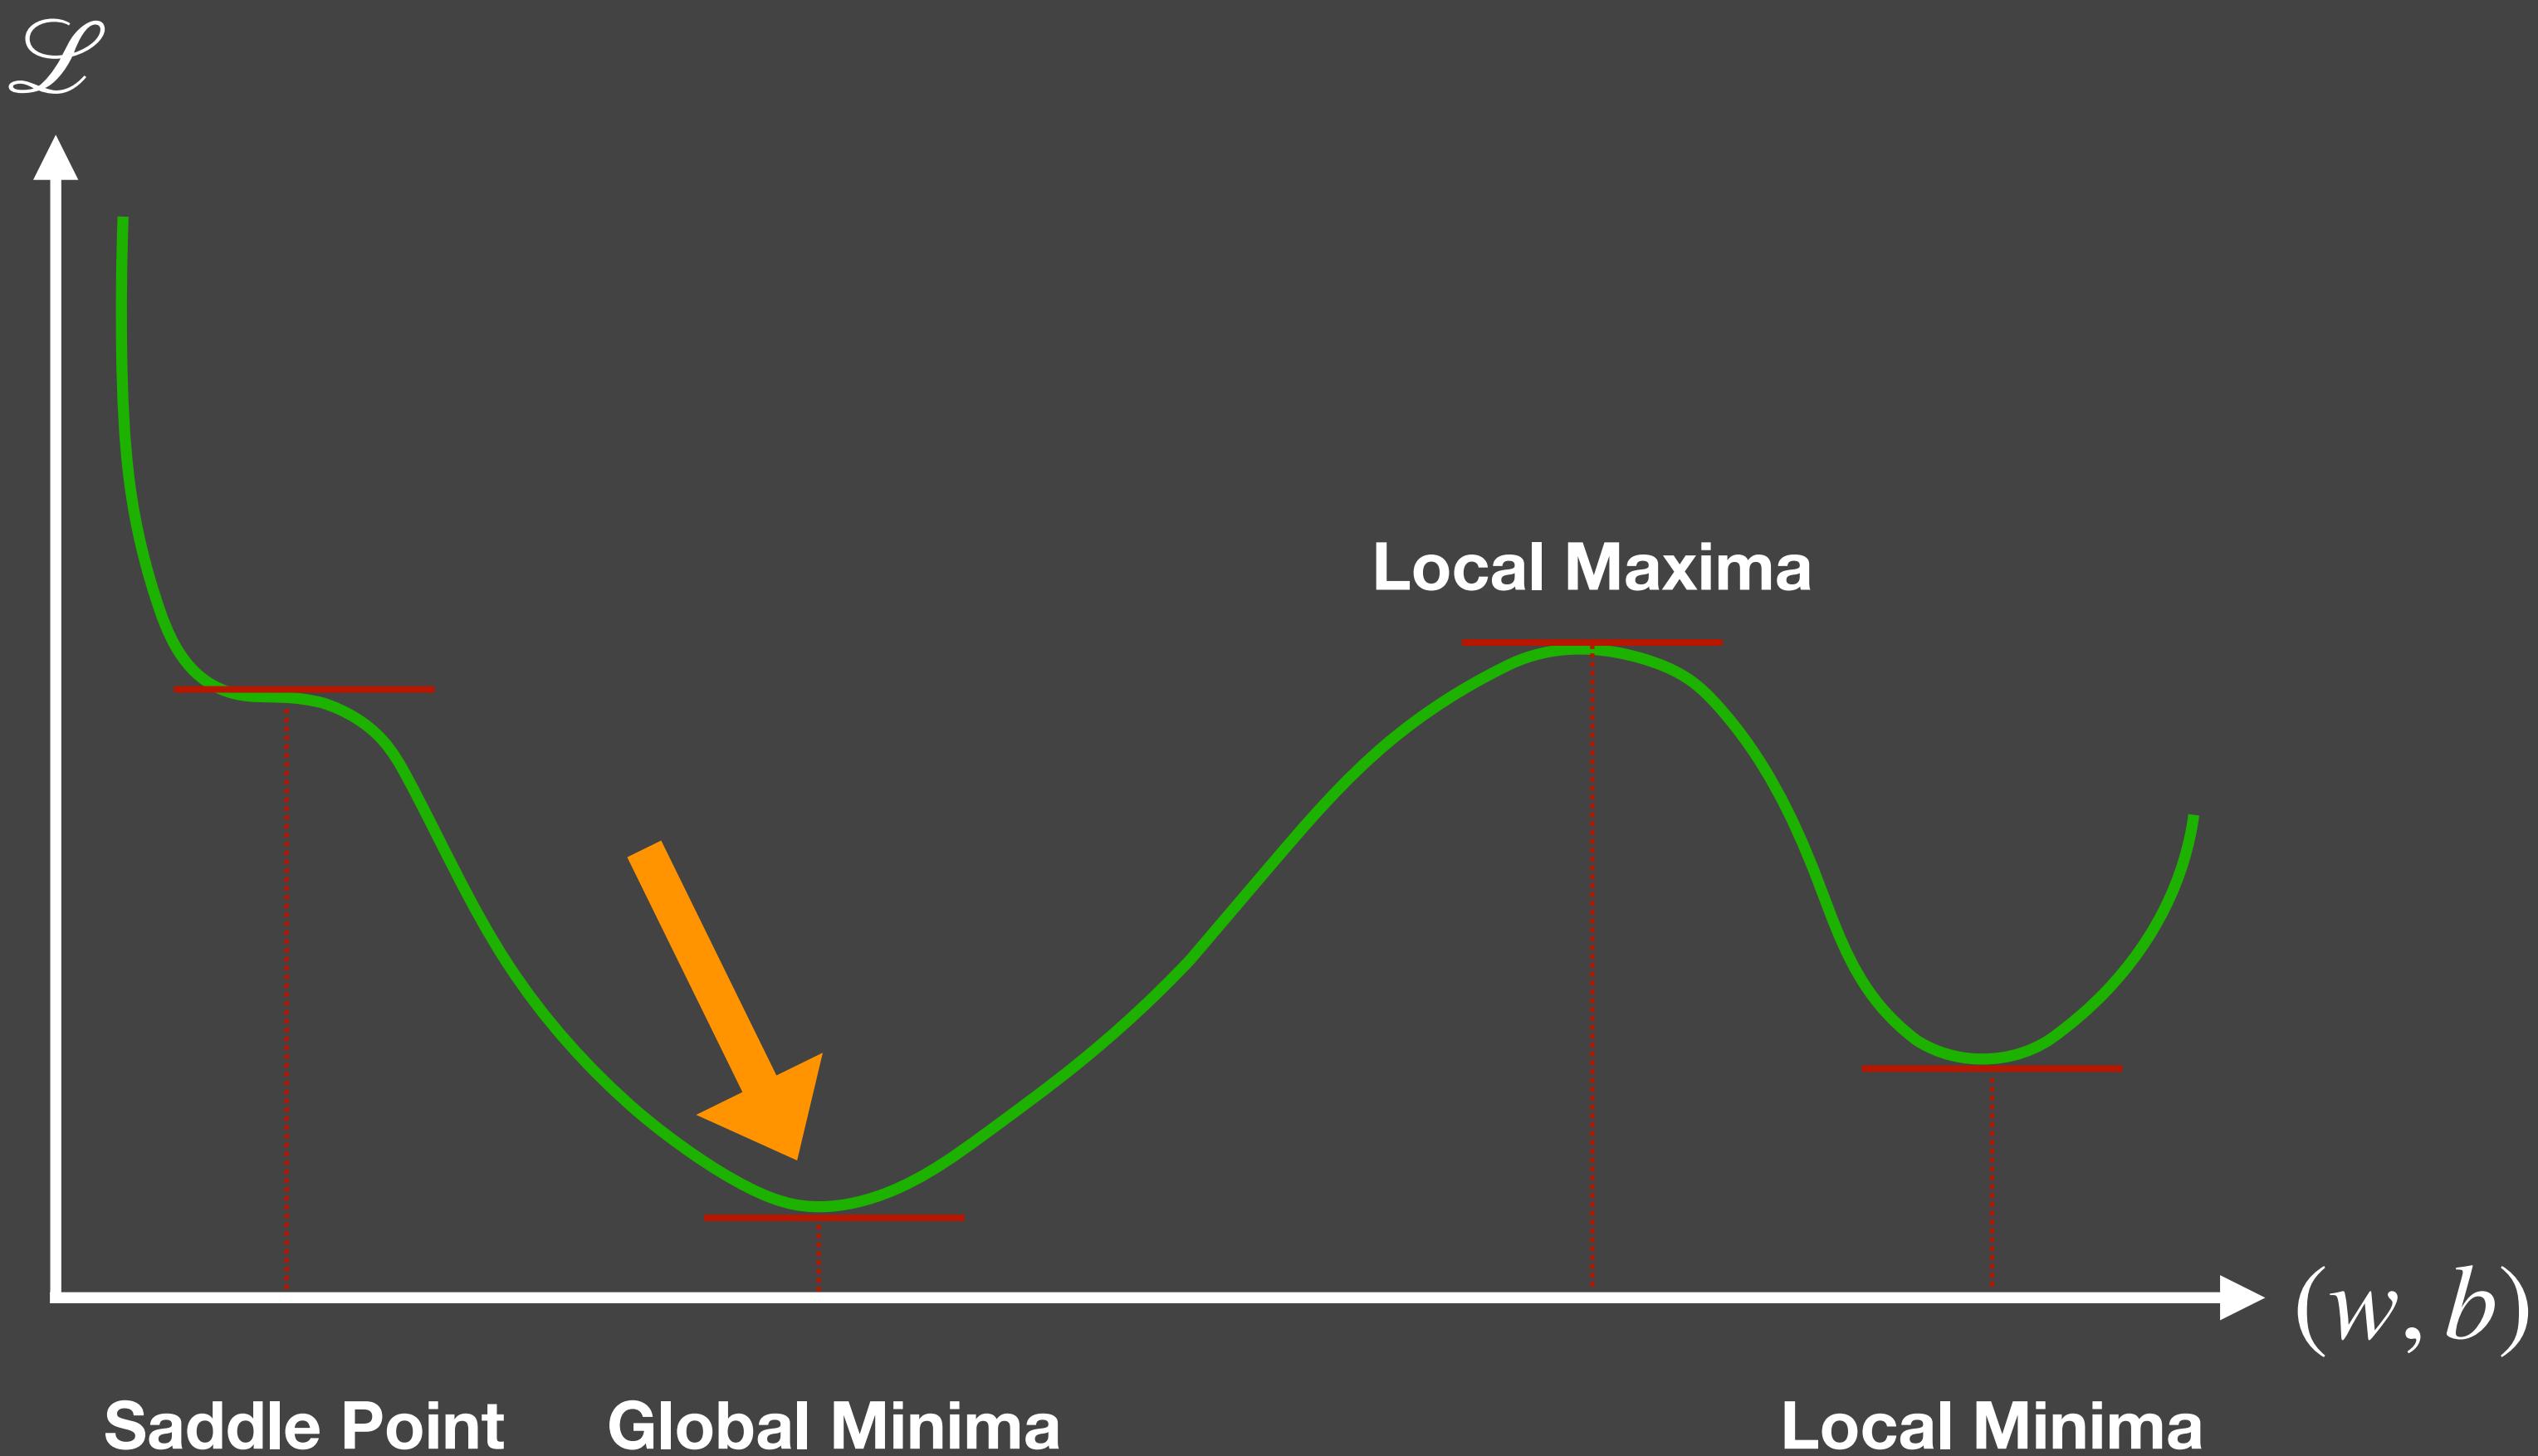


# Optimization Theory - Stationary Points

- Stationary Points
  - Minima, maxima, saddle points
  - Derivative is zero
- Normal Equation

$$\frac{\partial \mathcal{L}}{\partial w} = 0$$

$$\frac{\partial \mathcal{L}}{\partial b} = 0$$



# Optimization Theory - Differentiation and Gradients

- Univariate functions

- $f : R \rightarrow R$
- Single input, single output
- Differentiation with single variable

$$f(x) = x^2 \quad (1)$$

$$\frac{df(x)}{dx} = 2x \quad (2)$$

$$f(x) \approx f(a) + \frac{df(a)}{dx}(x - a) \quad (3)$$

# Optimization Theory - Differentiation and Gradients

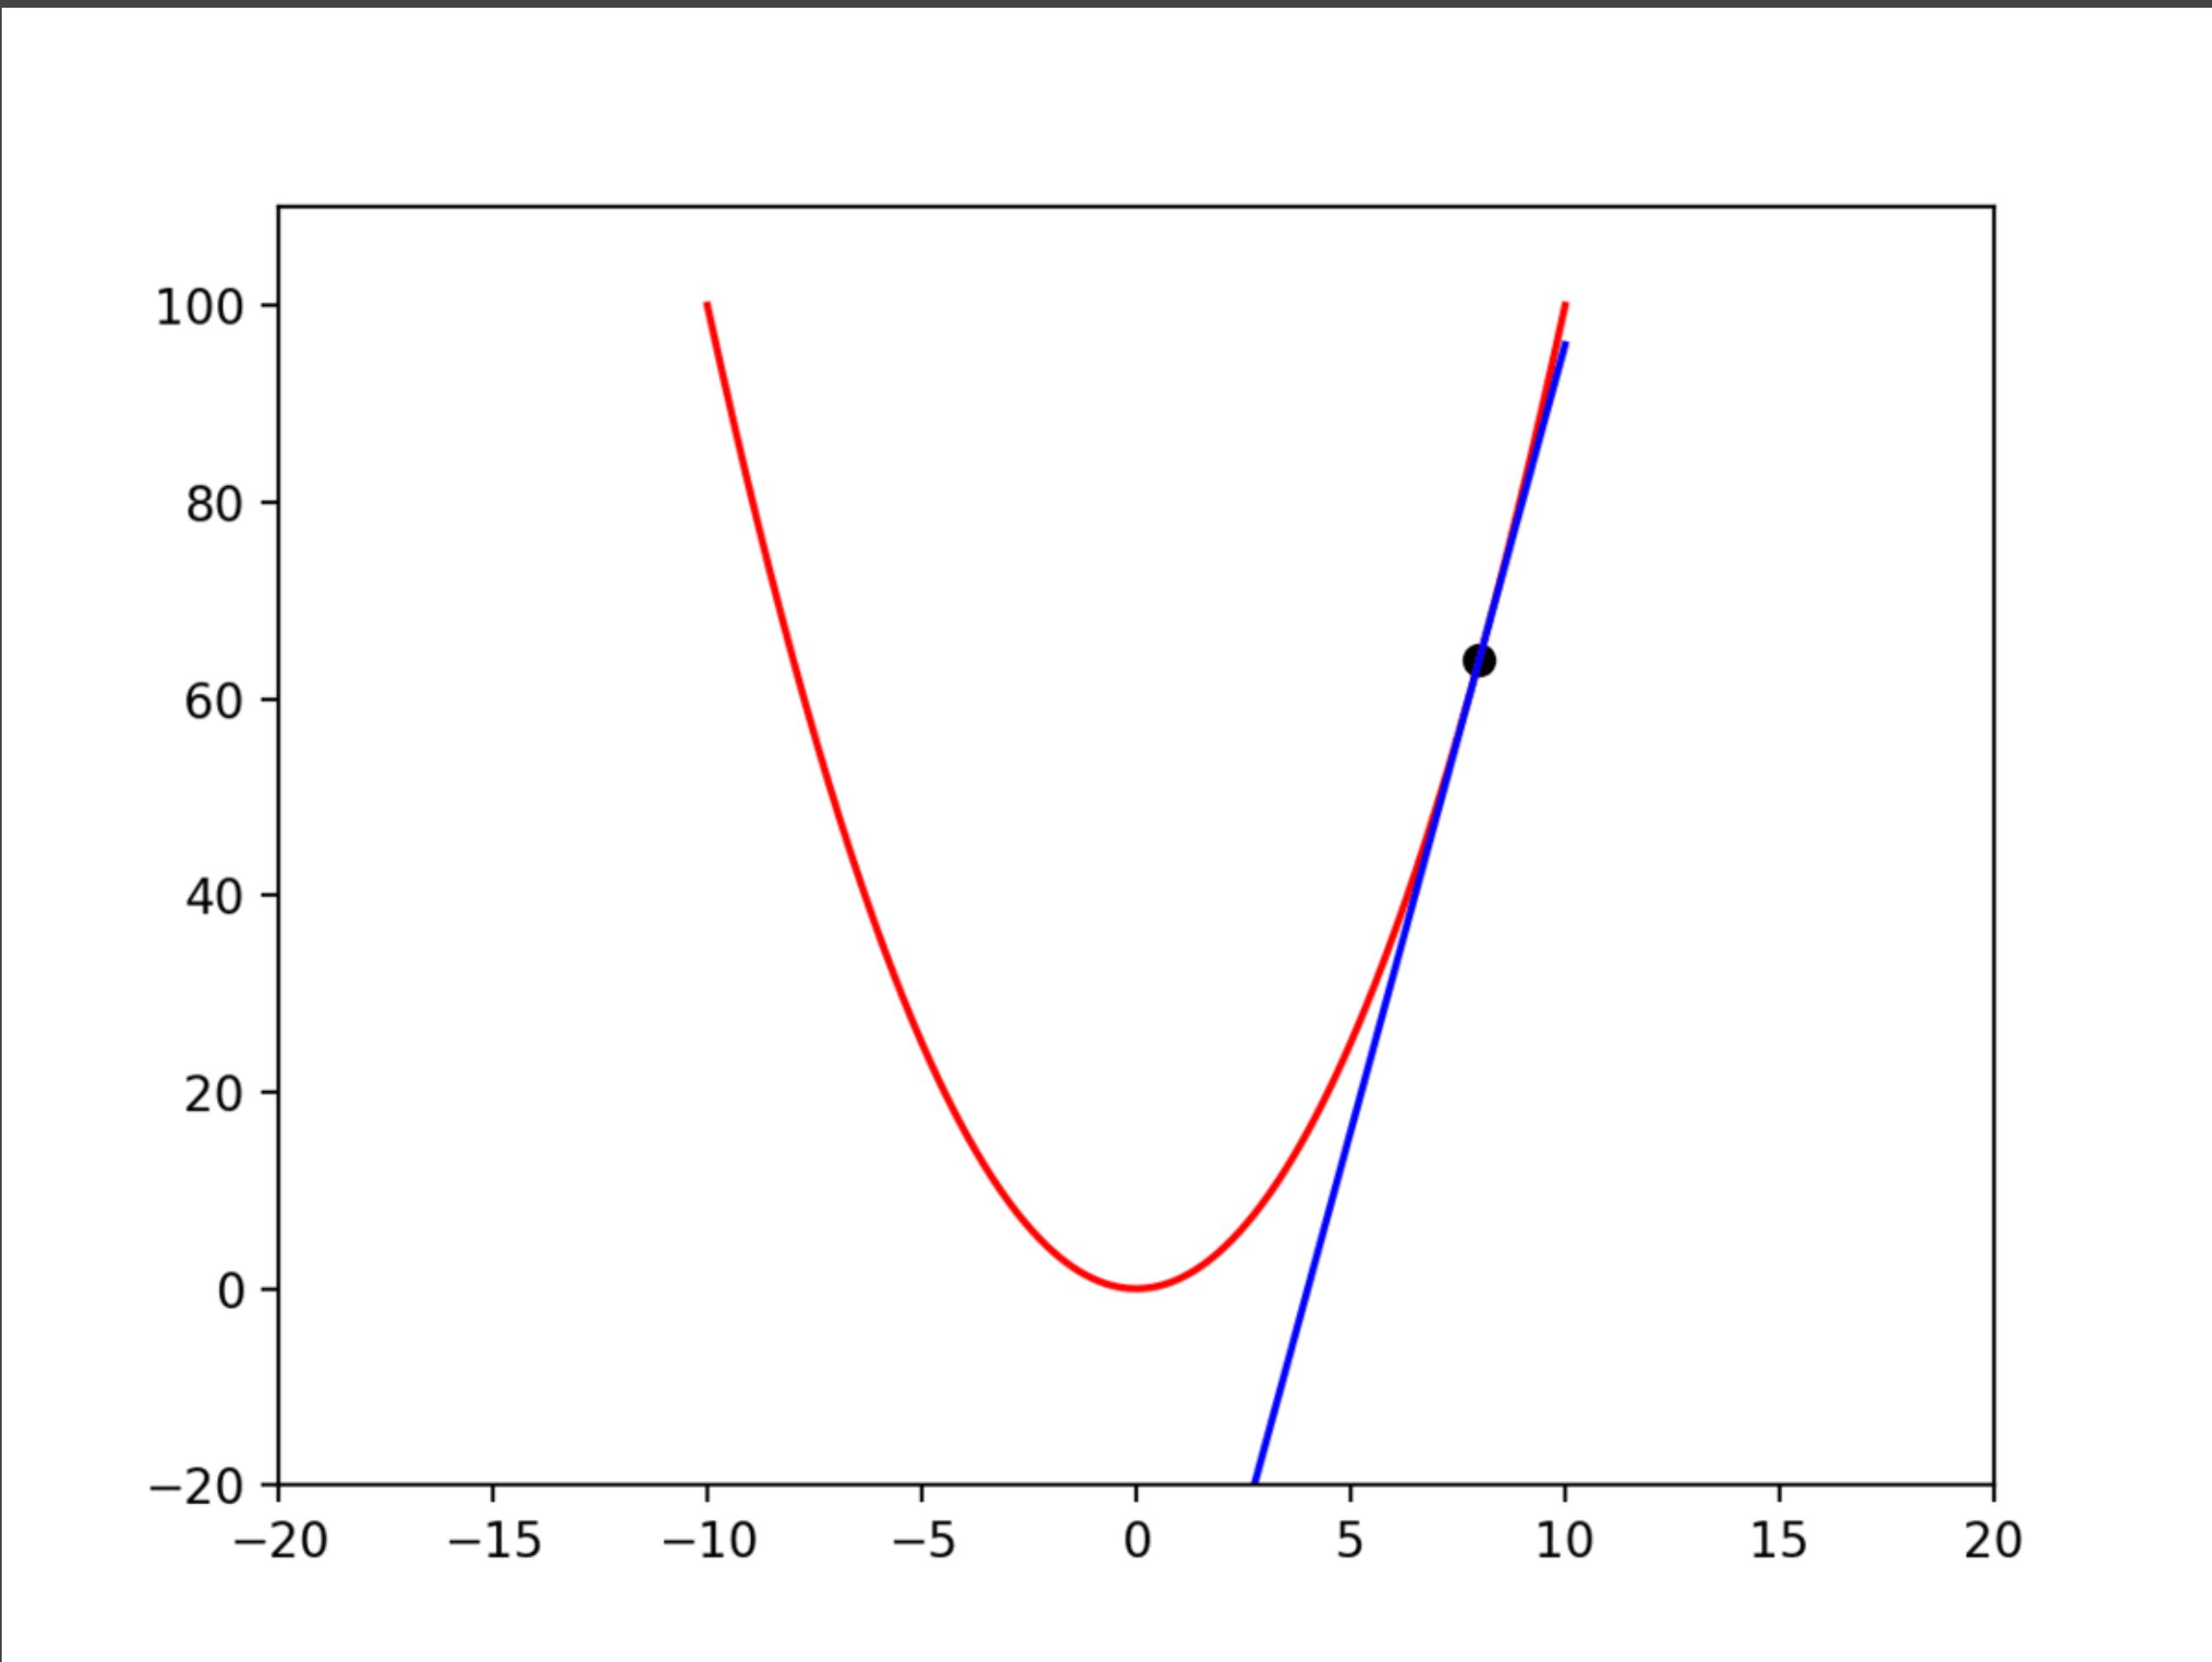
- Univariate functions

- $f: R \rightarrow R$
- Single input, single output
- Differentiation with single variable

$$f(x) = x^2 \quad (1)$$

$$\frac{df(x)}{dx} = 2x \quad (2)$$

$$f(x) \approx f(a) + \frac{df(a)}{dx}(x - a) \quad (3)$$



# Optimization Theory - Differentiation and Gradients

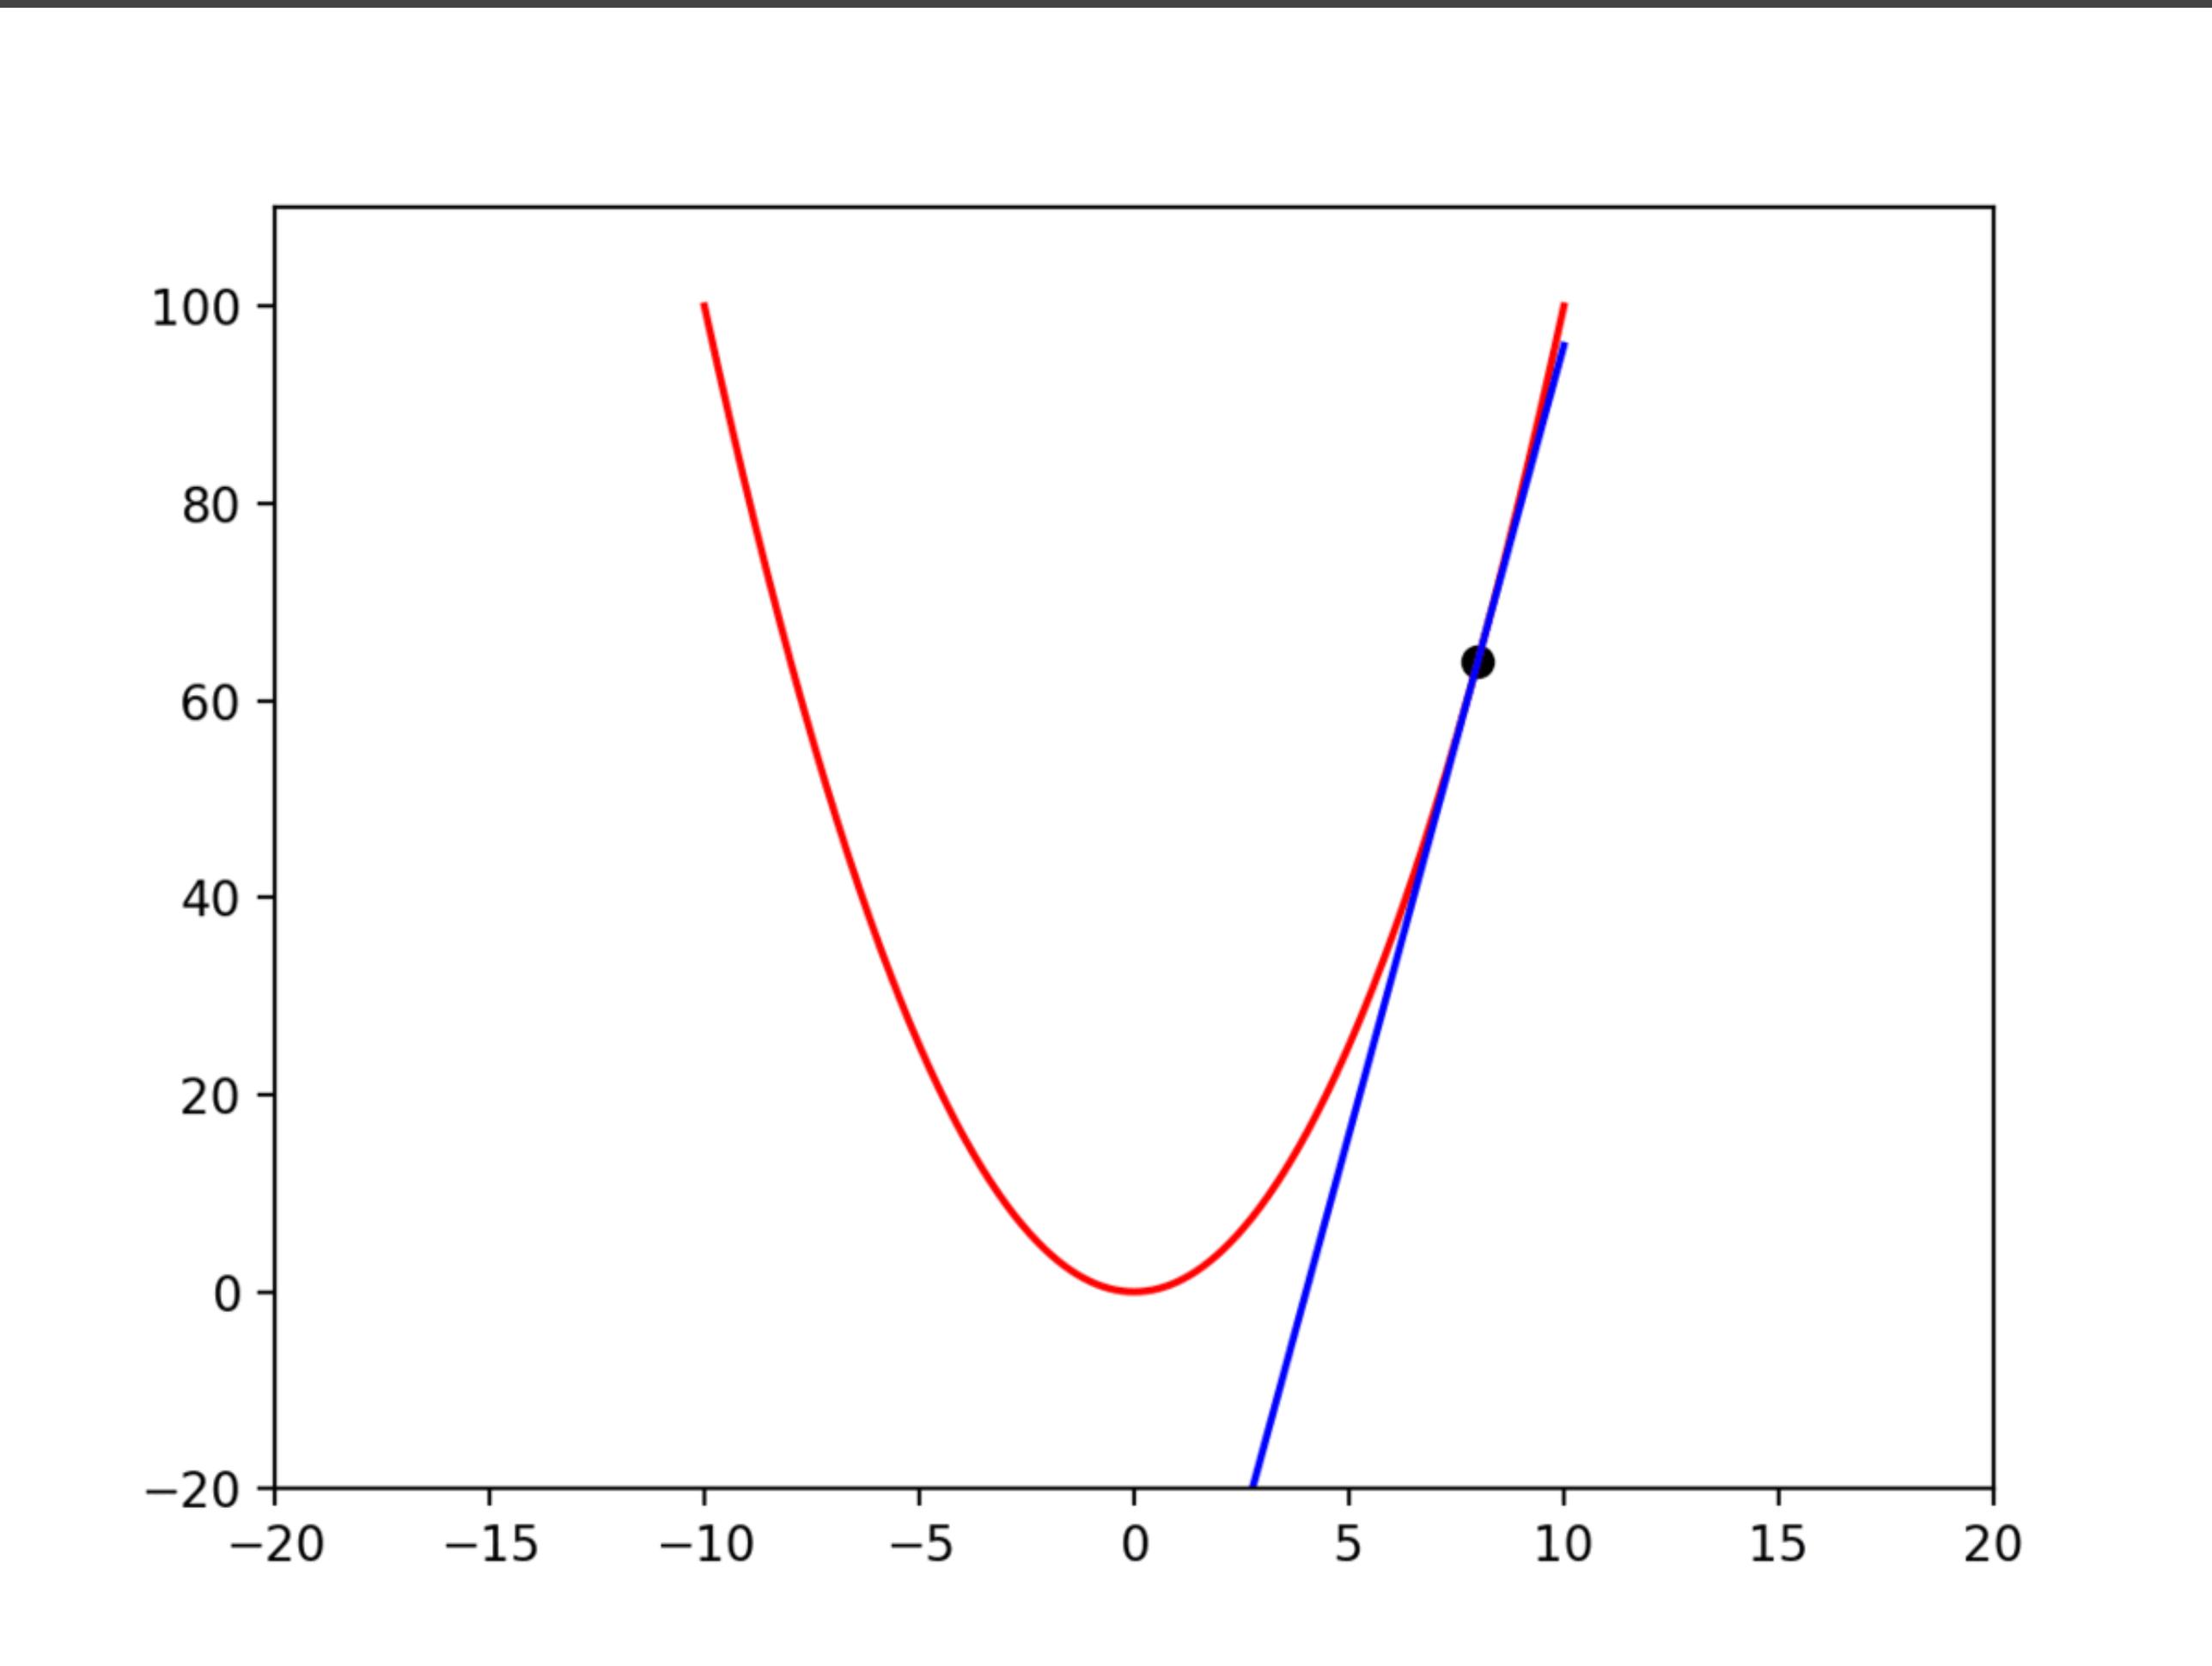
- Univariate functions

- $f: R \rightarrow R$
- Single input, single output
- Differentiation with single variable

$$f(x) = x^2 \quad (1)$$

$$\frac{df(x)}{dx} = 2x \quad (2)$$

$$f(x) \approx f(a) + \frac{df(a)}{dx}(x - a) \quad (3)$$



# Optimization Theory - Differentiation and Gradients

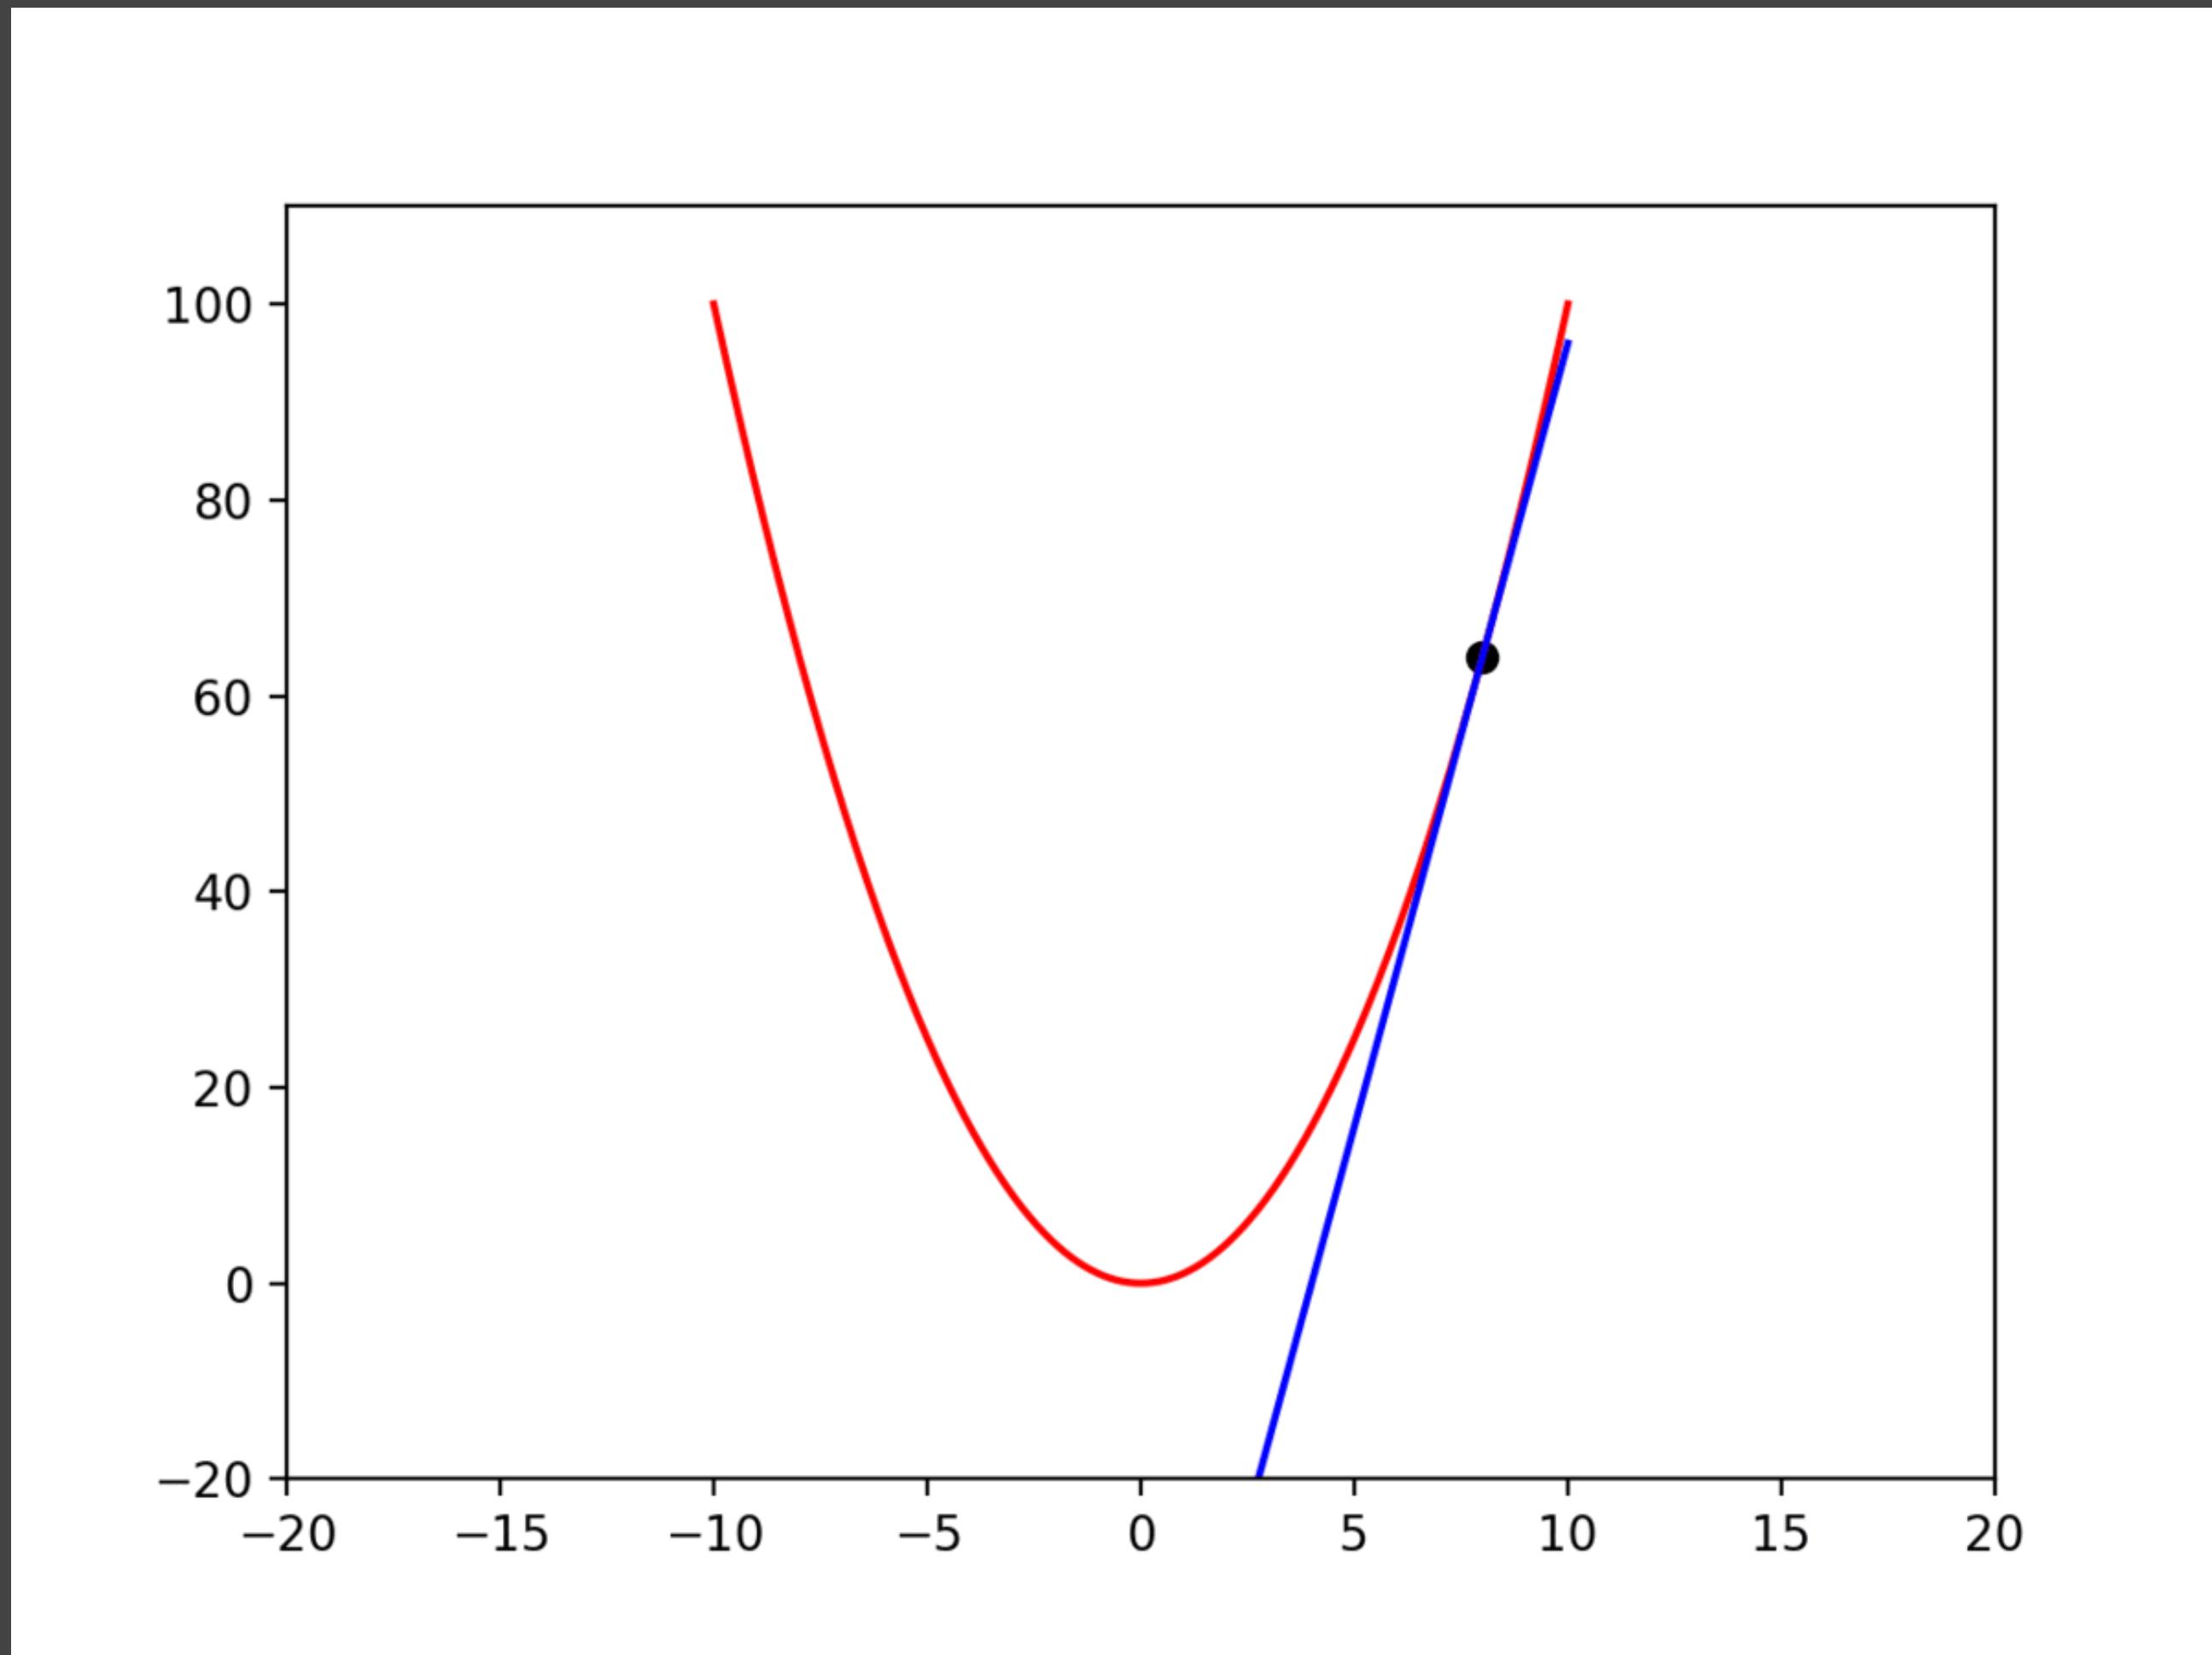
- Univariate functions

- $f: R \rightarrow R$
- Single input, single output
- Differentiation with single variable

$$f(x) = x^2 \quad (1)$$

$$\frac{df(x)}{dx} = 2x \quad (2)$$

$$f(x) \approx f(a) + \frac{df(a)}{dx}(x - a) \quad (3)$$



# Optimization Theory - Differentiation and Gradients

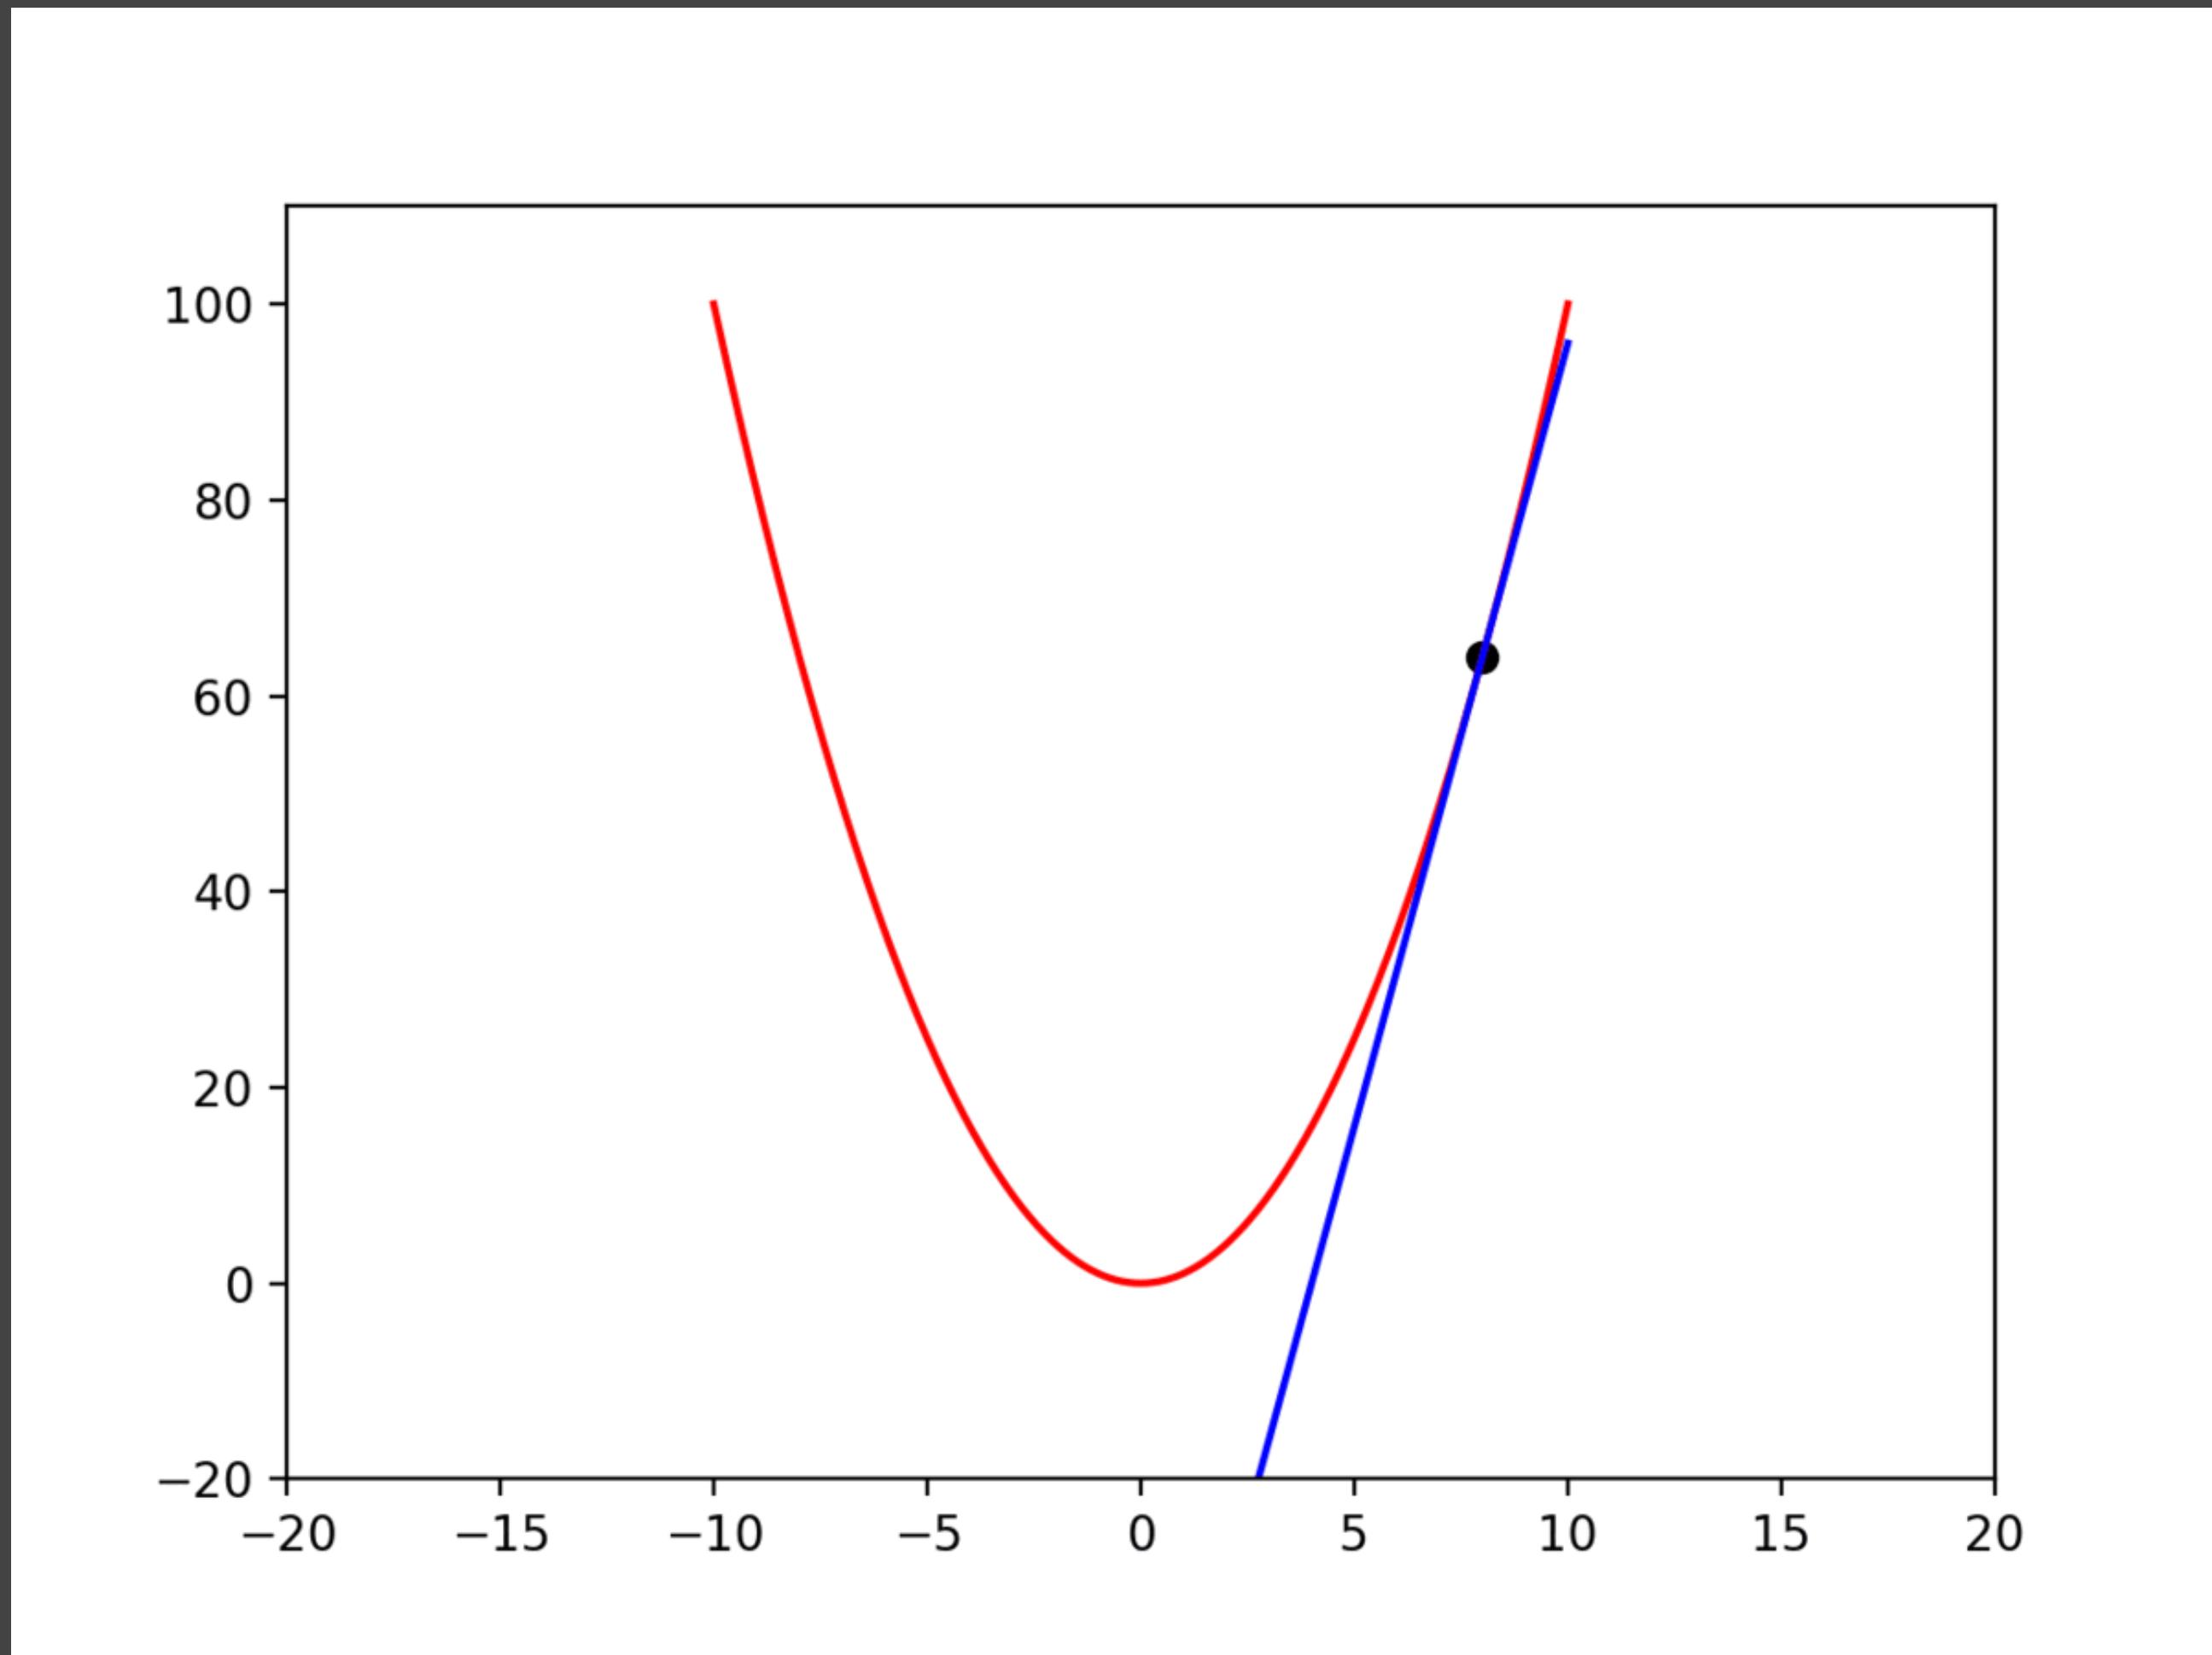
- Univariate functions

- $f: R \rightarrow R$
- Single input, single output
- Differentiation with single variable

$$f(x) = x^2 \quad (1)$$

$$\frac{df(x)}{dx} = 2x \quad (2)$$

$$f(x) \approx f(a) + \frac{df(a)}{dx}(x - a) \quad (3)$$



# Optimization Theory - Differentiation and Gradients

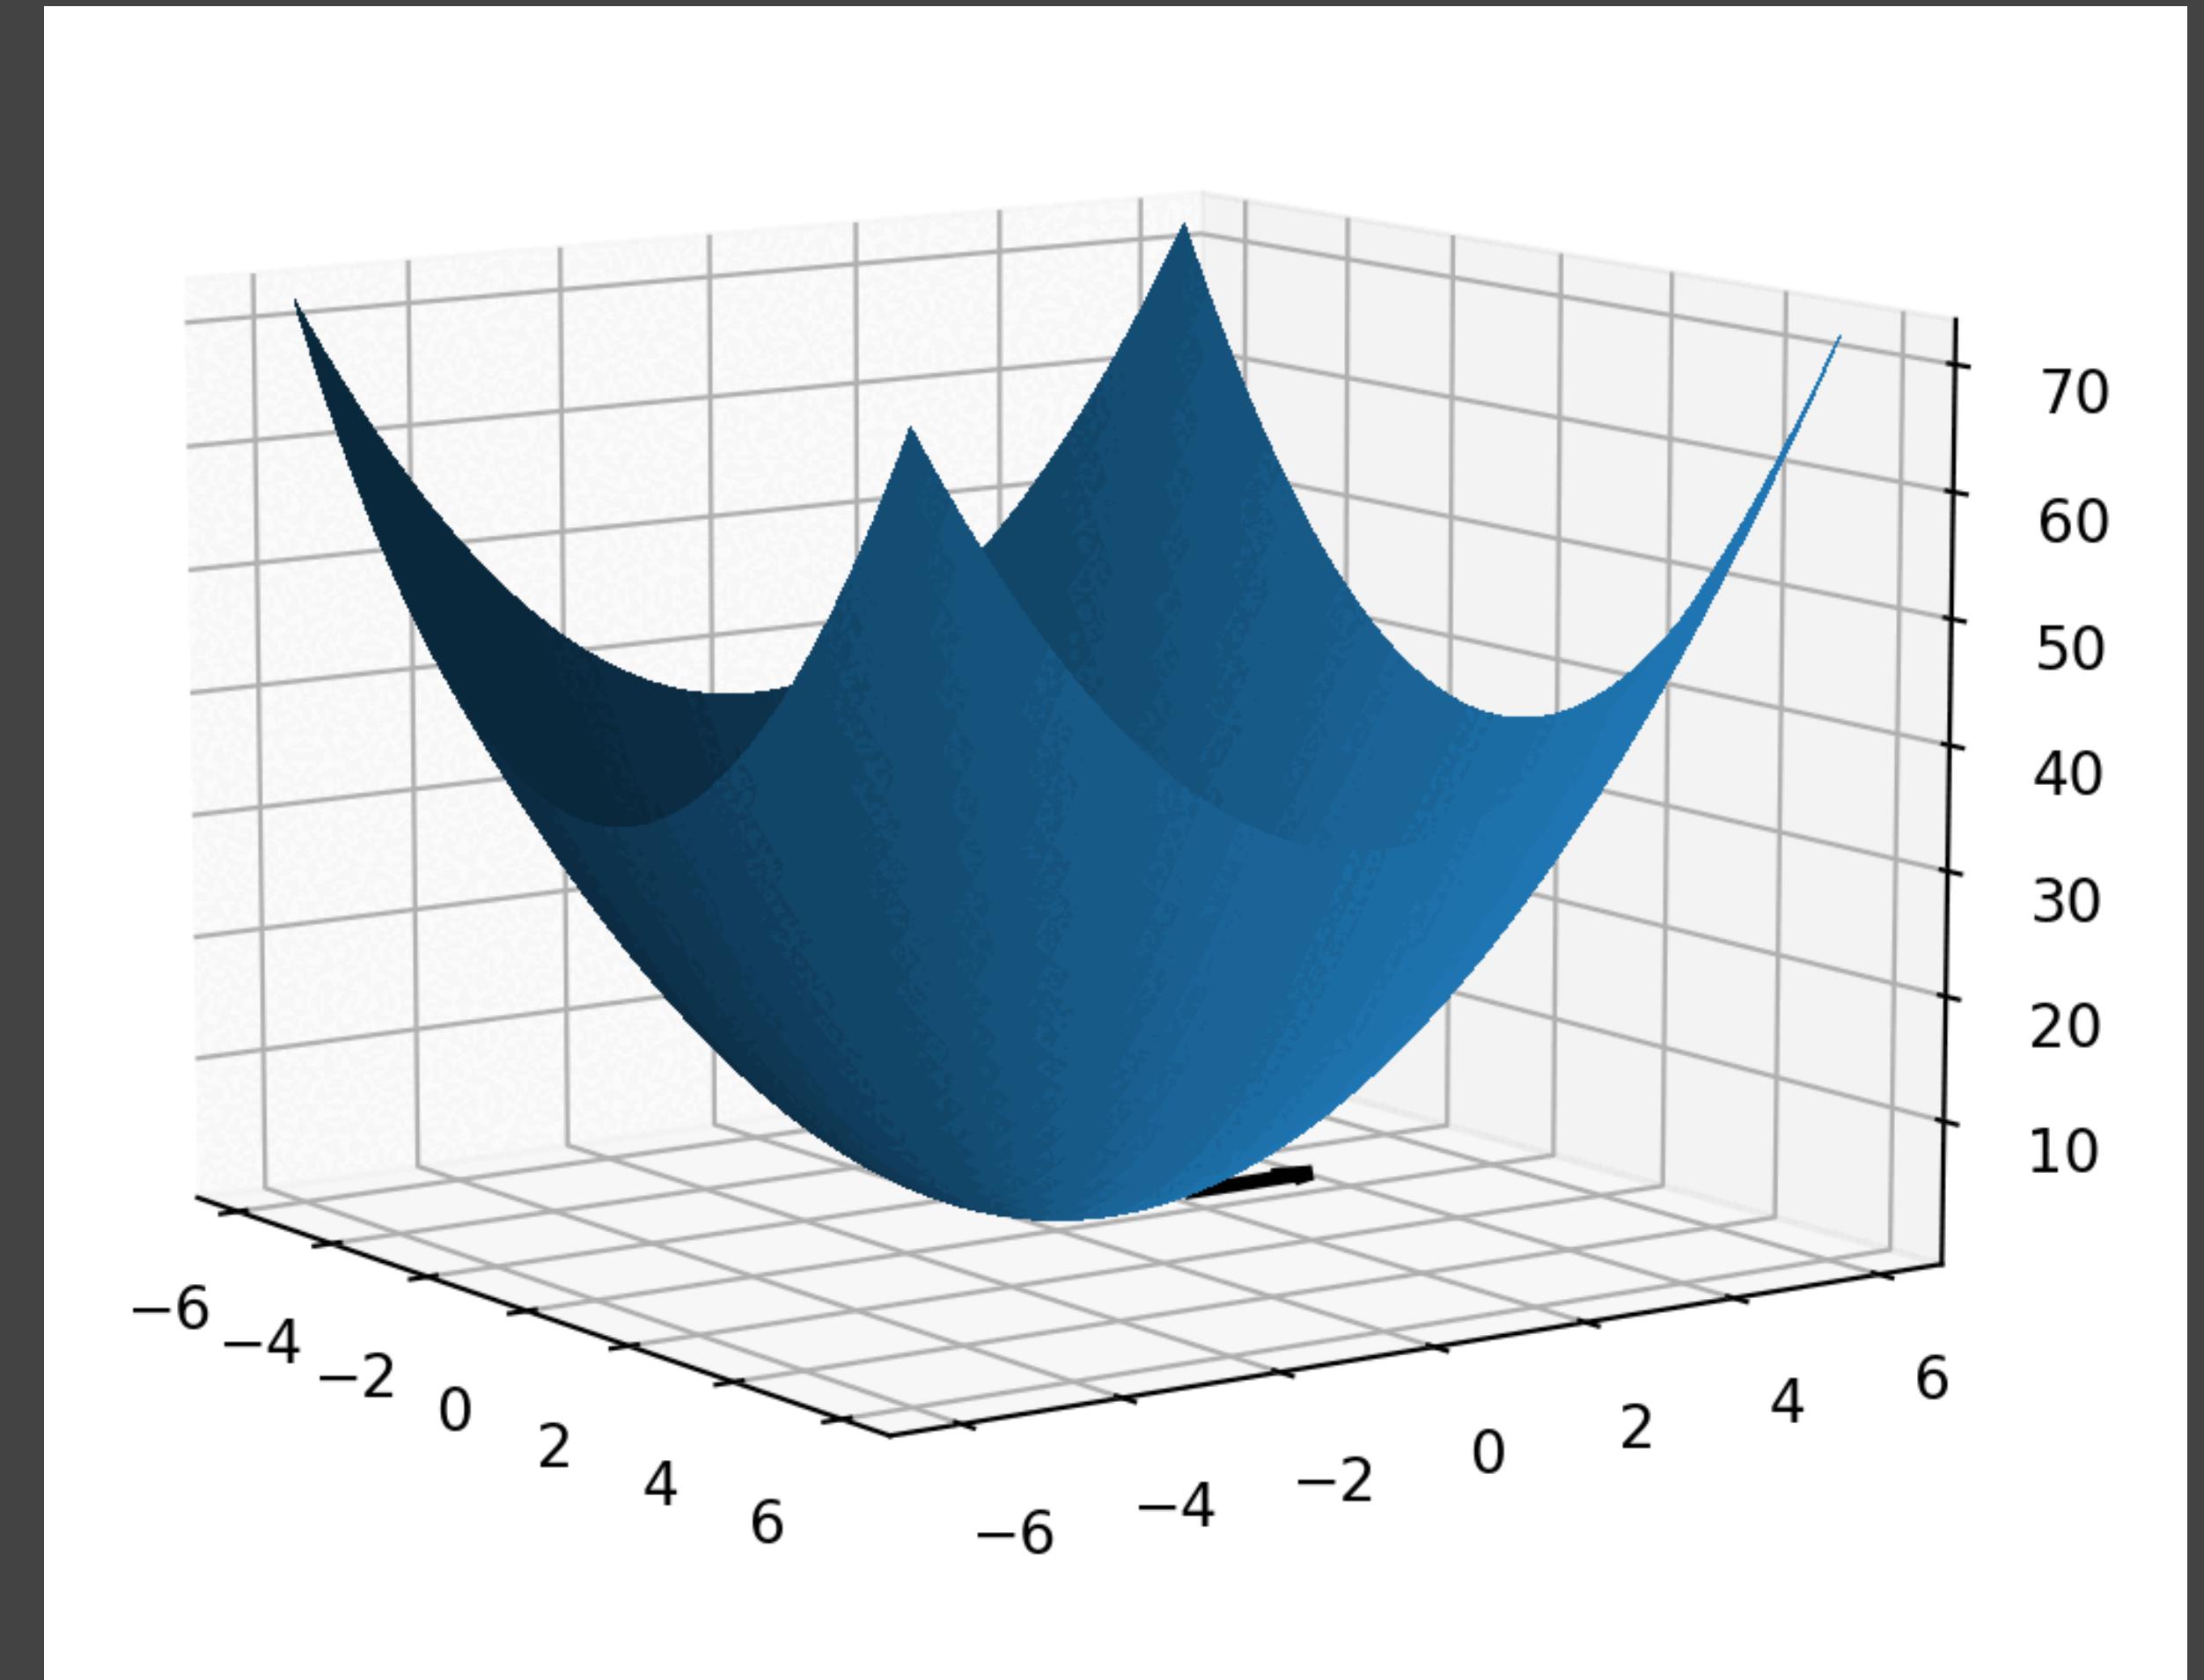
- Multivariate functions

- $g : R^2 \rightarrow R$
- Multiple input, single output
- Partial differentiation with each of multiple variables

$$g(x) = x^T \cdot x$$

$$\nabla g(x) = \left[ \frac{\partial g(x)}{\partial x_1}, \frac{\partial g(x)}{\partial x_2} \right]^T$$

$$\nabla g(x) = [2x_1, 2x_2]^T$$



# Optimization Theory - Differentiation and Gradients

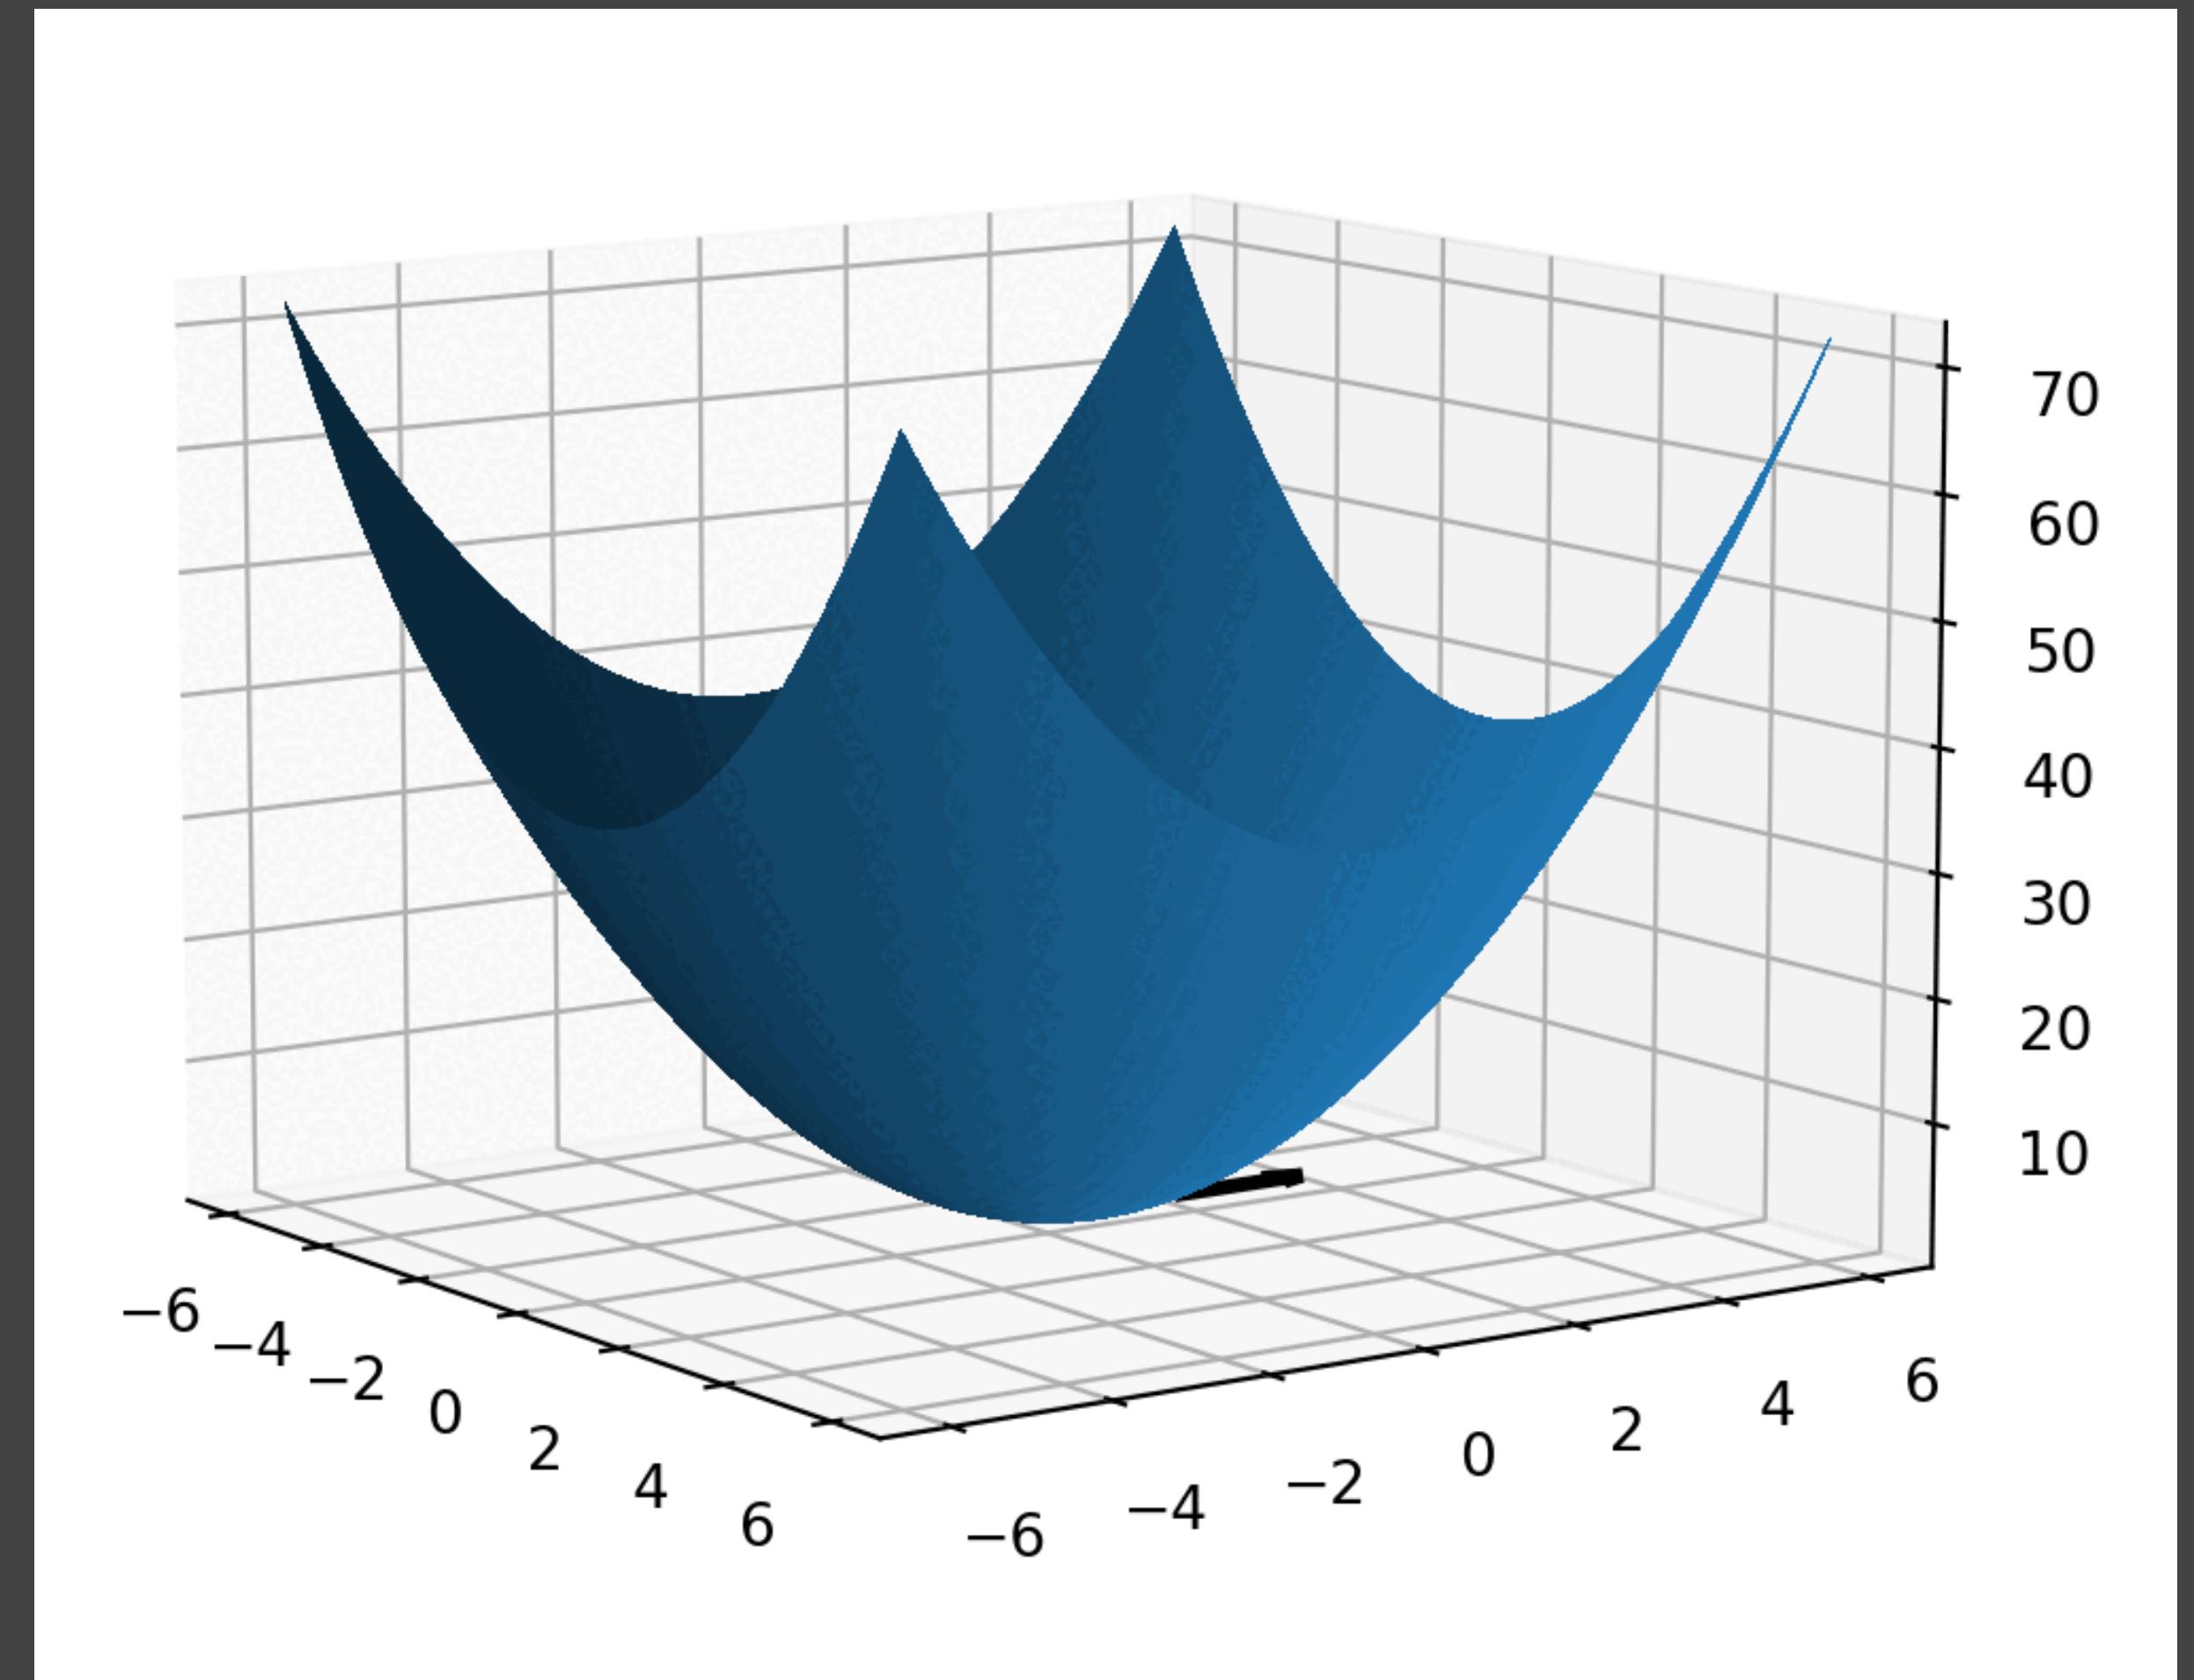
- Multivariate functions

- $g : R^2 \rightarrow R$
- Multiple input, single output
- Partial differentiation with each of multiple variables

$$g(x) = x^T \cdot x$$

$$\nabla g(x) = \left[ \frac{\partial g(x)}{\partial x_1}, \frac{\partial g(x)}{\partial x_2} \right]^T$$

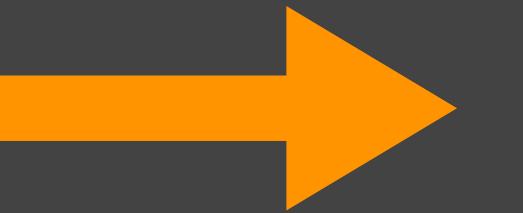
$$\nabla g(x) = [2x_1, 2x_2]^T$$



# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

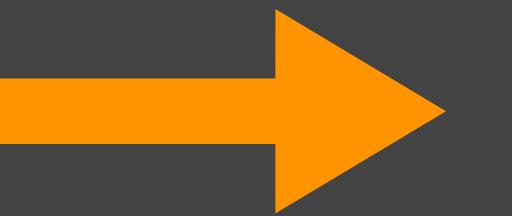
$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



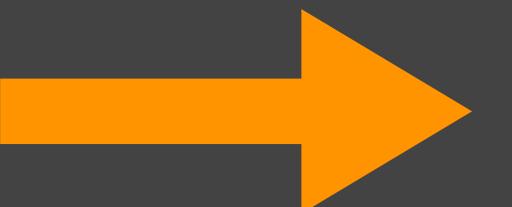
$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

- Gradient Descent

- Iterative optimization algorithm
- Based on gradient vector



*Init*  $\hat{w}^0 = [b^0, w^0]$   
*while*(stop condition)

$$\hat{w}^{k+1} = \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k)$$
$$k = k + 1$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



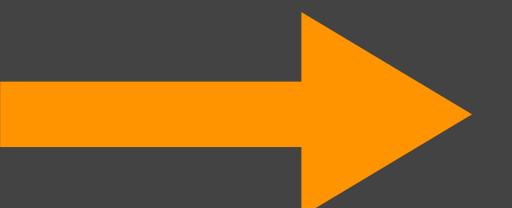
$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

- Gradient Descent

- Iterative optimization algorithm
- Based on gradient vector



$$Init \quad \hat{w}^0 = [b^0, w^0]$$

*while(stop condition)*

$$\begin{aligned}\hat{w}^{k+1} &= \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k) \\ k &= k + 1\end{aligned}$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



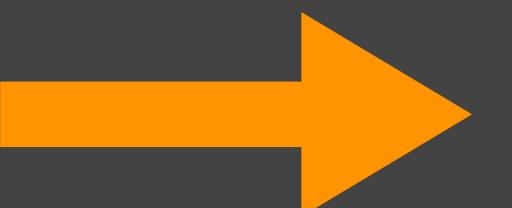
$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

- Gradient Descent

- Iterative optimization algorithm
- Based on gradient vector



$$Init \quad \hat{w}^0 = [b^0, w^0]$$

*while(stop condition)*

$$\hat{w}^{k+1} = \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k)$$
$$k = k + 1$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



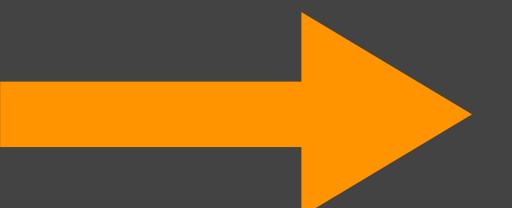
$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

- Gradient Descent

- Iterative optimization algorithm
- Based on gradient vector



$$Init \quad \hat{w}^0 = [b^0, w^0]$$

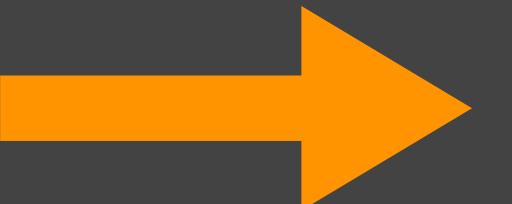
*while(stop condition)*

$$\hat{w}^{k+1} = \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k)$$
$$k = k + 1$$

# Optimization Theory - Gradient Descent

- Reformulation

- Bias and weight in one parameter vector
- Updating model and cost functions to remove bias extension in model and cost functions



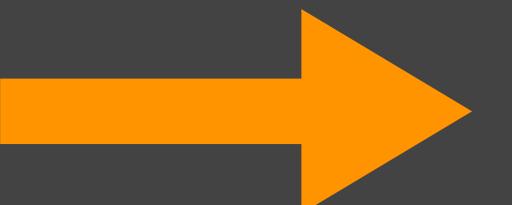
$$\hat{w} = [b, w]^T \quad \hat{x} = [1, x]^T$$

$$f_{\hat{w}}(x) = \hat{x}^T \cdot \hat{w}$$

$$\mathcal{L}(\hat{w}) = \frac{1}{N} \sum_{i=1}^N (f_{\hat{w}}(x^i) - y^i)^2$$

- Gradient Descent

- Iterative optimization algorithm
- Based on gradient vector



$$Init \quad \hat{w}^0 = [b^0, w^0]$$

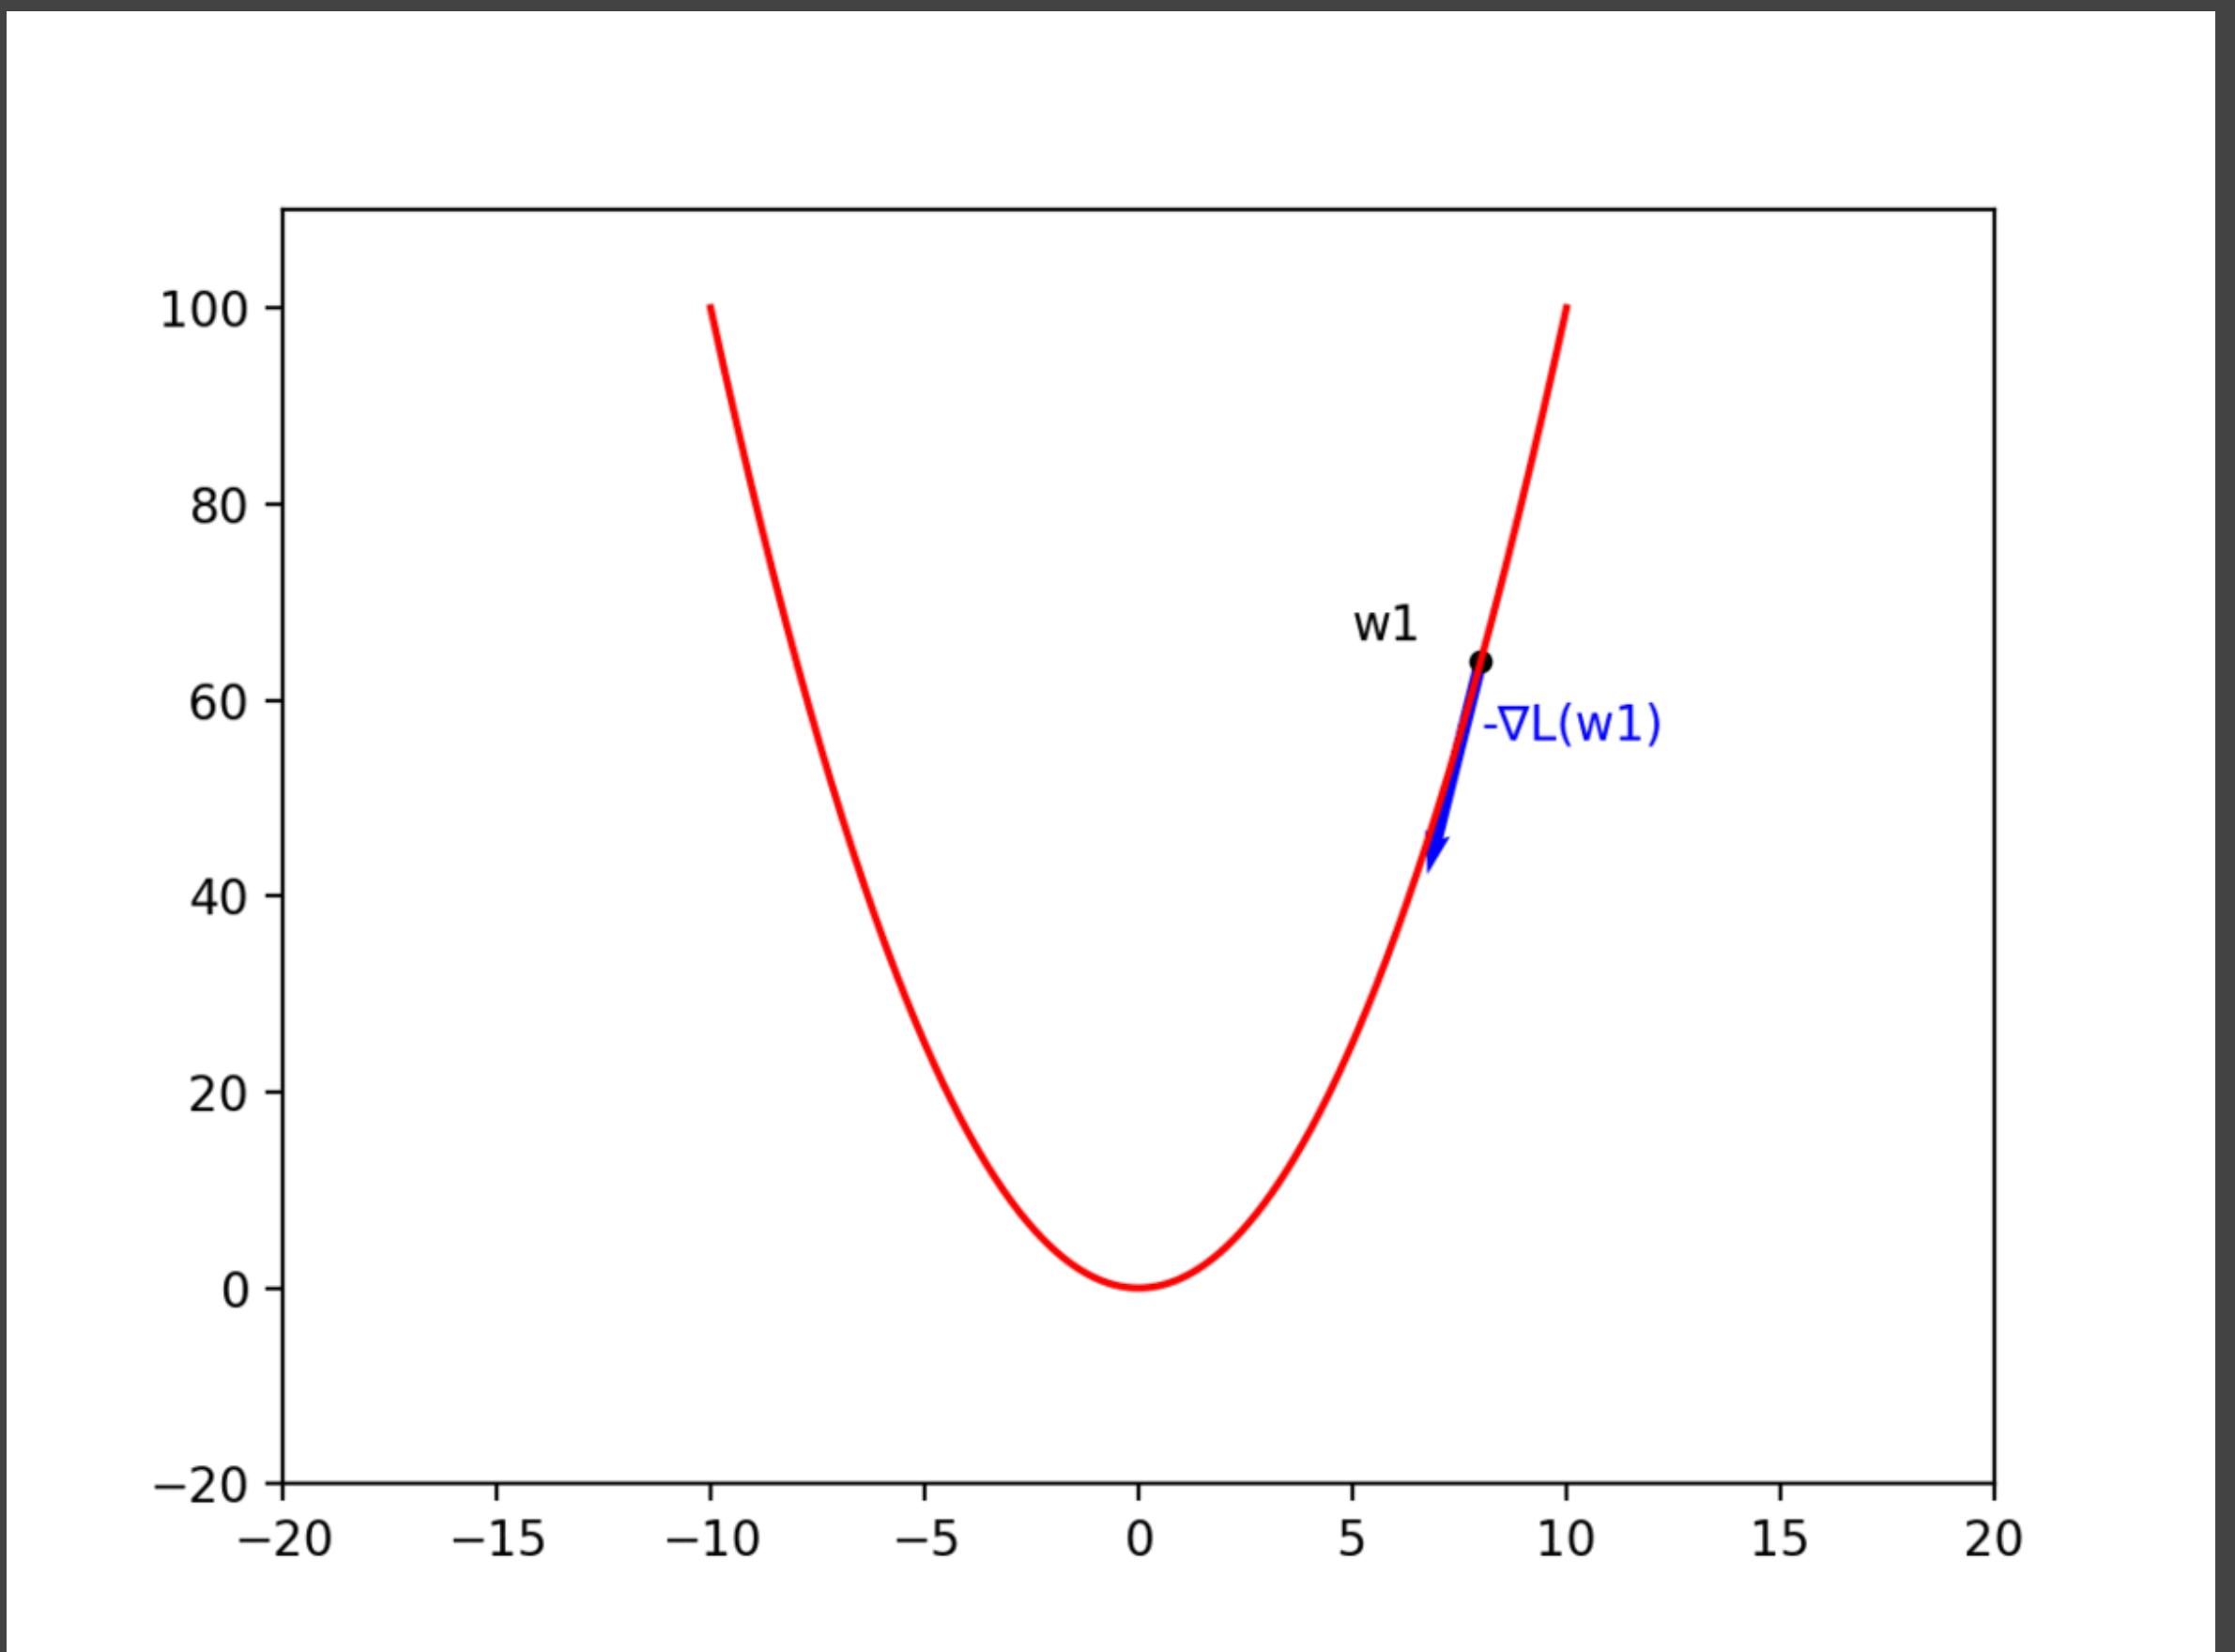
*while(stop condition)*

$$\hat{w}^{k+1} = \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k)$$
$$k = k + 1$$

# Optimization Theory - Gradient Descent

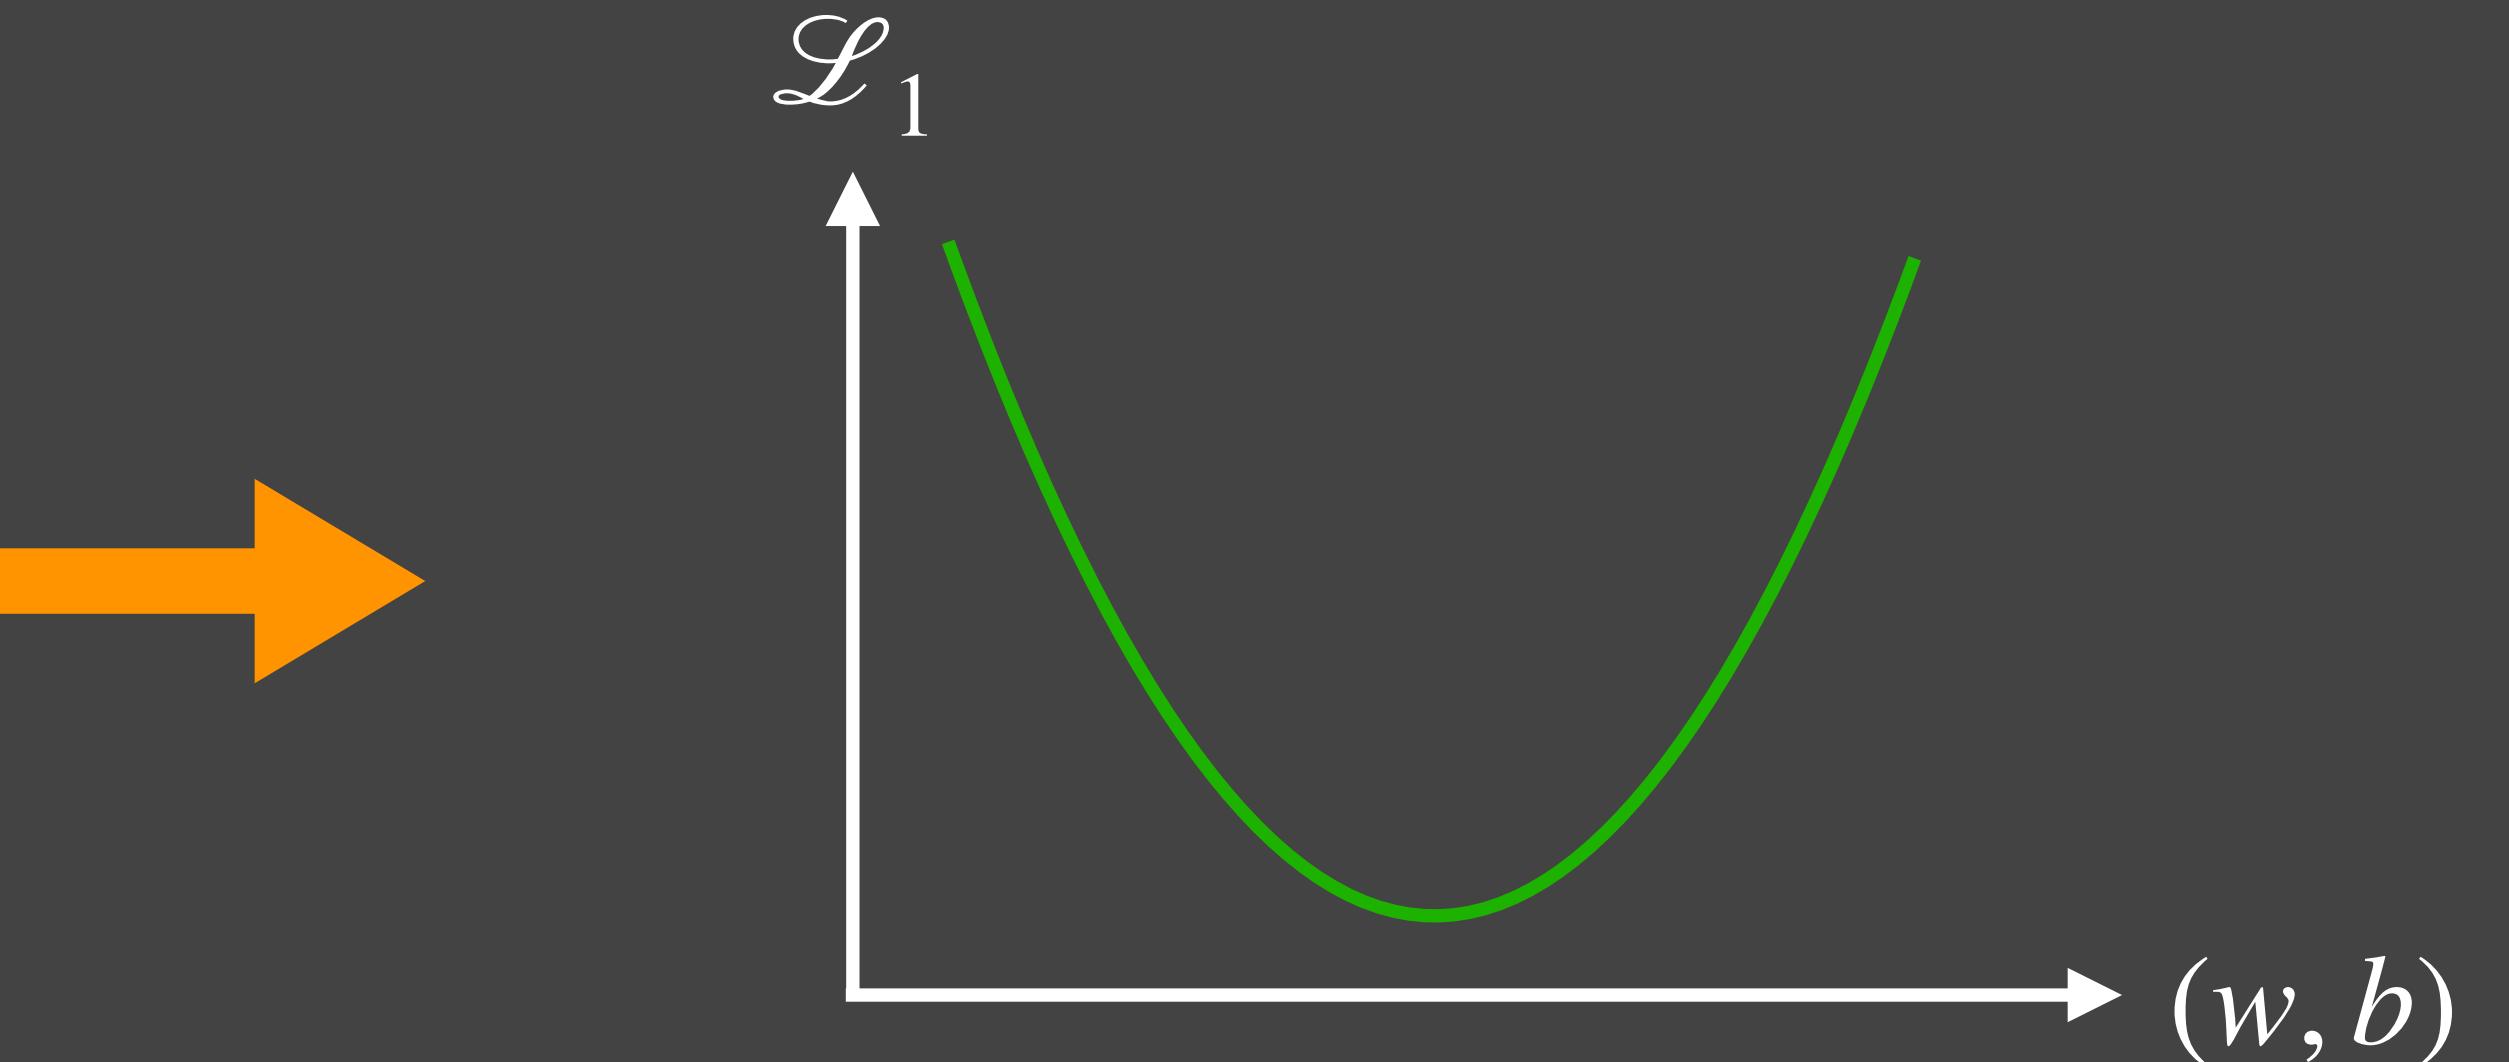
- Gradient descent optimization
  - ▶ Parabola in red, our loss function
  - ▶  $W_1$  as initial set of parameters and starting point
  - ▶ Moving along negative direction of gradient vector
  - ▶ Closer to global minima, smaller gradient vector
  - ▶ Step length to control how fast we converge

$$\hat{w}^{k+1} = \hat{w}^k - \alpha \nabla \mathcal{L}(\hat{w}^k)$$



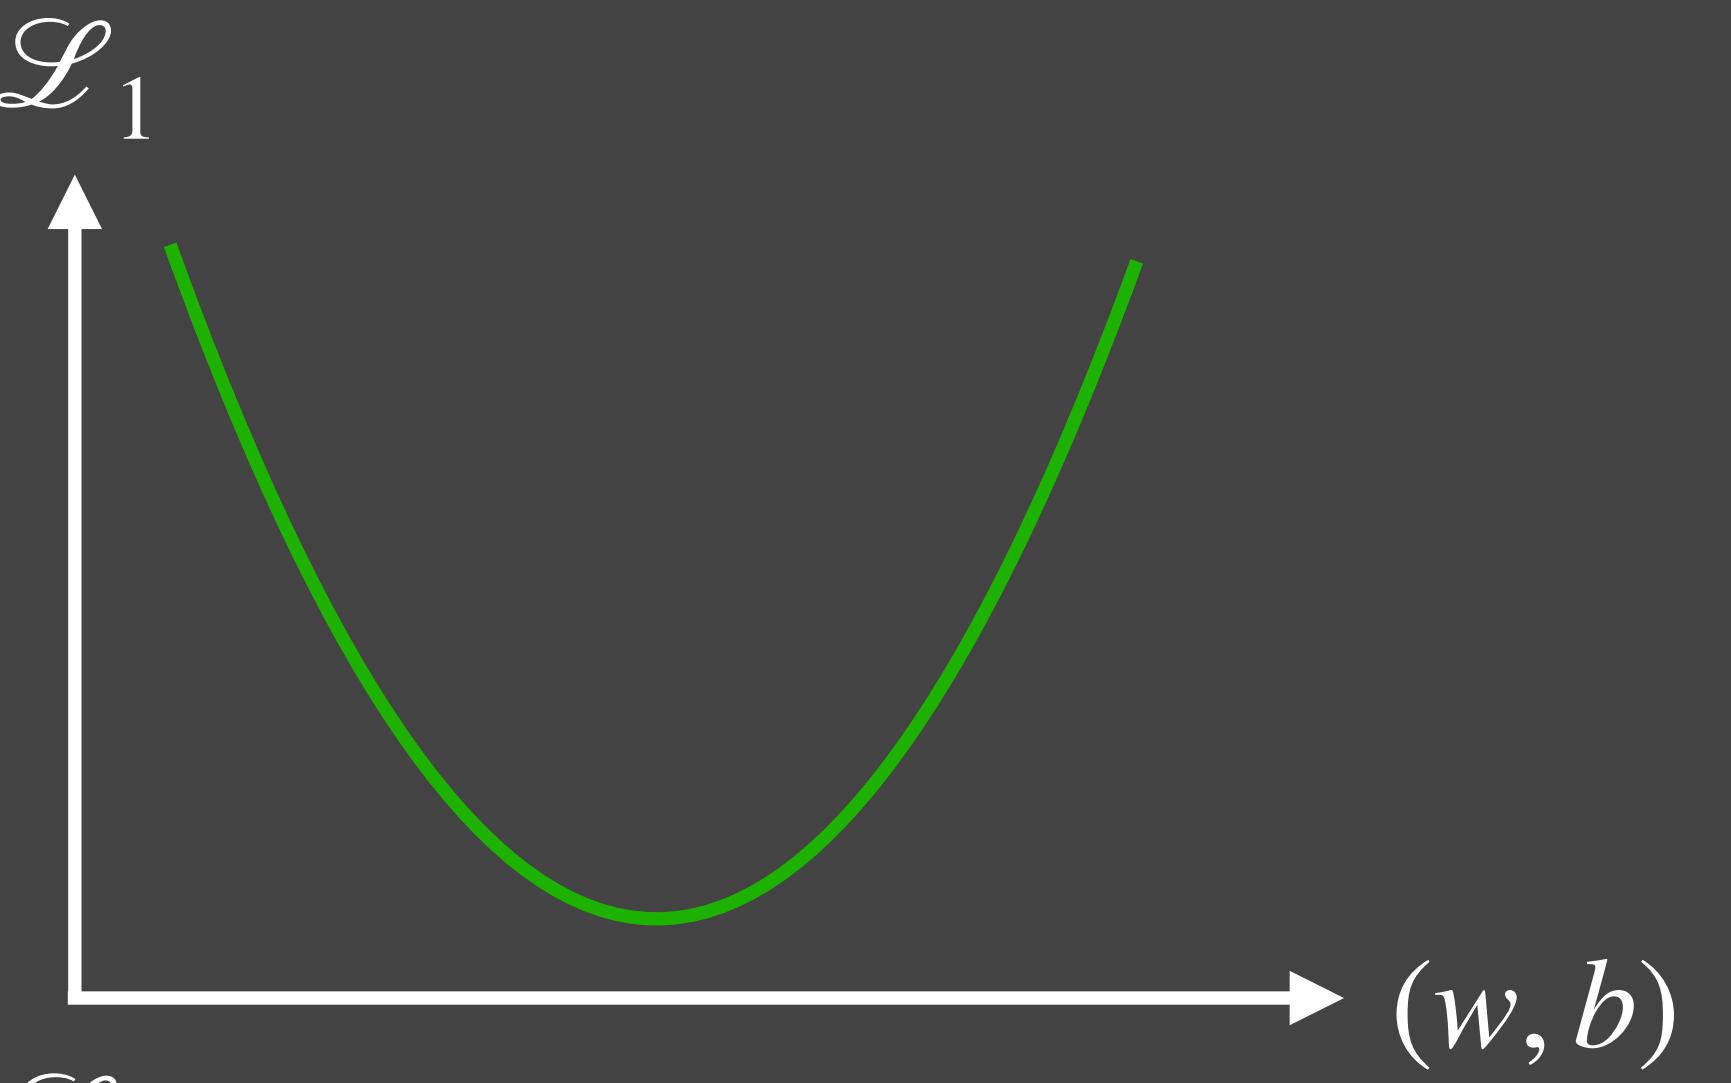
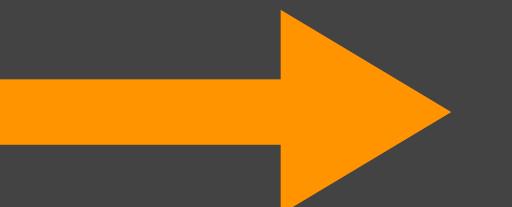
# Optimization Theory - Convexity of Loss Functions

- Unique Minima
  - Global Minima
  - Optimal parameter set
  - Convex function

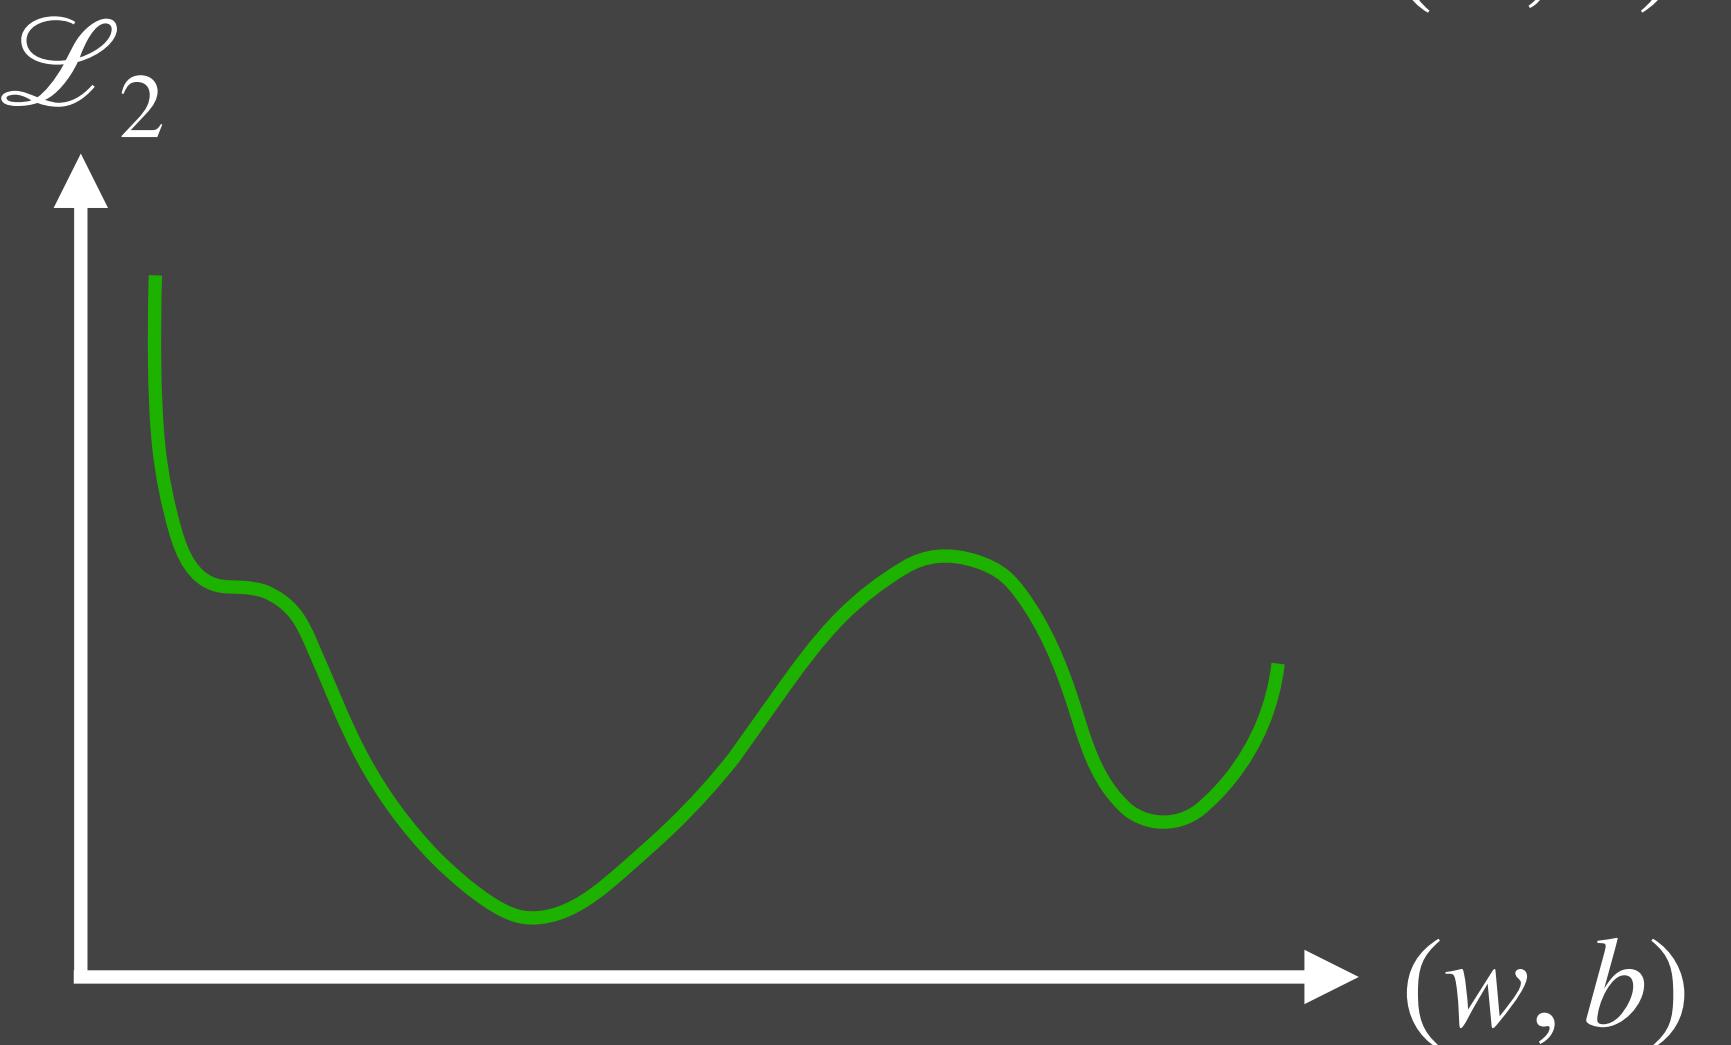
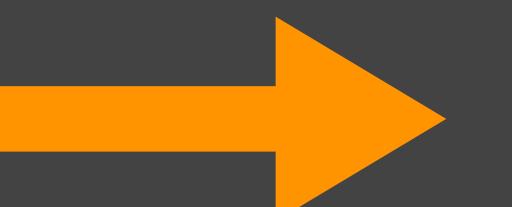


# Optimization Theory - Convexity of Loss Functions

- Unique Minima
  - Global Minima
  - Optimal parameter set
  - Convex function

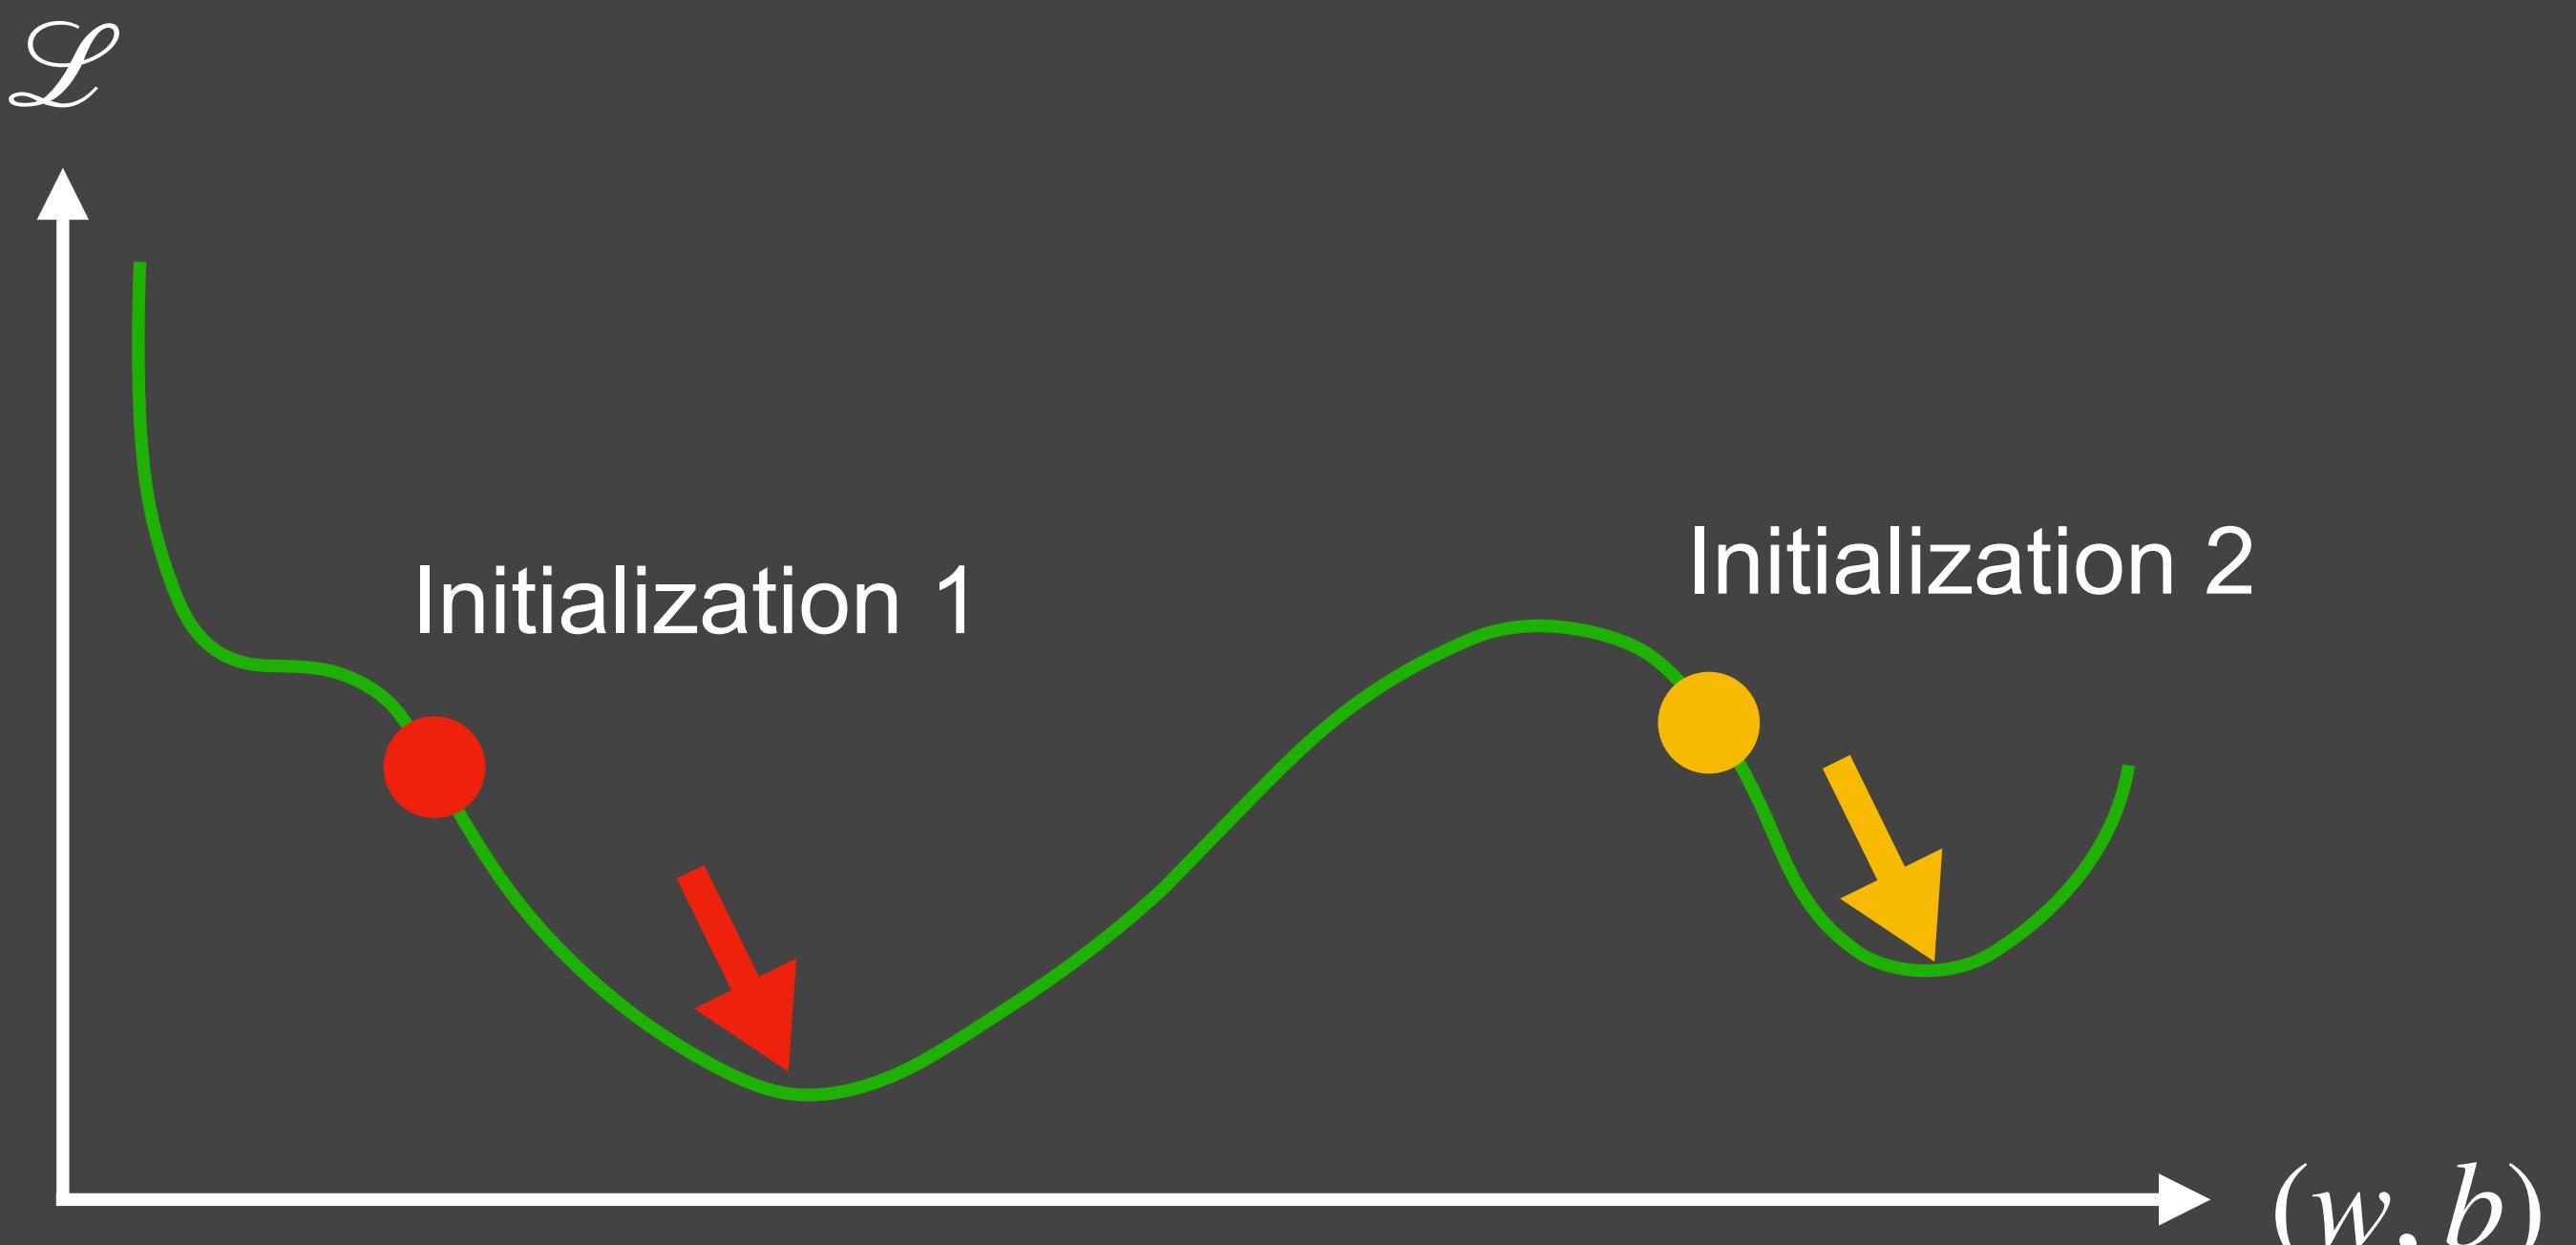


- Multiple Minimas
  - Global minima + Local minimas
  - 1 Optimal + Many sub-optimal solutions
  - Non-convex function



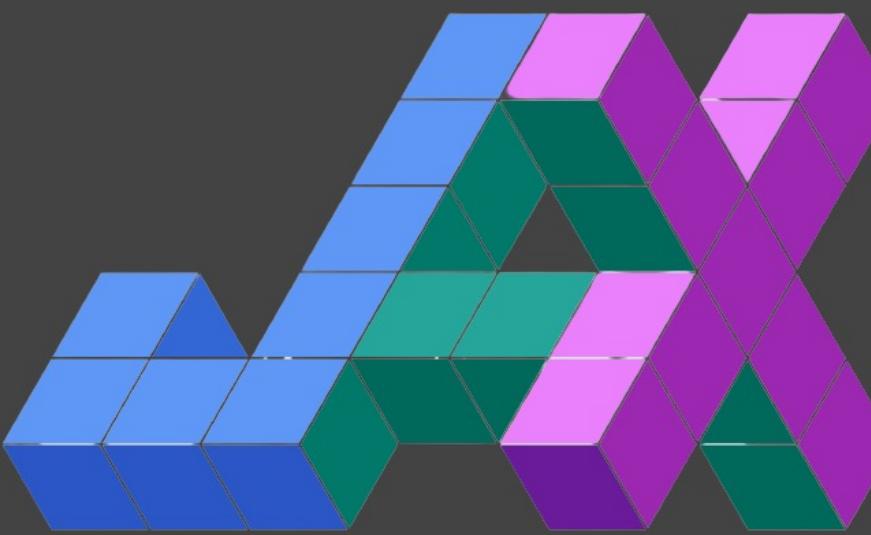
# Optimization Theory - Convexity of Loss Functions

- Gradient descent on non-convex functions
  - Solution changes depending on initialization
  - Executing algorithm multiple times
- Convex optimization problems
  - Linear regression
  - Linear classification



# JAX - Fundamentals

- Numerical computing library
  - Differentiable NumPy + Autograd
  - No usage of Tensor
  - XLA Compilation
- High level functionalities
  - Auto Vectorization
  - JIT Compilation



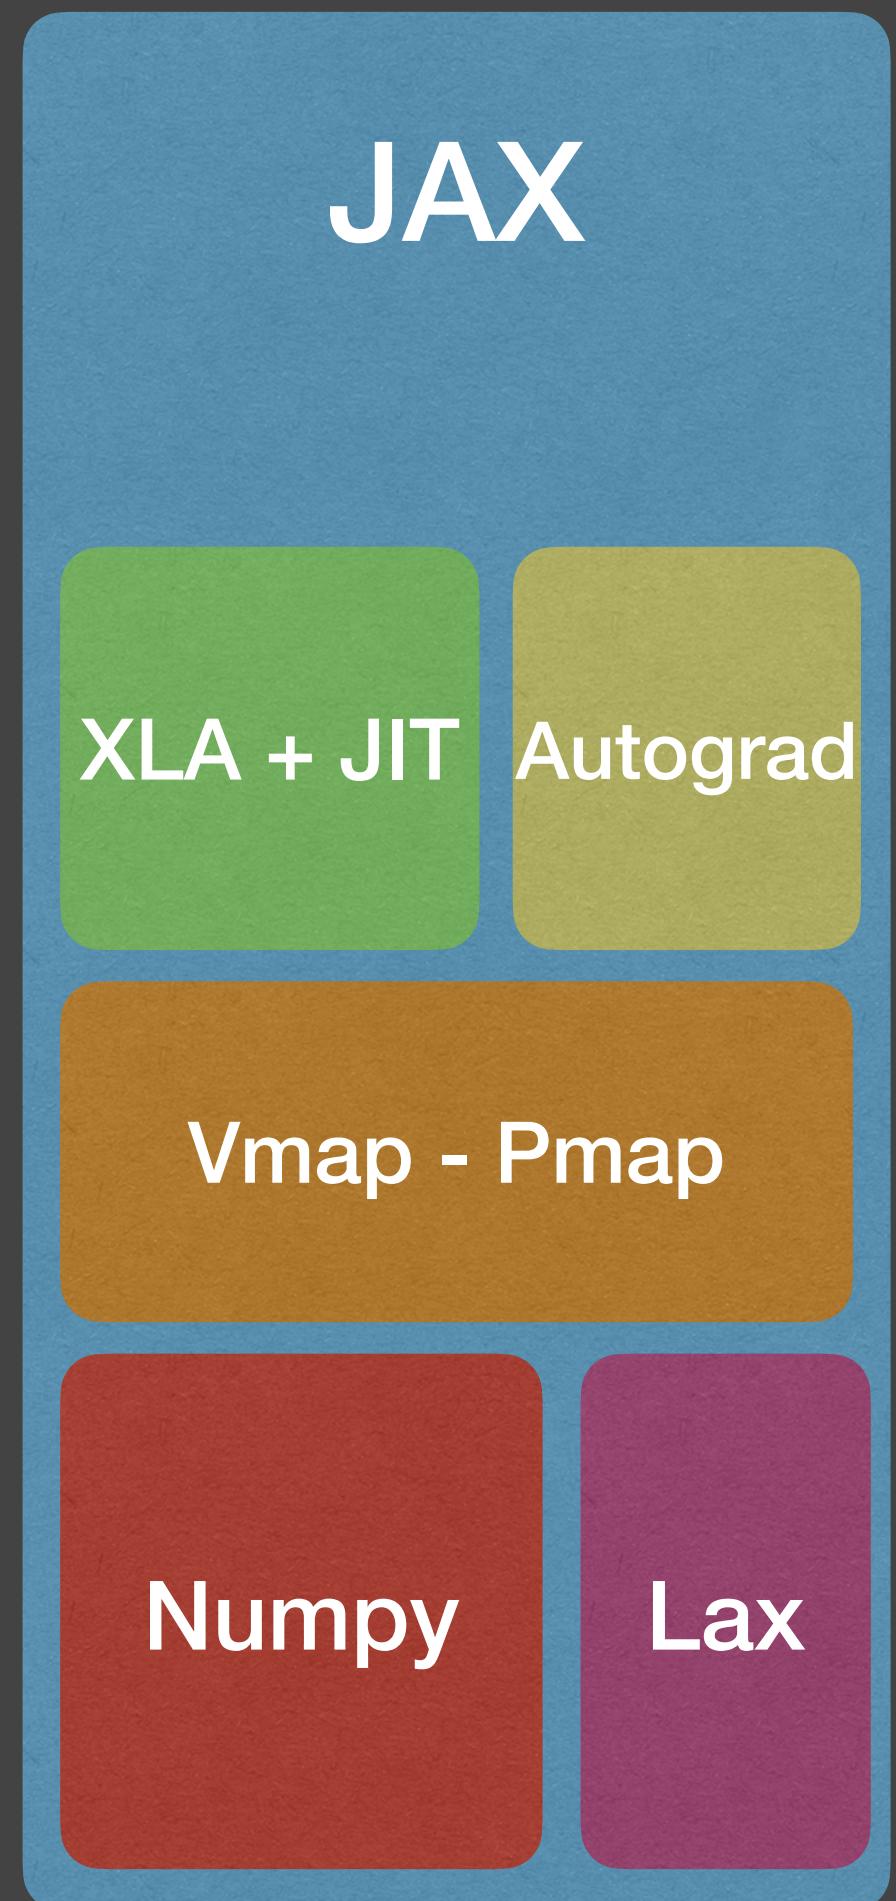
# JAX - Numpy

## JAX - Numpy

- Embedded in JAX
- Immutable
- Accelerator Support
- Vmap and JIT Support

## Standard Numpy

- Standalone Library
- Mutable
  - No GPU or TPU
  - No Vmap or JIT



# JAX - Composable Function Transformations

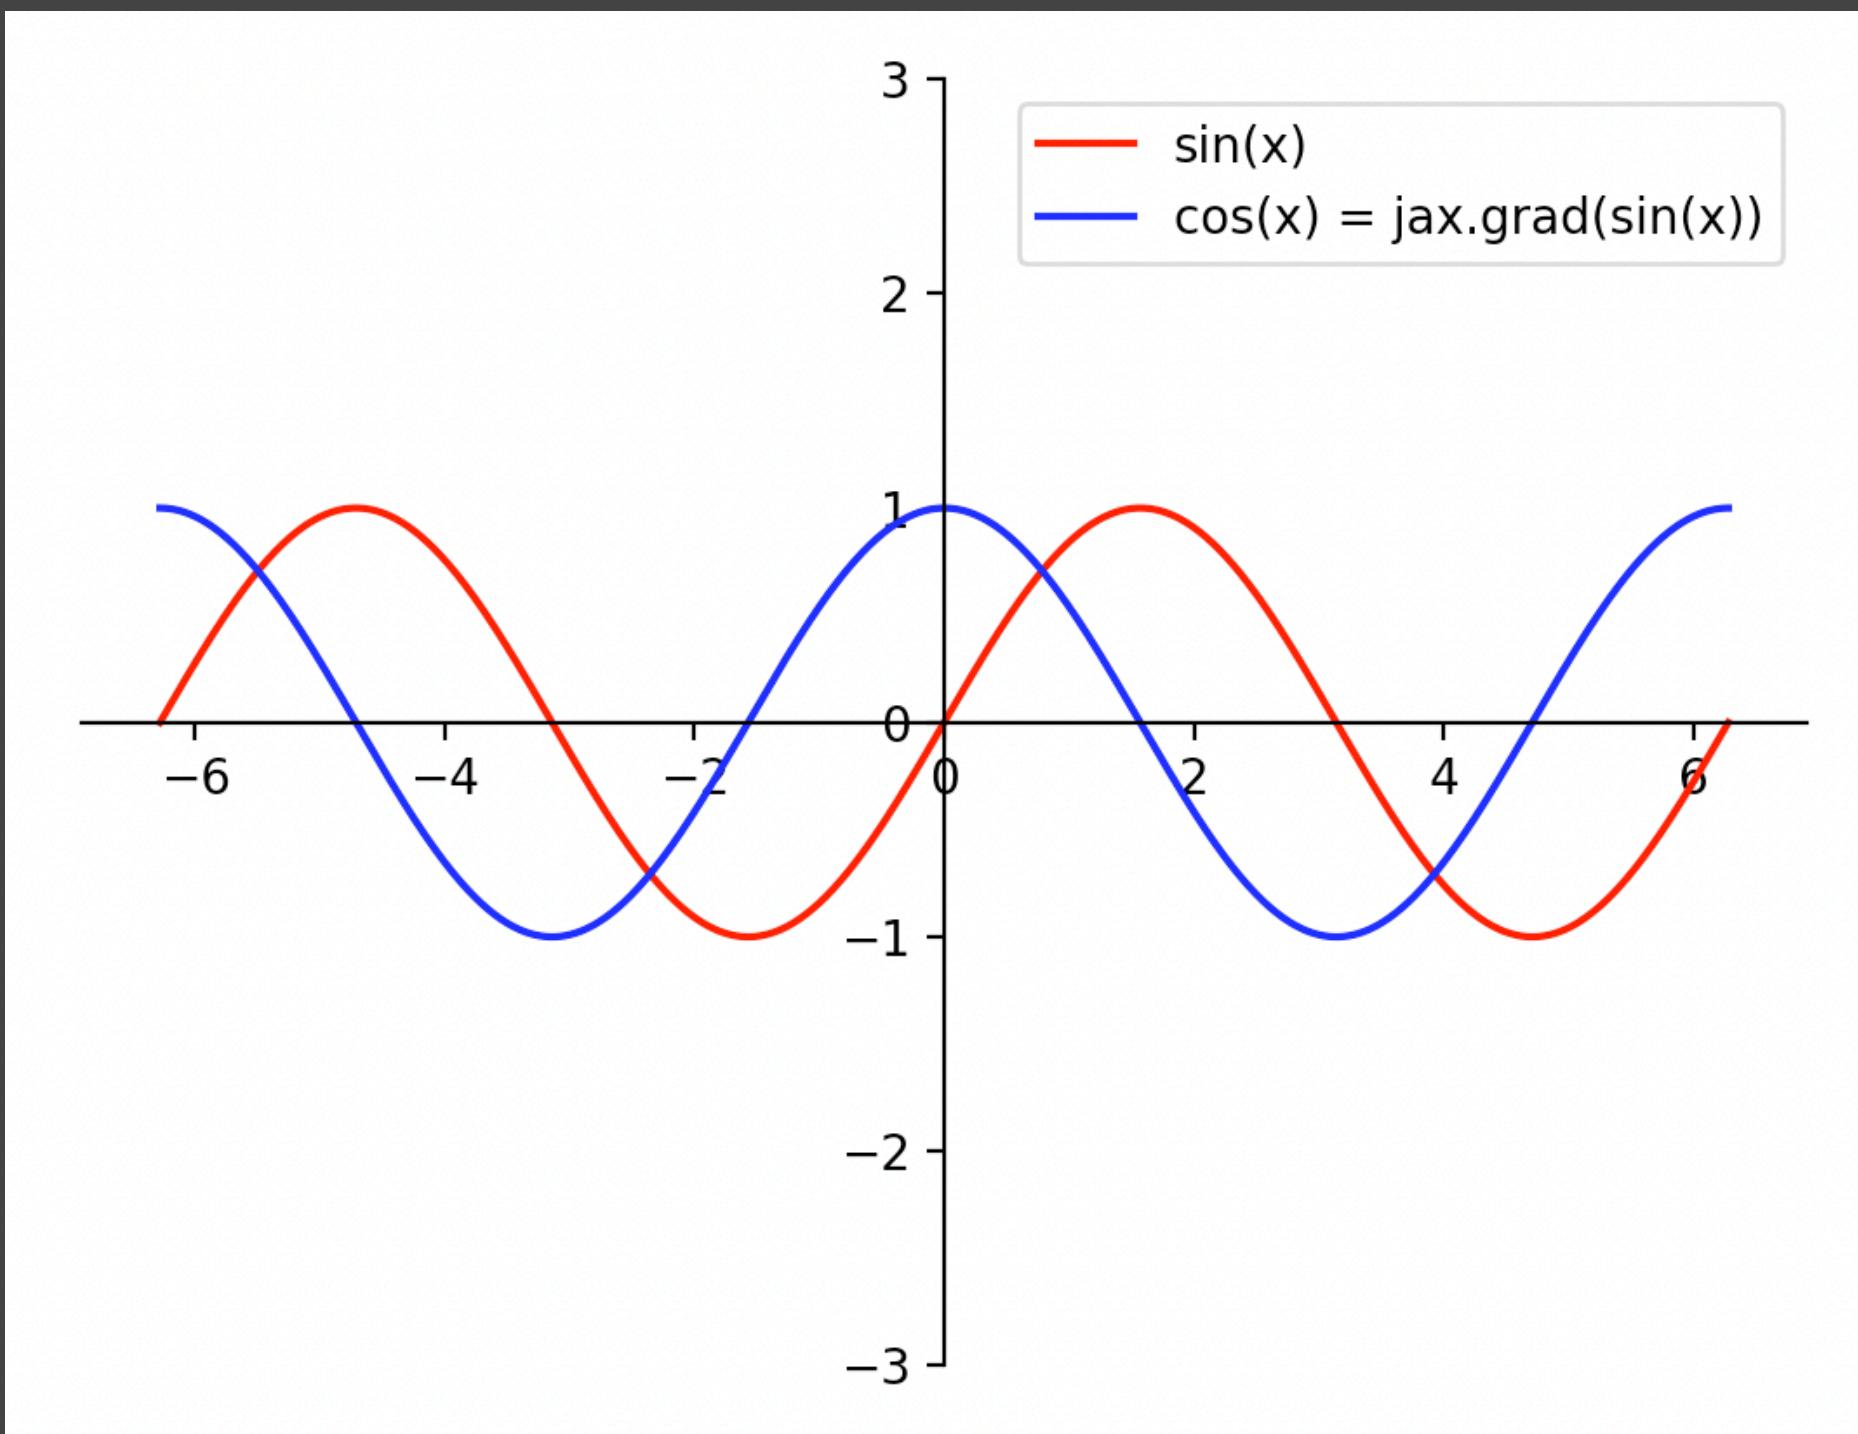
- Auto-Gradient
  - Differentiating a python function
  - Returning derivative as another function
  - `jax.grad( . )`

$$f(x) = \sin(x)$$

$$\frac{df(x)}{dx} = \cos(x)$$

$$\text{jax.grad} \approx \frac{d}{dx}$$

```
x = 0
y = jax.numpy.sin(x)
cos = jax.grad(jax.numpy.sin)
y = cos(x)
```



# JAX - Composable Function Transformations

- Auto-Gradient

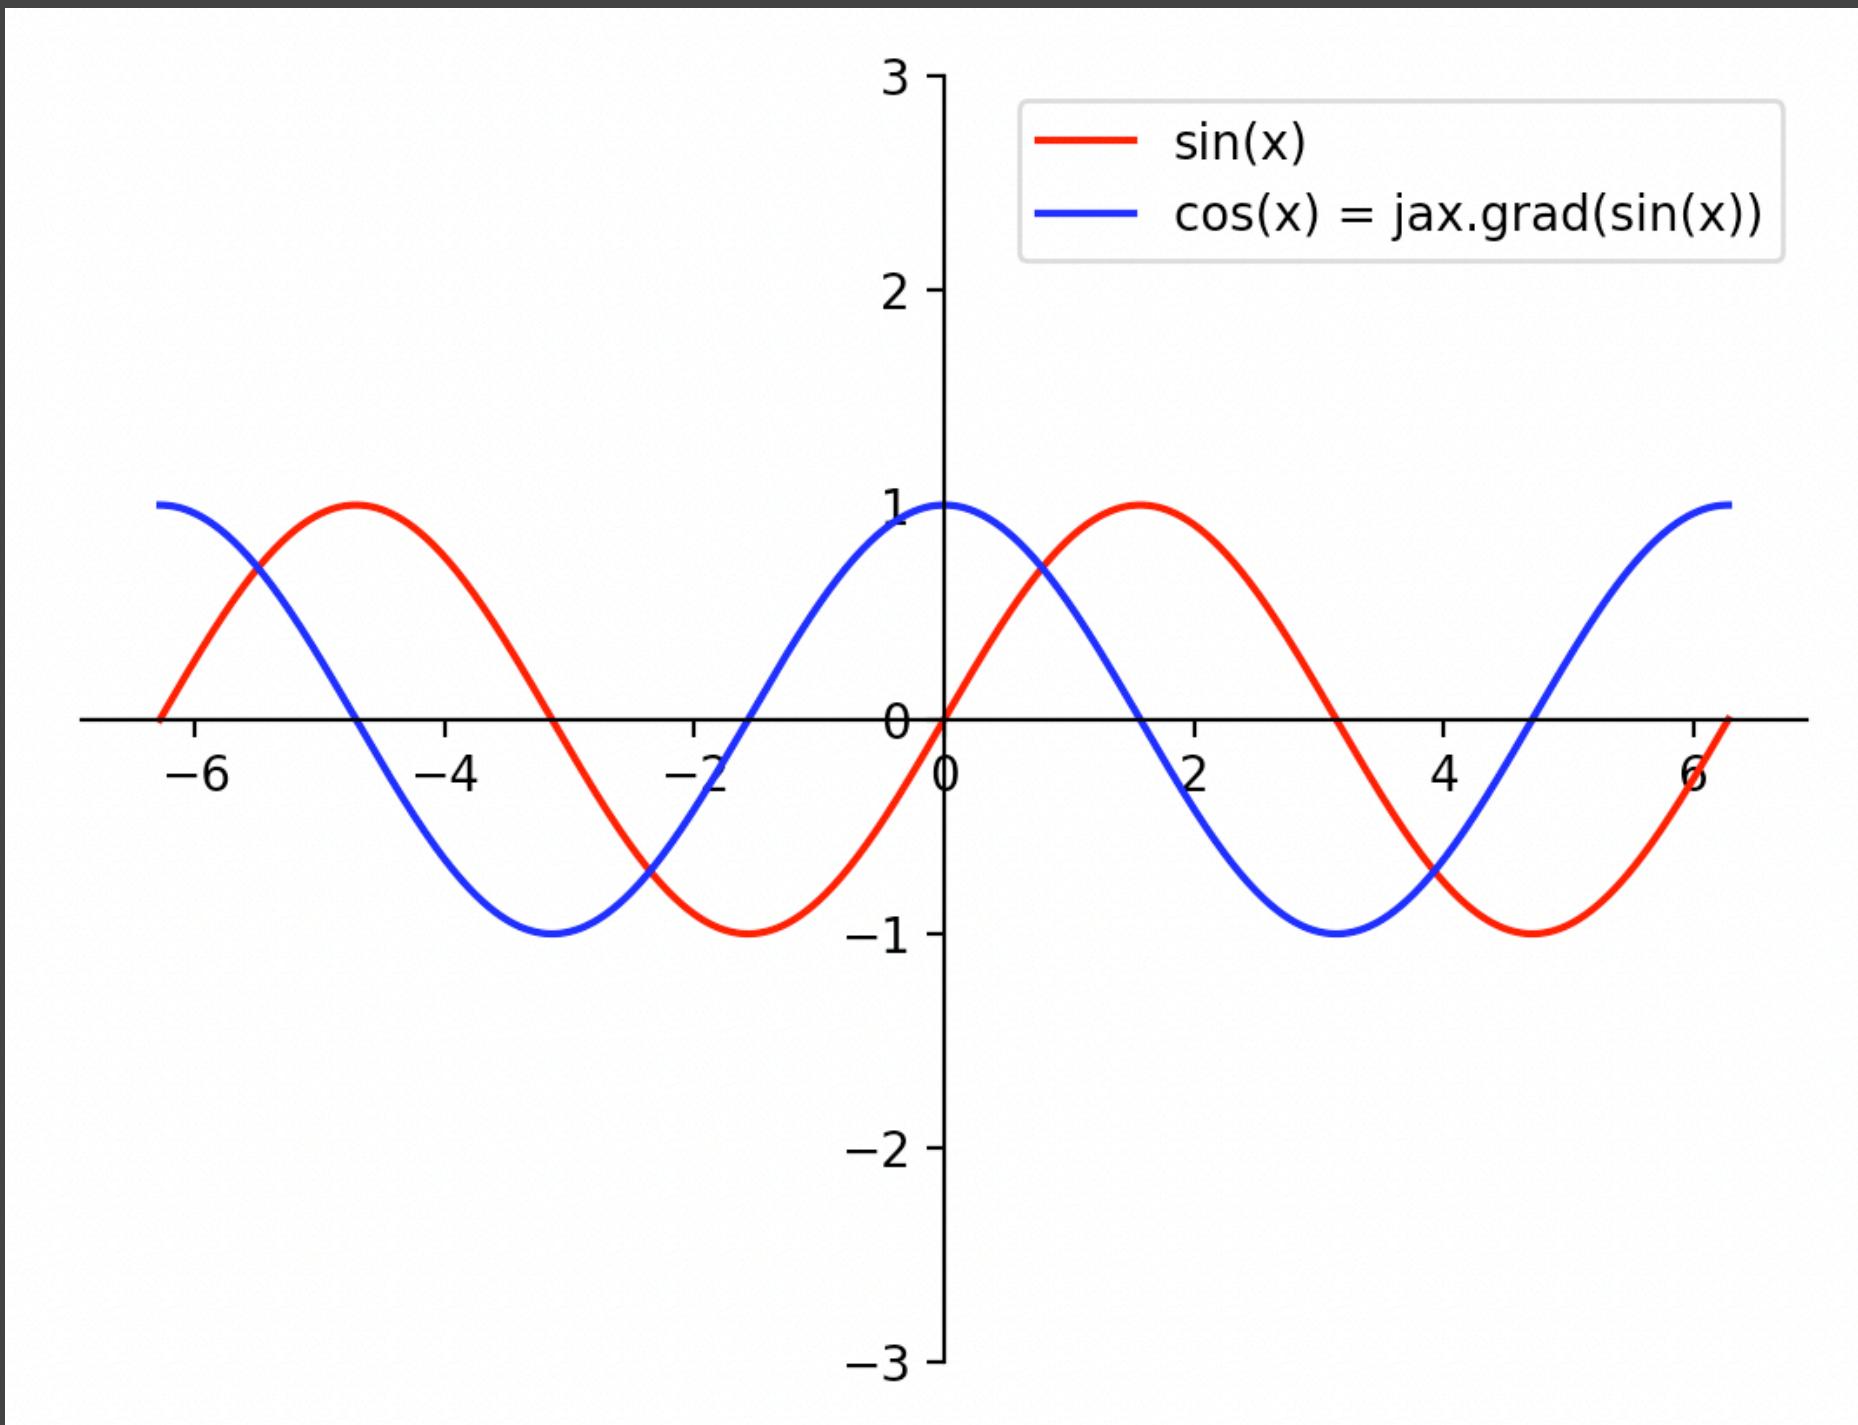
- Differentiating a python function
- Returning derivative as another function
- $\text{jax.grad}(\cdot)$

$$f(x) = \sin(x)$$

$$\frac{df(x)}{dx} = \cos(x)$$

$$\text{jax.grad} \approx \frac{d}{dx}$$

```
x = 0
y = jax.numpy.sin(x)
cos = jax.grad(jax.numpy.sin)
y = cos(x)
```



# JAX - Composable Function Transformations

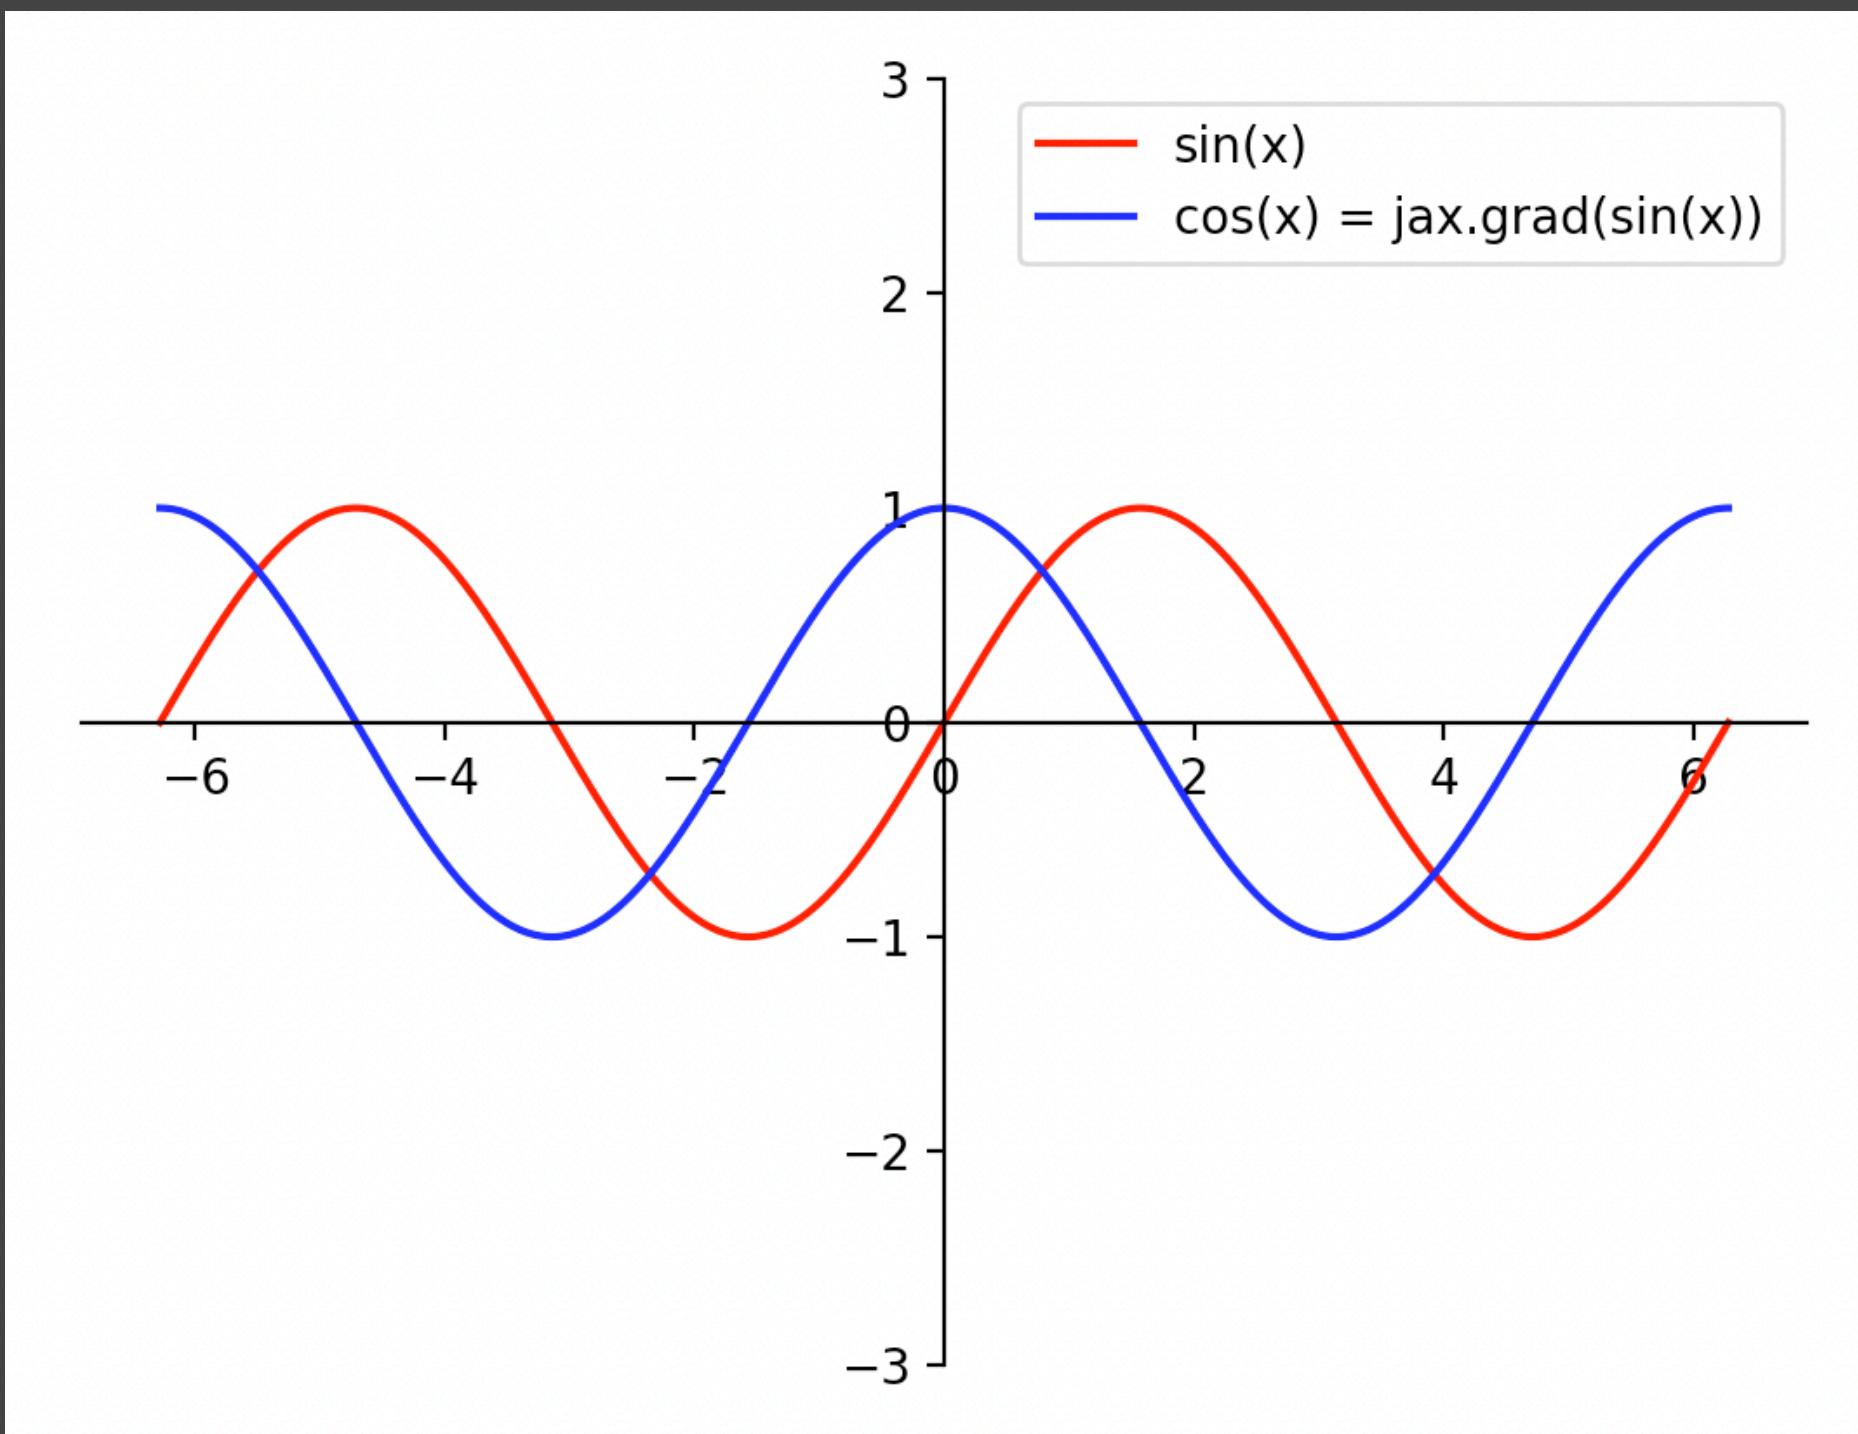
- Auto-Gradient
  - Differentiating a python function
  - Returning derivative as another function
  - `jax.grad( . )`

$$f(x) = \sin(x)$$

$$\frac{df(x)}{dx} = \cos(x)$$

$$jax.grad \approx \frac{d}{dx}$$

```
x = 0
y = jax.numpy.sin(x)
cos = jax.grad(jax.numpy.sin)
y = cos(x)
```



# JAX - Composable Function Transformations

- Auto-Vectorization

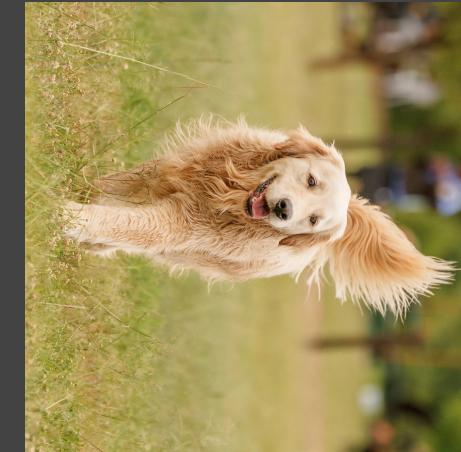
- Rotating multiple images at the same time
- Different rotation angle for each image
- *jax.vmap( . )*

```
vec_rotate = jax.vmap(pix.rotate, in_axes=(0, 0), out_axes=0)
rotated_imgs = vec_rotate(imgs, radians)
```

```
rotated_img = pix.rotate(img, radian)
```



90°



90°



180°

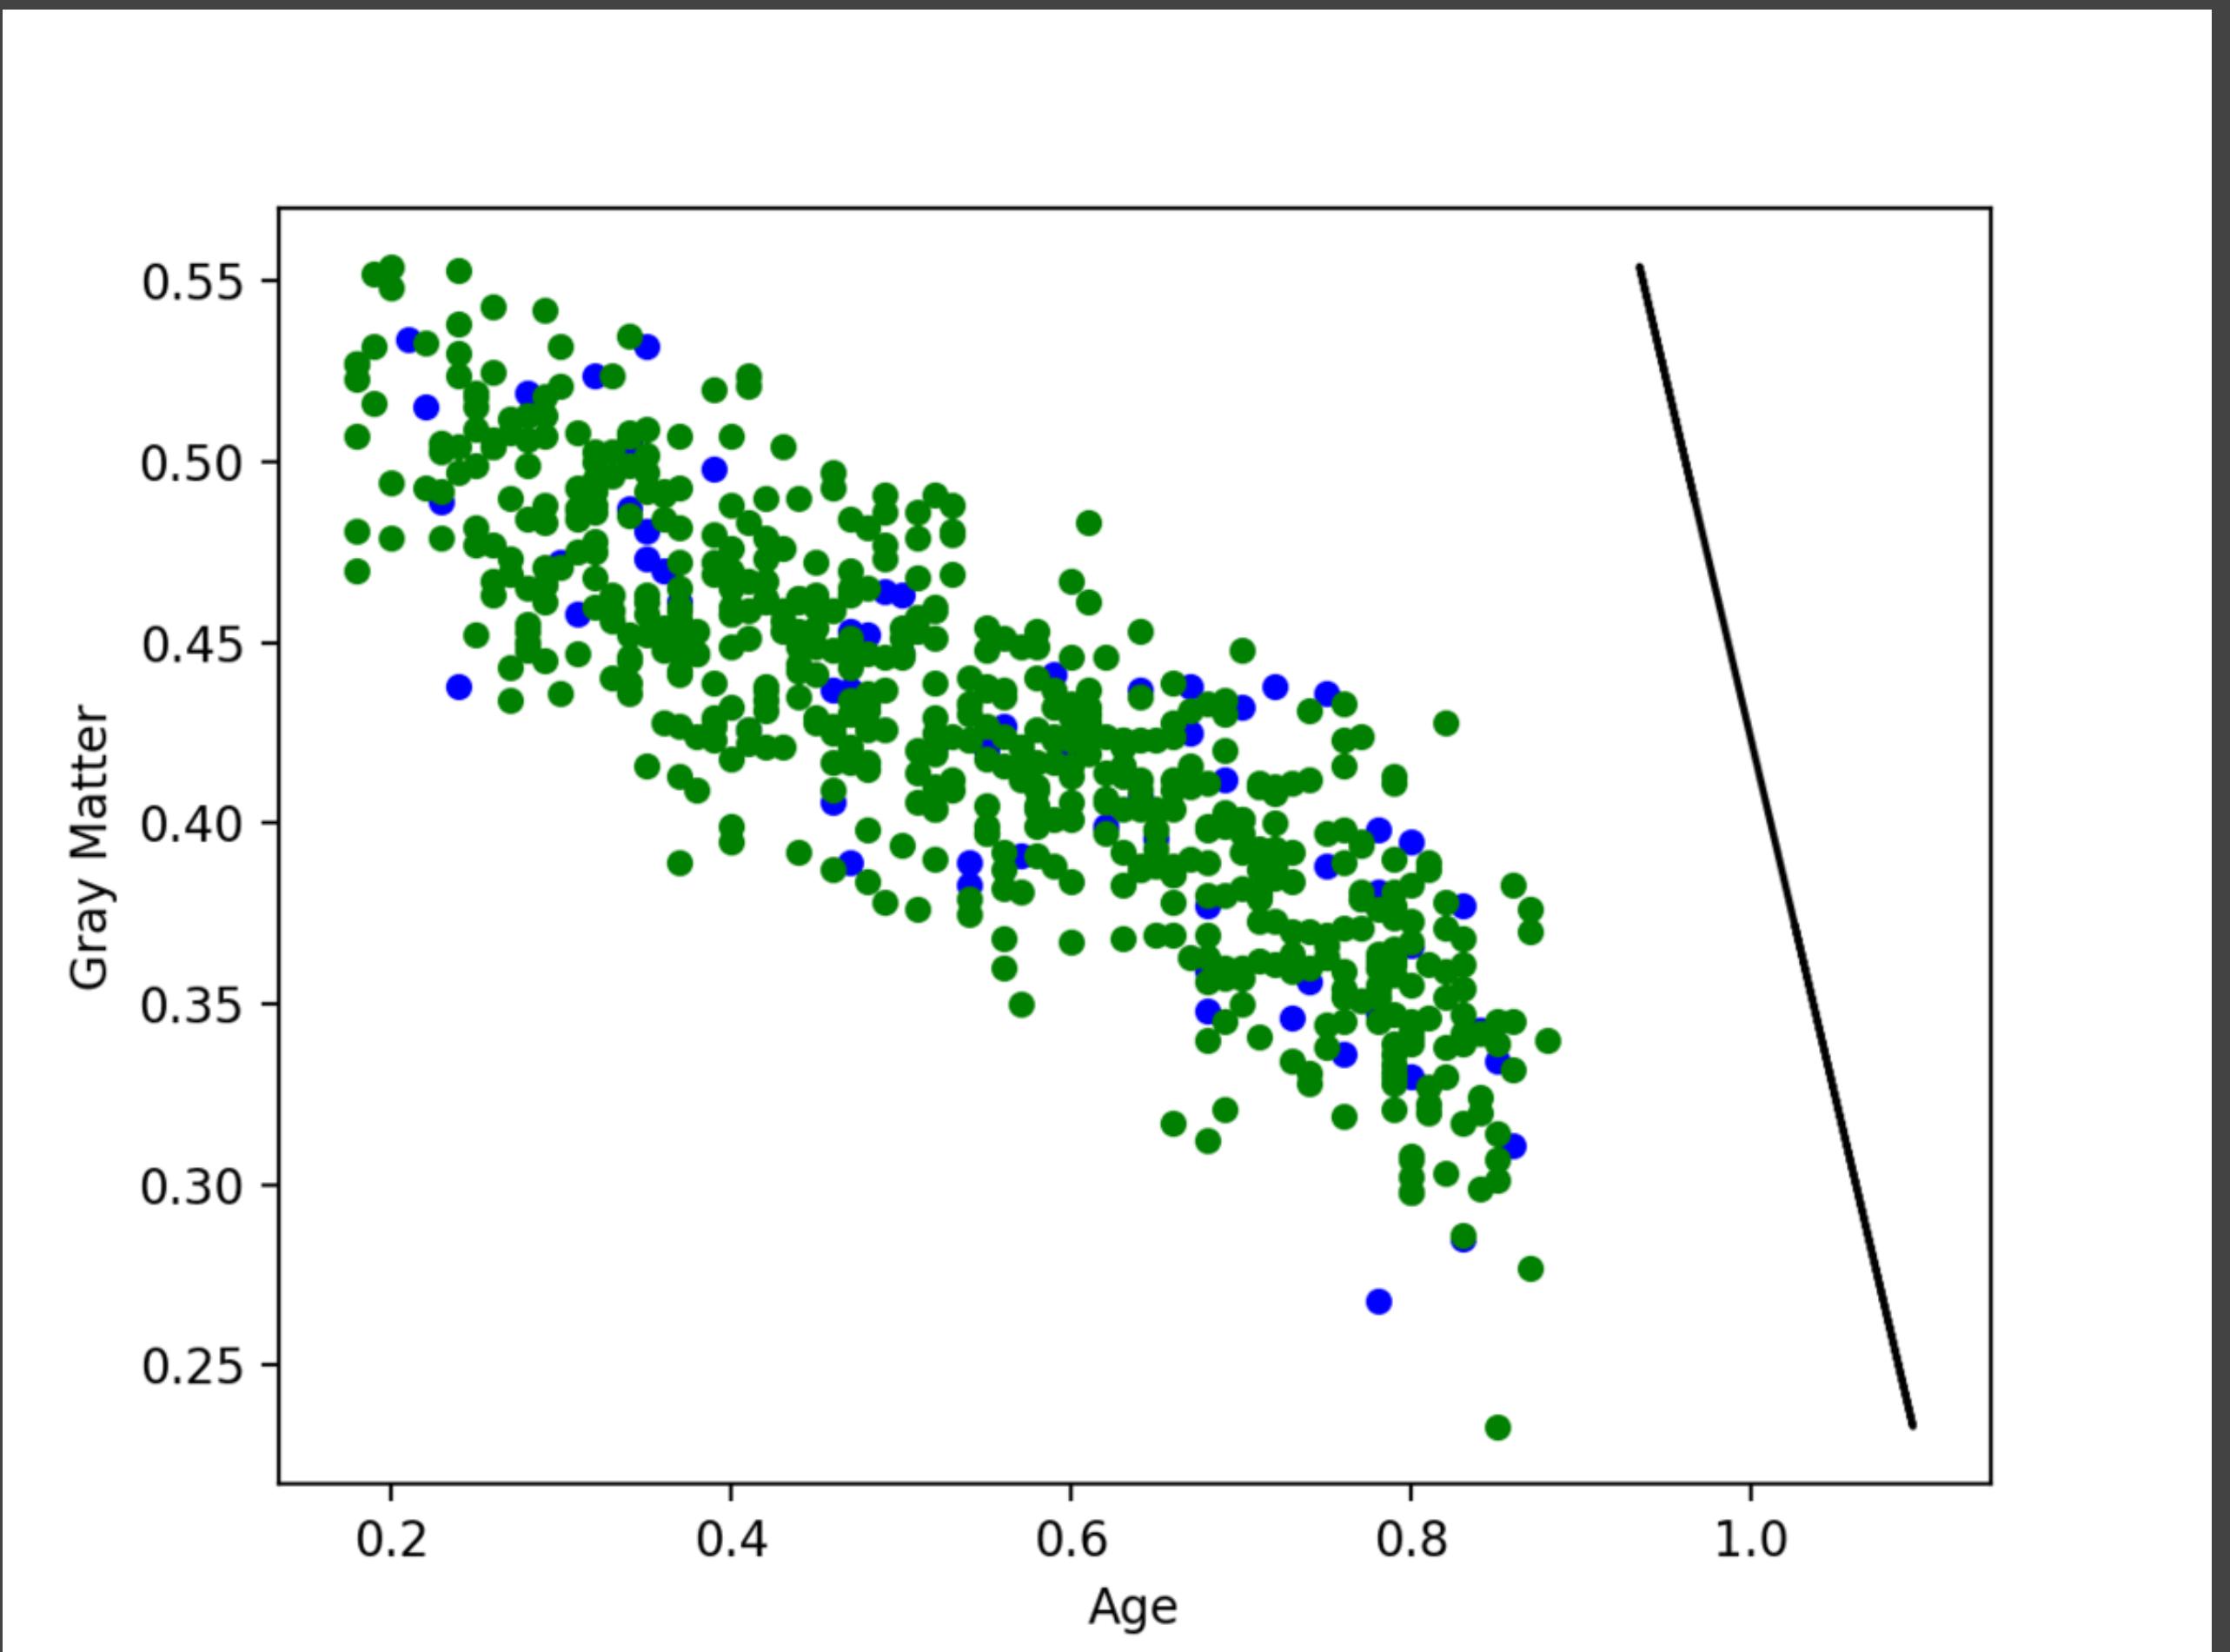


270°



# Coding Session - Brain Age Prediction

- Data Representation
  - Only 1 feature → Gray Matter
  - Samples → Training Data
  - Samples → Test Data
  - Older people with less gray matter
- Training Process
  - 15000 number of iterations
  - Min MSE Loss: 0.0091
  - Max R2 Score: 0.7344
- Testing Process
  - Min MSE Loss: 0.0105
  - Max R2 Score: 0.6980



# References and Resources

- Additional JAX Resources and Tutorials
  - <https://www.kaggle.com/code/goktugguvercin/introduction-to-jax>
- Coding Session - Resource Colab File and Dataset
  - <https://drive.google.com/drive/folders/1xbc07s6rLYEcb4Iy0MP8diB3U39uvfoz?usp=sharing>
  - <https://github.com/GoktugGuvercin/Machine-Learning-and-Optimization-Theory>